

Computer Vision and Deep Learning

Lecture 4

Today's agenda

- Short history of of neural networks
- Convolutional neural networks
- How to evaluate a classifier

Classifier evaluation

Train, dev and test sets

- Training set
 - Used to train the model, determines what the network learns
- Development (*dev*) set or validation set
 - Used to evaluate the performance of your models and determine which ones work best
- Test set
 - Used to get an unbiased estimate of the final performance of the model

Sometimes it might be ok to not have a test set (only train and dev sets)

How to split your data into train/dev/test set?

- Before deep learning, when available data was relatively limited (e.g. 10000 images)

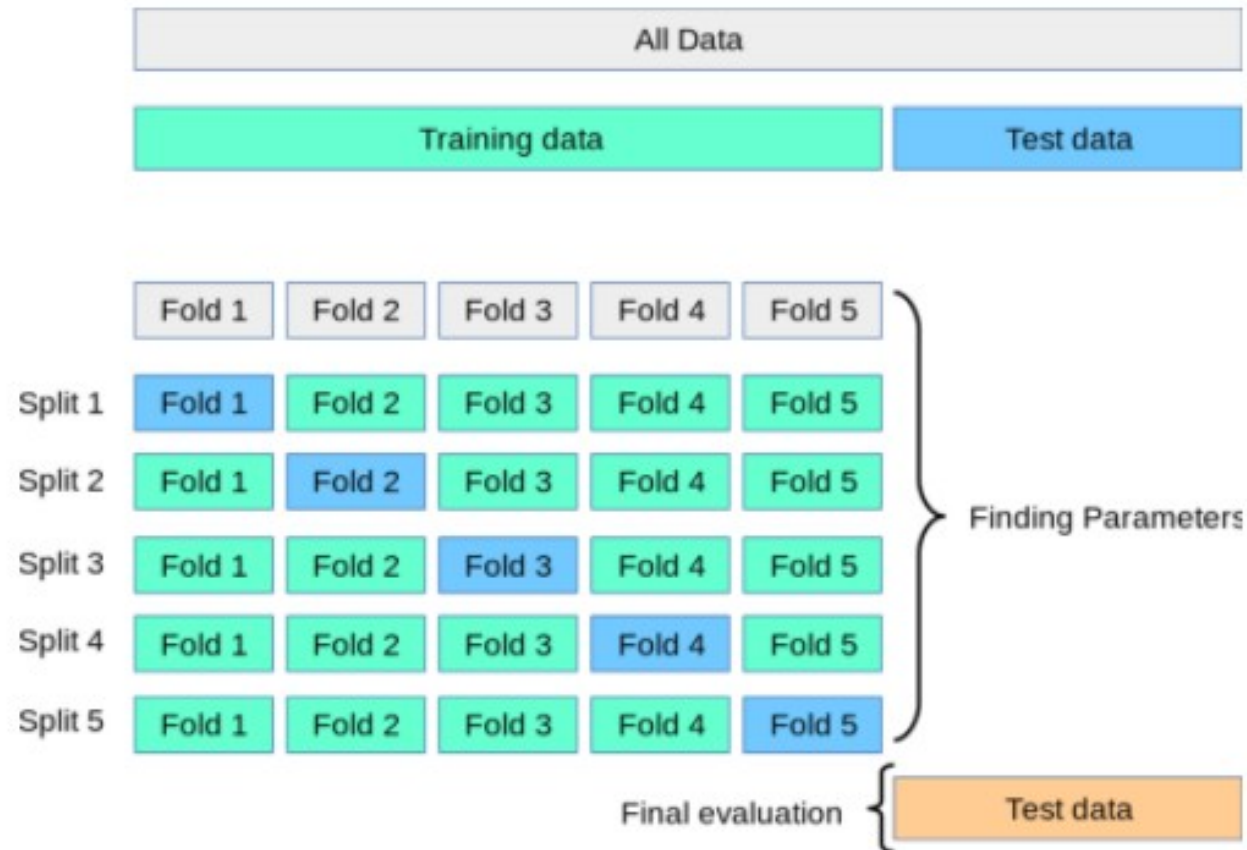


7000 train, 3000 dev




6000 train, 2000 dev,
2000 test

K-fold validation



Stratified k-fold validation




How to split your data into train/dev/test set?

- Large scale datasets
 - ImageNet > 14 million images
 - VGGFace 2 > 3.3 million images
 - JFT > 300 million images
- E.g. 1 million images
 - 980000 train set, 10000 dev set, 10000 test set -> 98% train, 1% dev, 1% test
- E.g. 4 million images
 - 3980000 train set, 10000 dev set, 10000 test set -> 99.5% train, 0.25% dev, 0.25% test

Convolutional neural networks

Understanding the visual cortex



Hubel and Wiesel, 1959

https://www.youtube.com/watch?v=IOHayh06LJ4&ab_channel=PaulLester



Nobel Prize for Physiology or Medicine in 1981:
David Hubel and Torsten Wiesel

Simple cells:

orientation, position

Complex cells:


orientation, motion, direction

“Hypercomplex” cells:

orientation, motion, direction, length

Understanding the visual cortex

- Nearby cells in the cortex represented and processed nearby regions in the visual field
- Hubel and Wiesel hypothesized that the visual cortex can be described by a hierarchical organization of simple cells that fed into complex cells which have more complicated activations and can form higher level representations




Backpropagation, 1986

Learning representations by back-propagating errors

David E. Rumelhart*, **Geoffrey E. Hinton†**
& **Ronald J. Williams***


* Institute for Cognitive Science, C-015, University of California,
San Diego, La Jolla, California 92093, USA

† Department of Computer Science, Carnegie-Mellon University,
Pittsburgh, Philadelphia 15213, USA



Gradient-Based Learning Applied to Document Recognition, *Yann LeCun, Leon Bottou, Yoshua Bengio, and Patrick Haffner, 1998*

1989 – the original form of LeNet




40004 75216
14189-2087 23505
96203 14310
44151 05153

Deep Big Simple Neural Nets Excel on Handwritten Digit Recognition

Dan Ciresan, 2010

- One of the very first implementations of GPU Neural nets
 - forward and backward pass implemented of an artificial neural network (up to 9 layers) implemented on an NVIDIA GTX 280 graphic processor



2012 ImageNet Classification with Deep Convolutional
Neural Networks, *Alex Krizhevsky, Ilya Sutskever, Geoffrey E.
Hinton*



Convolutional neural networks

Convolutional layers


- Preserve the spatial information
- Convolve the filter over the entire input volume
 - The filter has the **same depth** as the input



filter
5x5x3




Activation map
28x28x1



Convolutional layers


- Preserve the spatial information
- Convolve the filter over the entire input volume
 - The filter has the **same depth** as the input



At each position in the input volume:
Multiply (element-wise, across all channels)
the filter and a small patch (5x5x3) in this
input volume and add a bias term

Convolutional layers

- Preserve the spatial information
- Convolve the filter over the entire input volume
 - The filter has the **same depth** as the input



Input volume
32x32x3



k filters



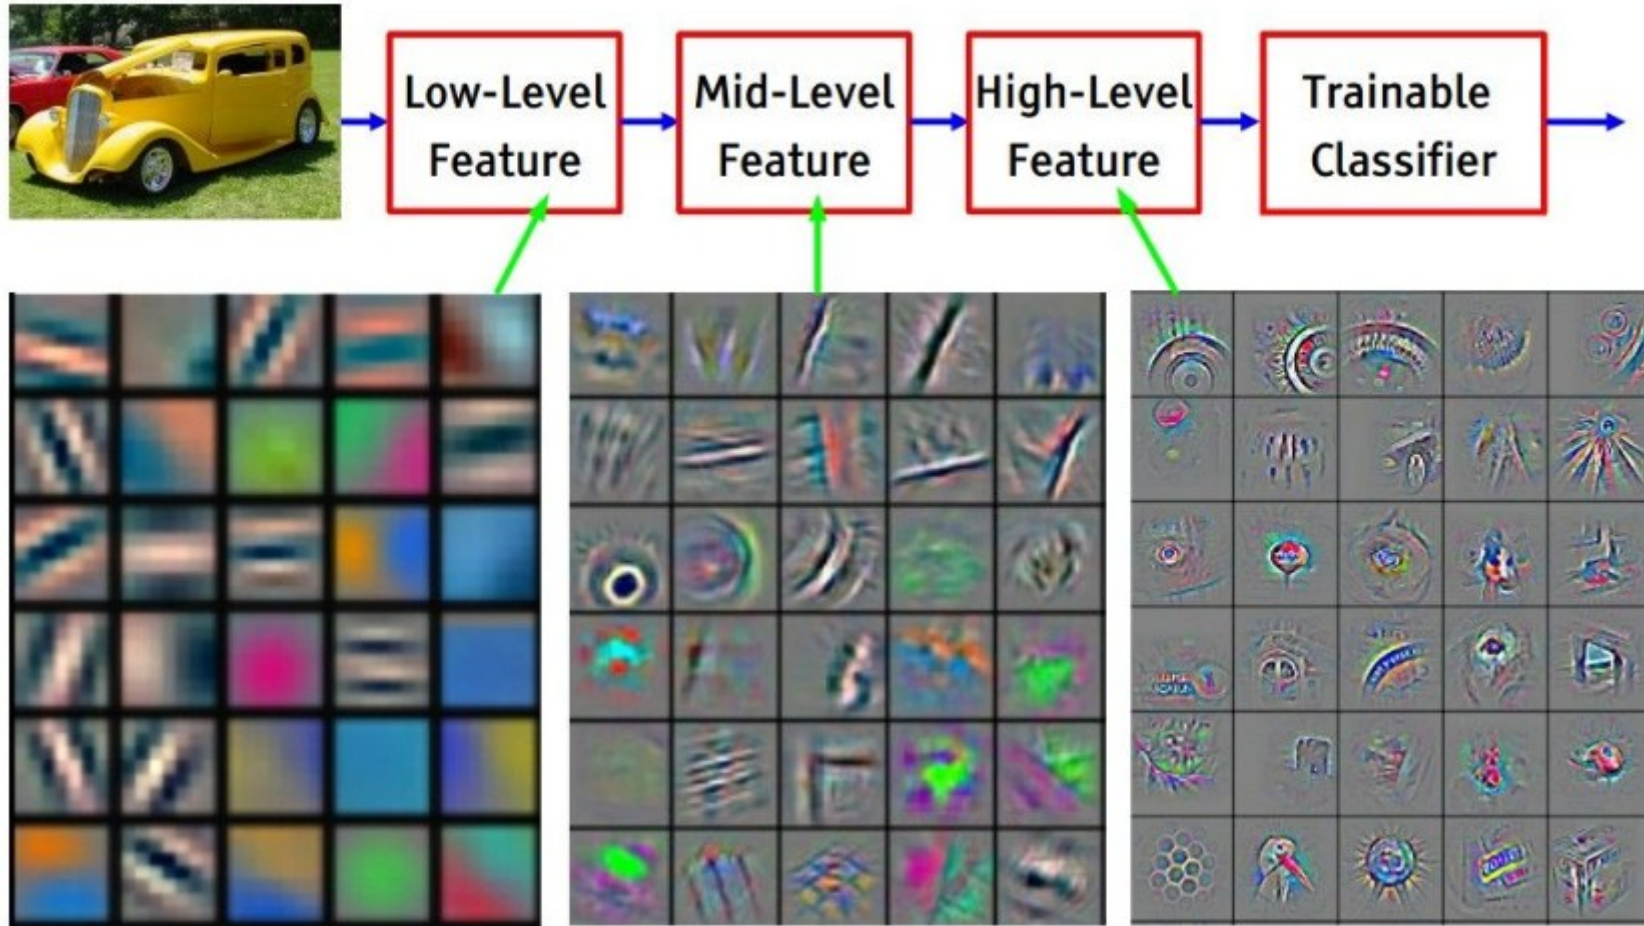
k activation maps
28x28x1

Convolutional layers

- Neurons in an activation map:
 - Each neuron is connected to a small region in the input
 - All of neurons in the activation map share parameters
- Receptive field of a neuron
- k filters \rightarrow k different neurons all looking at the same region in the input volume



k activation maps
28x28x1






https://medium.com/@chriskevin_80184/feature-maps-ee8e11a71f9e

<https://www.youtube.com/watch?v=AgkflQ4IGaM>

Convolutional layers

Stride


- Stride – the amount by which the filter shifts
- Stride 1



Convolutional layers

Stride

- Stride – the amount by which the filter shifts
- Stride 2



Convolutional layers

Padding

0	0	0	0	0	0	0
0						0
0						0
0						0
0						0
0						0
0	0	0	0	0	0	0

padding: 1

Convolve an input of spatial size 7x7 with a 3x3 filter : output spatial size?

Convolve an input of spatial size 7x7 with a 3x3 filter, and applying 0 padding to the input: output spatial size?

To preserve size spatially:

CONV layers with stride 1, filters of size $F \times F$, and zero-padding with $(F-1)/2$

Why convolutions

- Parameter sharing
 - A feature detector that's useful in one part of an image is probably useful in other parts of the image
 - Translation invariance
- Sparsity of connections
 - The output volume depends only on a small (filter size) subset of the input volume

Convolutional layers

$(\ , \)(\ , \)$

= filter depth

– input width, – output width

– input height, – input height

F – filter size

P – padding

S - stride

$$W_o = \frac{W_I - F + 2P}{S} + 1$$

$$H_o = \frac{H_I - F + 2P}{S} + 1$$

Conv2D layer

Conv2D class

```
tf.keras.layers.Conv2D(  
    filters,  
    kernel_size,  
    strides=(1, 1),  
    padding="valid",  
    data_format=None,  
    dilation_rate=(1, 1),  
    groups=1,  
    activation=None,  
    use_bias=True,  
    kernel_initializer="glorot_uniform",  
    bias_initializer="zeros",  
    kernel_regularizer=None,  
    bias_regularizer=None,  
    activity_regularizer=None,  
    kernel_constraint=None,  
    bias_constraint=None,  
    **kwargs  
)
```



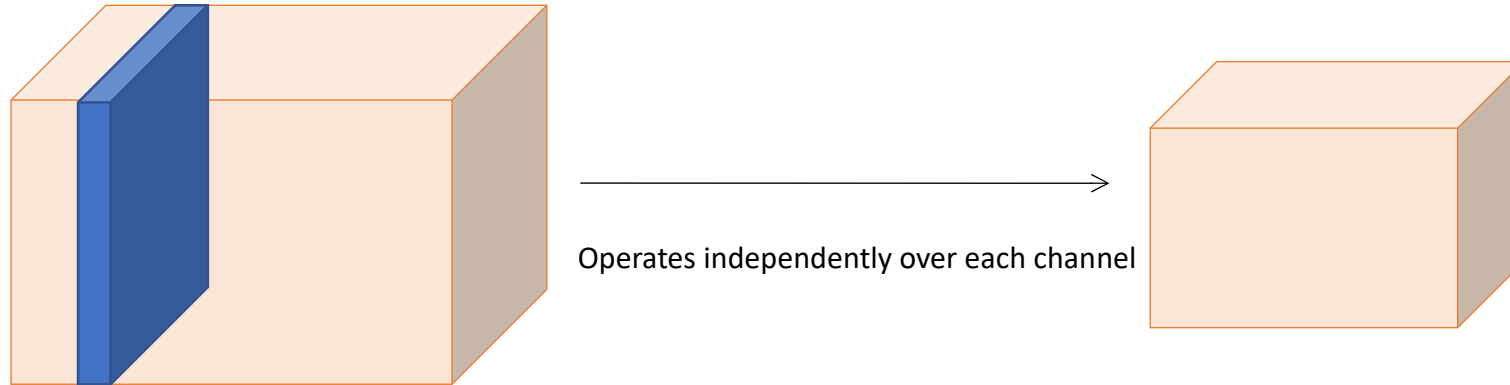
Padding: "same" or "valid"

2D convolution layer (e.g. spatial convolution over images).


Pooling layers

- Operate **independently** over each channel in the input
- Reduces the input size, creating smaller (and more manageable) representations
- Common pooling layers:
 - max pooling: takes the maximum value within each “patch” in the feature map
 - average pooling: takes the average value of each “patch” in the feature map
- It **does not contain** any **learnable weights**

Pooling layers



Pooling – reduces spatial dimension



Pooling layers

Examples



Max pooling



Avg pooling

Pooling layers

Example: max pooling layer, F 2, stride 1

1	3	6	6
20	9	8	4
2	1	4	5
1	12	13	10

20	9	8
20	9	8
12	13	13

Pooling layers

Example: max pooling layer, F 2, stride 2

1	3	6	6
20	9	8	4
2	1	4	5
1	12	13	10

20	8
12	13

Pooling layer

- Parameters
 - Filter size (spatial extent): F
 - Stride: S
- Input: $W_i \times H_i \times D$
- Output: $W_o \times H_o \times D$
- It has no learnable parameters

$$W_o = \frac{W_i - F}{S} + 1$$

$$H_o = \frac{H_i - F}{S} + 1$$

MaxPooling2D layer

MaxPooling2D class

```
tf.keras.layers.MaxPooling2D(  
    pool_size=(2, 2), strides=None, padding="valid", data_format=None, **kwargs  
)
```

Max pooling operation for 2D spatial data.


Downsamples the input representation by taking the maximum value over the window defined by `pool_size` for each dimension along the features axis. The window is shifted by `strides` in each dimension. The resulting output when using "valid" padding option has a shape(number of rows or columns) of: $\text{output_shape} = (\text{input_shape} - \text{pool_size} + 1) / \text{strides}$

The resulting output shape when using the "same" padding option is: $\text{output_shape} = \text{input_shape} / \text{strides}$

AveragePooling2D

Fully connected layers

- They don't preserve spatial information
 - Linear unit followed by a linearity
- Just in regular NN, contain several neurons that are connected to the entire input volume
 - Each neuron "sees" the entire input volume



- If we have an input volume of $27 \times 27 \times 5$, what will be size of this volume if we apply a padding of 2?
- How many parameters (including the bias) does a convolutional layer with 10 filters of size 5×5 have?
- How many parameters (including the bias) does a convolutional layer with 10 filters of size 5×5 have if the input size is $32 \times 32 \times 3$? What if we use a stride of 2?
- Given an input volume that is $63 \times 63 \times 16$ that is convolved with 32 filters that are each 7×7 , using a stride of 2 and no padding. What is the output volume?
- Given a RGB image of size 300×300 , and you use a classical neural network with the first hidden layer of 100 neurons (each one fully connected to the input). How many parameters does this hidden layer have?

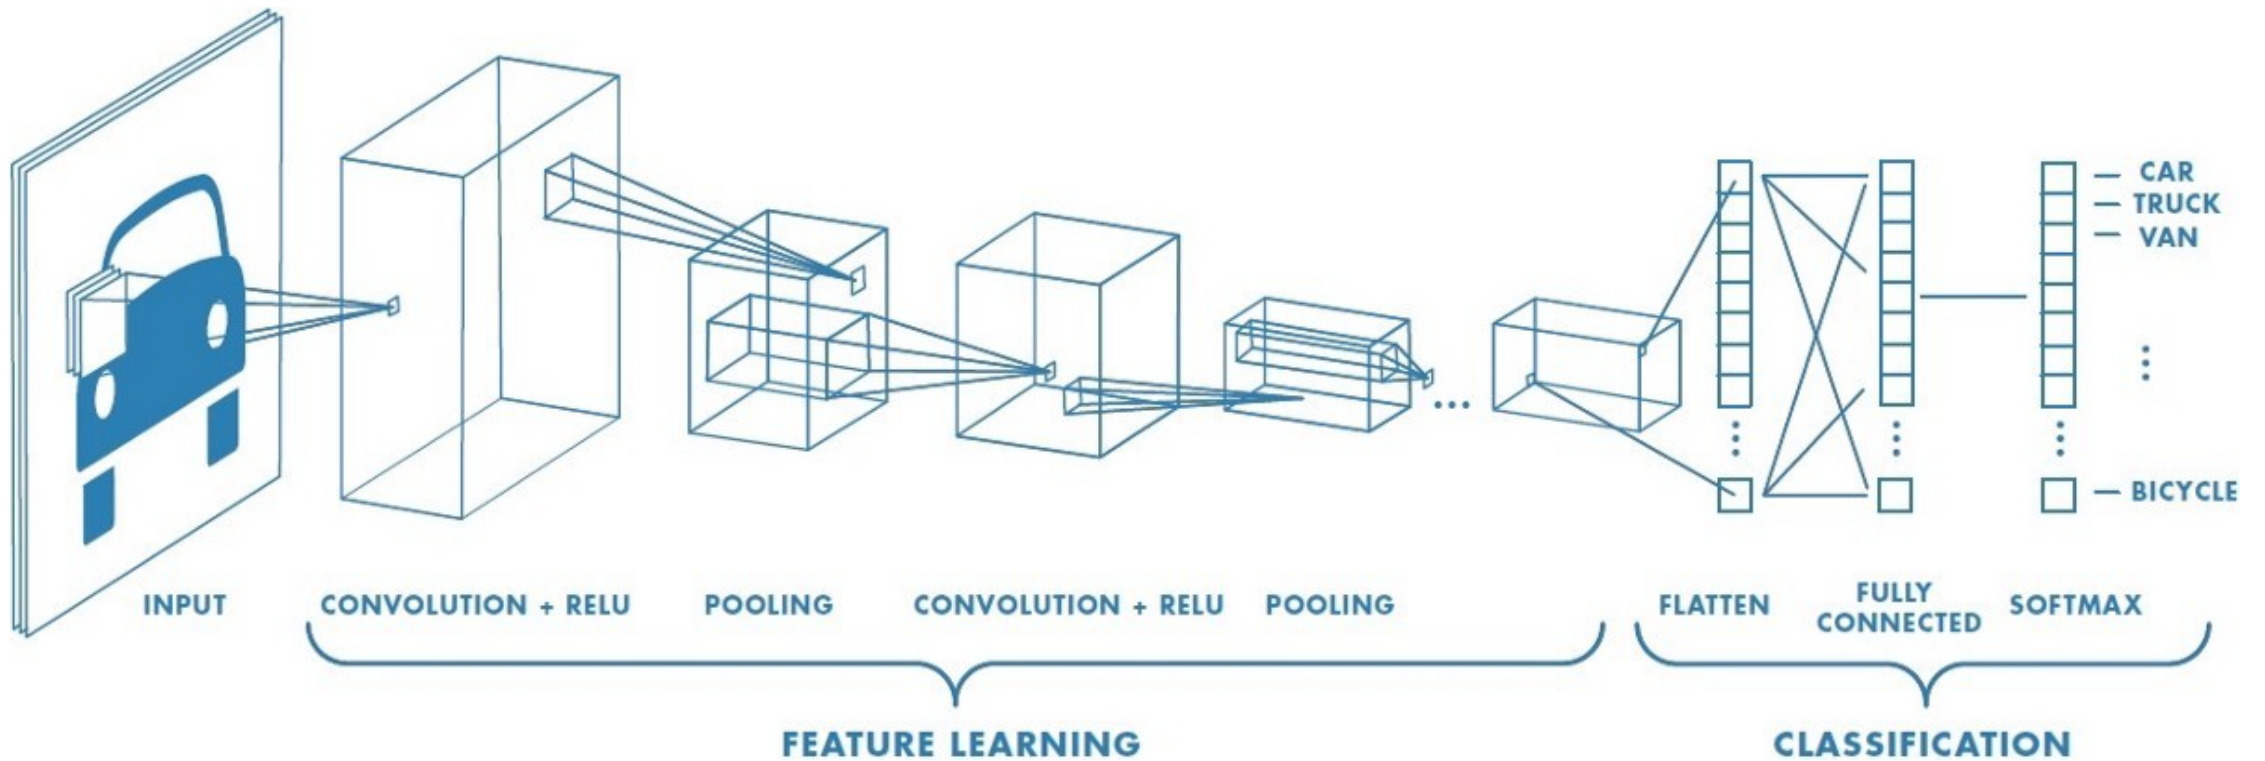
Dense layer

Dense class

```
tf.keras.layers.Dense(  
    units,  
    activation=None,  
    use_bias=True,  
    kernel_initializer="glorot_uniform",  
    bias_initializer="zeros",  
    kernel_regularizer=None,  
    bias_regularizer=None,  
    activity_regularizer=None,  
    kernel_constraint=None,  
    bias_constraint=None,  
    **kwargs  
)
```

Just your regular densely-connected NN layer.


Typical neural network architecture



Typical neural network architecture

- Several CONV, POOL and FC layers stacked together
 - [CONV-ReLU-POOL]*N – [FC]*K – softmax
 - Recent neural networks change this paradigm
- The trend is to reduce the filter sizes and to increase the depth of the networks
- Another trend is to avoid using POOL and FC layers, and use only CONV layers

Alexnet example



AlexNet Network - Structural Details													
Input			Output			Layer	Stride	Pad	Kernel size		in	out	# of Param
227	227	3	55	55	96	conv1	4	0	11	11	3	96	34944
55	55	96	27	27	96	maxpool1	2	0	3	3	96	96	0
27	27	96	27	27	256	conv2	1	2	5	5	96	256	614656
27	27	256	13	13	256	maxpool2	2	0	3	3	256	256	0
13	13	256	13	13	384	conv3	1	1	3	3	256	384	885120
13	13	384	13	13	384	conv4	1	1	3	3	384	384	1327488
13	13	384	13	13	256	conv5	1	1	3	3	384	256	884992
13	13	256	6	6	256	maxpool5	2	0	3	3	256	256	0
						fc6			1	1	9216	4096	37752832
						fc7			1	1	4096	4096	16781312
						fc8			1	1	4096	1000	4097000
Total												62,378,344	

“AlexNet input starts with 227 by 227 by 3 images. And if you read the paper, the paper refers to 224 by 224 by 3 images. But if you look at the numbers, I think that the numbers make sense only of actually 227 by 227.”

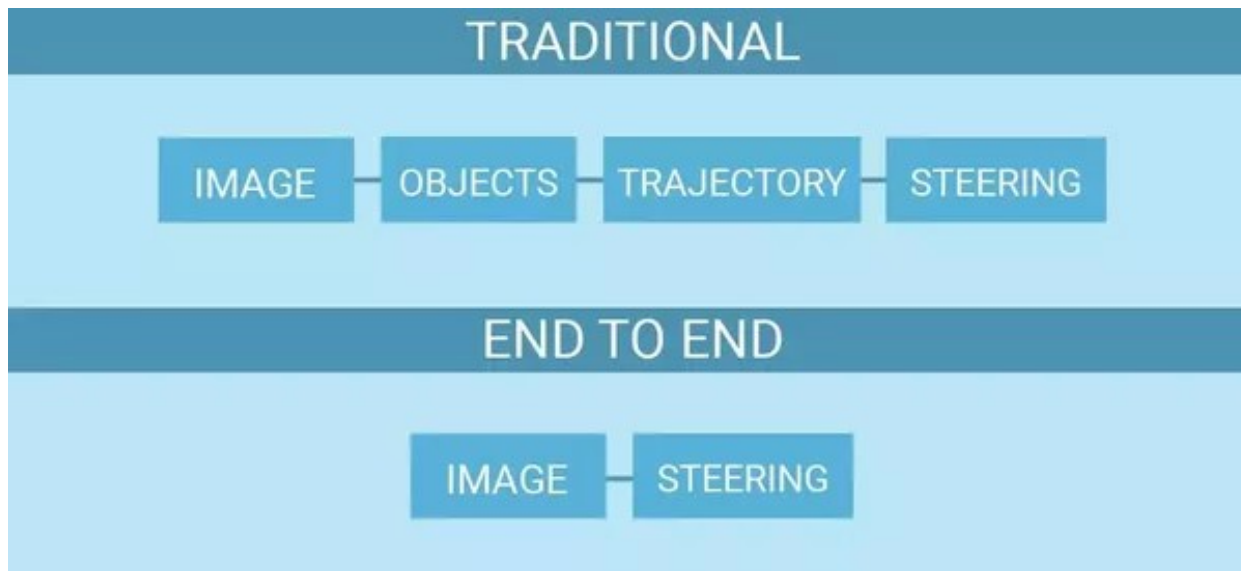
Playground

<https://cs.stanford.edu/people/karpathy/convnetjs/demo/mnist.html>

Learning from multiple tasks

- Transfer learning
 - Task A and task B have the same input
 - More data for task A than for task B
 - Low level features from task A could be helpful for learning task B
- Multi task learning
 - Training on a set of tasks that could benefit from having shared low level features
 - Amount of data for each task is quite similar. Can train a big enough network to do well on all the tasks.

End to end deep learning



Advantages:

- Lets the data speak
- No handcrafted features


Disadvantages

- Needs huge amount of data
- Excludes potentially useful hand-designed components


Data pre-processing

Importance of feature scaling


Gradient descent
without scaling




Gradient descent
after scaling variables



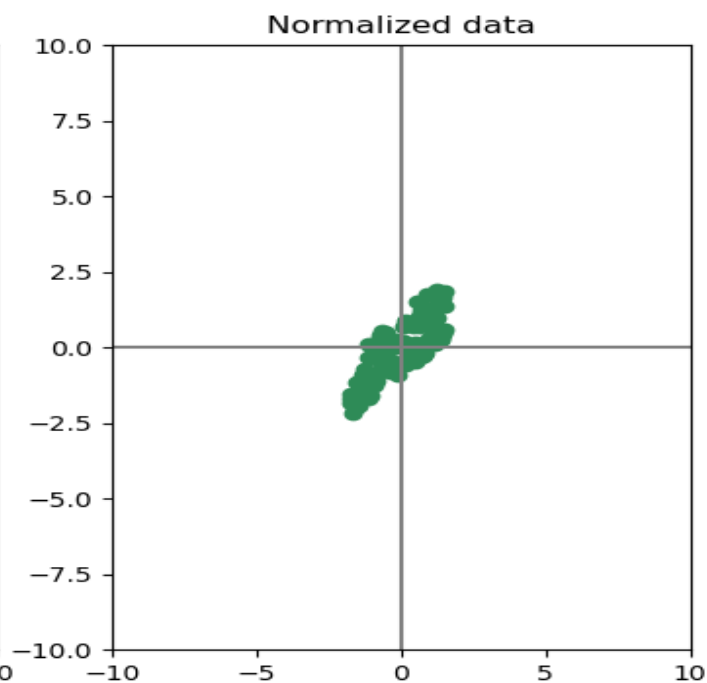
Mean subtraction, scaling



X



$X -= \text{np.mean}(X, \text{axis}=0)$



$X /= \text{np.std}(X, \text{axis}=0)$

Pre-processing for Images

- Zero center: subtract the mean across every individual feature in the data
 - Mean image
 - Mean across each channel
- Optional: normalize the data such that dimensions are approximately the same scale

https://github.com/keras-team/keras-applications/blob/master/keras_applications/imagenet_utils.py