

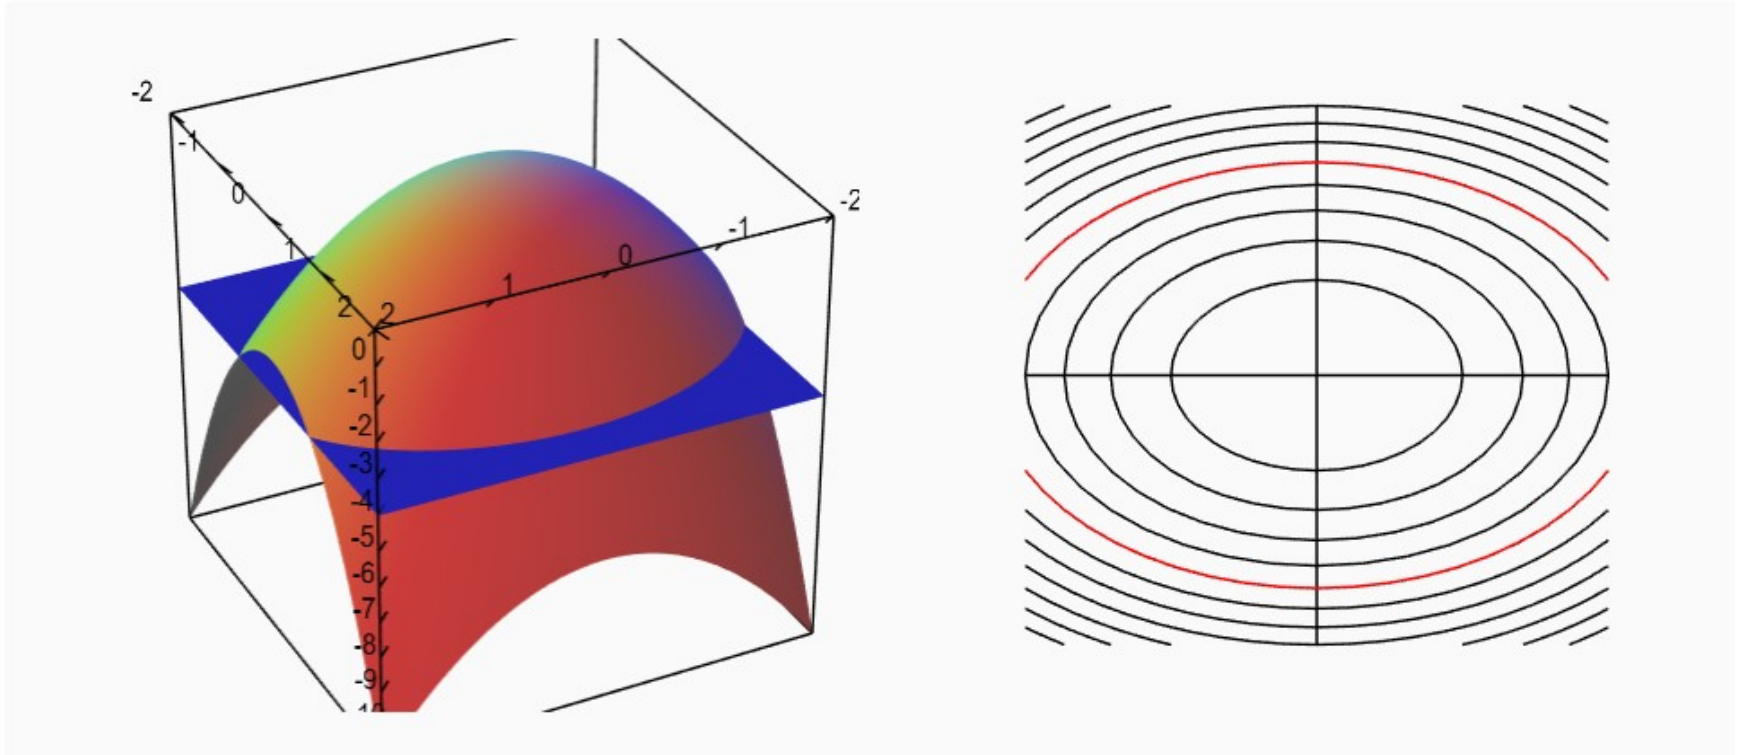
Computer Vision and Deep Learning

Lecture 6

Today's agenda

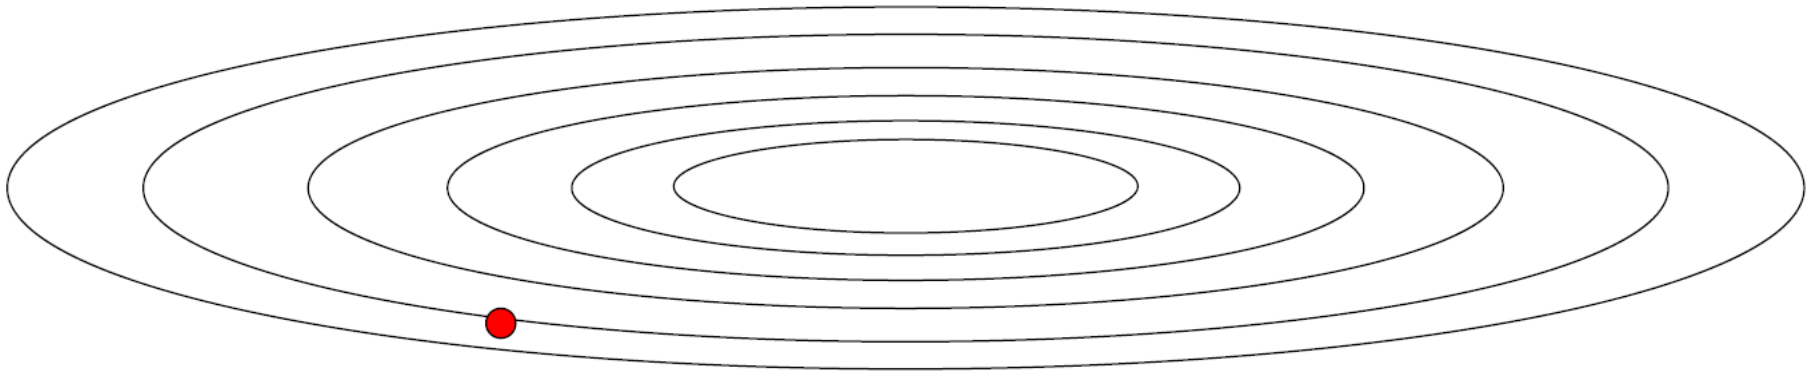
- Optimization algorithms
- Convolutional neural networks:
 - History
 - Case studies
- How to read a research paper?

Level sets of a surface



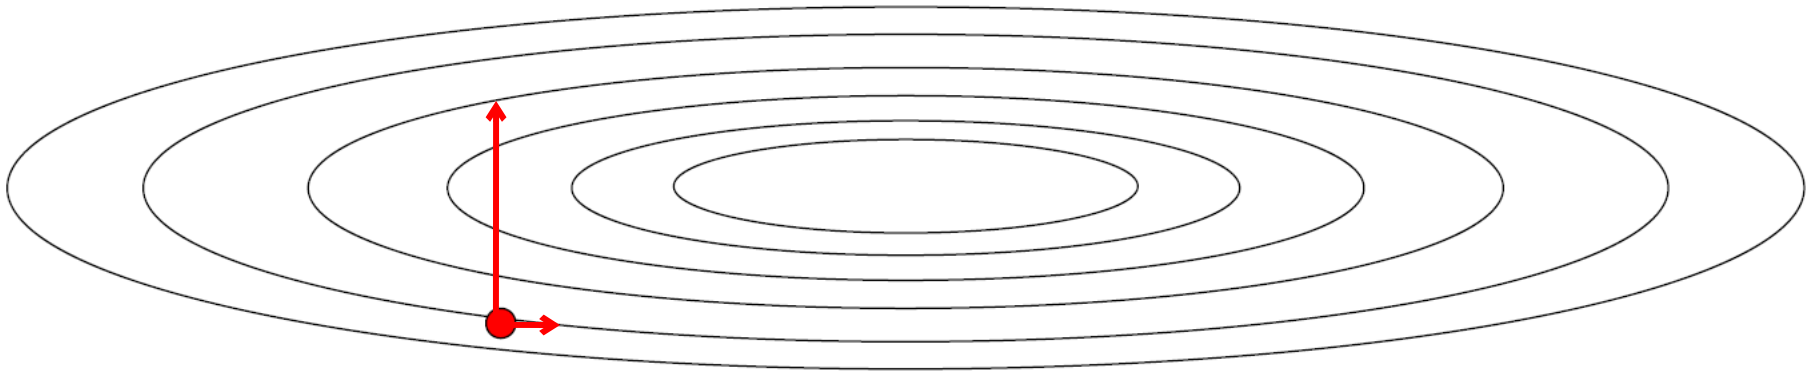
Problems with gradient descent

What if loss changes abruptly on one direction and slower in the other direction?



Problems with gradient descent

What if loss changes abruptly on one direction and slower in the other direction?

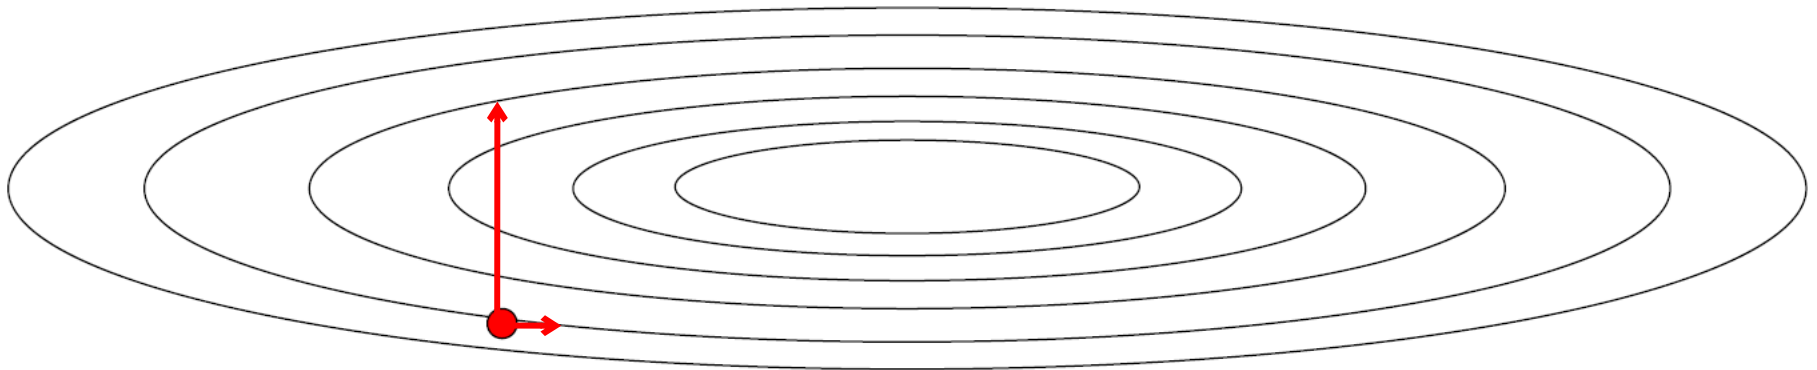


Problems with gradient descent

What if loss changes abruptly on one direction and slower in the other direction?

- Small gradient horizontally
- Large gradient vertically

Slow progress along the horizontal direction, jitter along the vertical direction (the steeper one)



Gradient descent with momentum

Compute an exponentially weighted average of the gradients and use this average to update the parameters of the network

Gradient descent with momentum

Update rule

Gradient descent update:

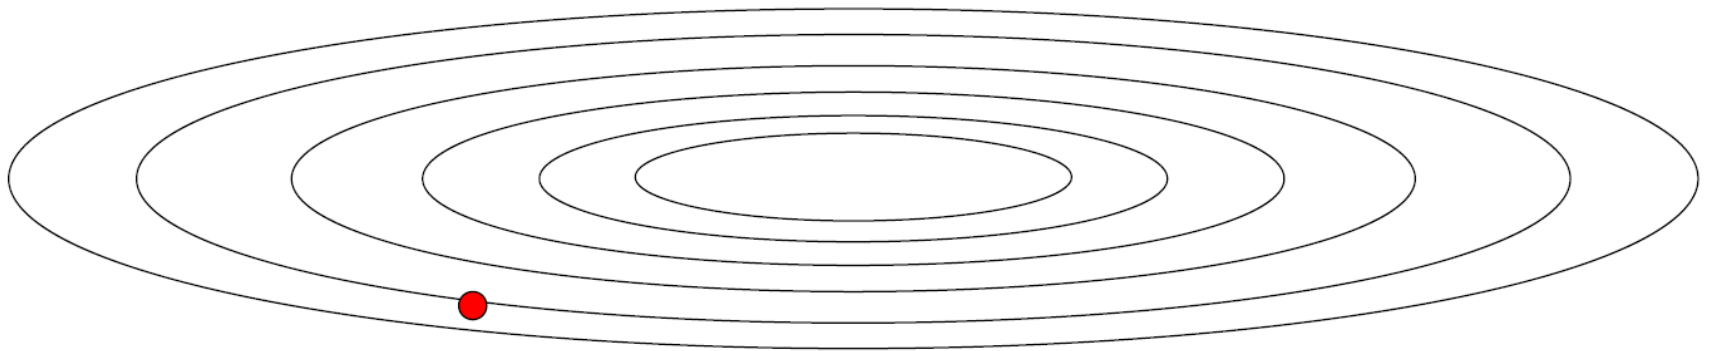


Momentum update:

– hyper-parameter; common value 0.9

Gradient descent with momentum

Take more straightforward part (damp oscillations)



Slower learning vertically, faster learning horizontally

Gradient descent with momentum

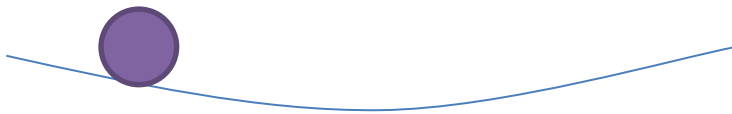
- Intuition
 - Ball rolling down the loss function with friction
 - Gradient: the force the ball is feeling ($F = ma$)

$$\mathbf{v}_{dW} = \beta \cdot \mathbf{v}_{dW} + (1 - \beta) \cdot dW$$

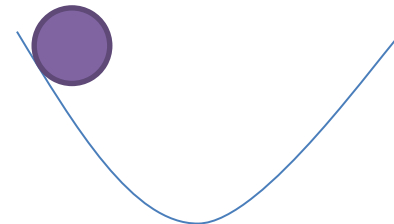
– friction
Hyper-parameter

velocity

acceleration



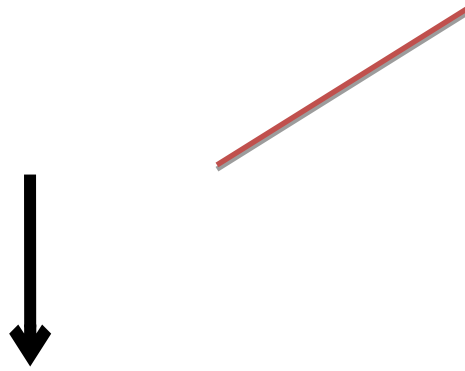
Shallow, consistent direction: build up the velocity across the dimension



Steep direction: attenuate velocity (quickly changing sign); oscillate to the “middle”

Gradient descent with momentum

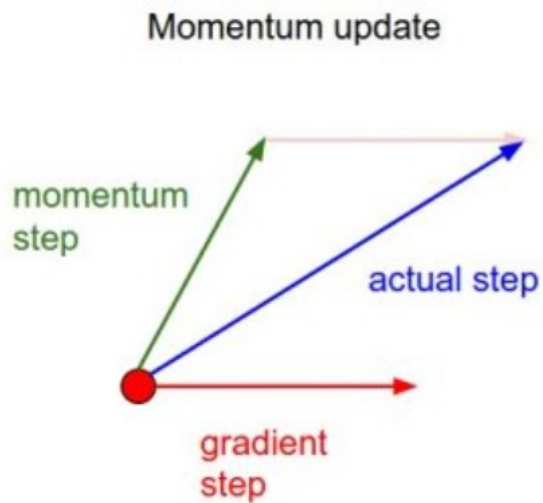
Update rule - variant



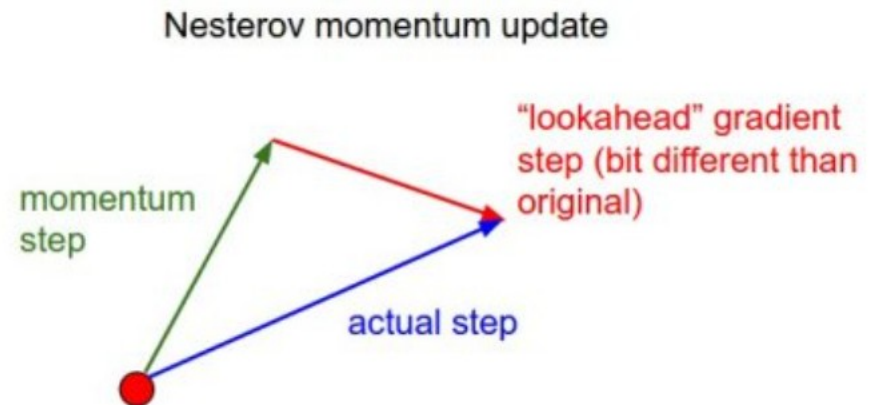
– hyper-parameter; common value 0.9

Nesterov momentum

$$v_{dW} = \beta \cdot v_{dW} + dW$$



$$v_{t+1} = \mu v_t - \eta \nabla l(\theta)$$
$$\theta_{t+1} = \theta_t + v_{t+1}$$



$$v_{t+1} = \mu v_t - \eta \nabla l(\theta + \mu v_t)$$
$$\theta_{t+1} = \theta_t + v_{t+1}$$

Image source: <https://cs231n.github.io/neural-networks-3/#sgd>

<https://dominikschmidt.xyz/nesterov-momentum/>

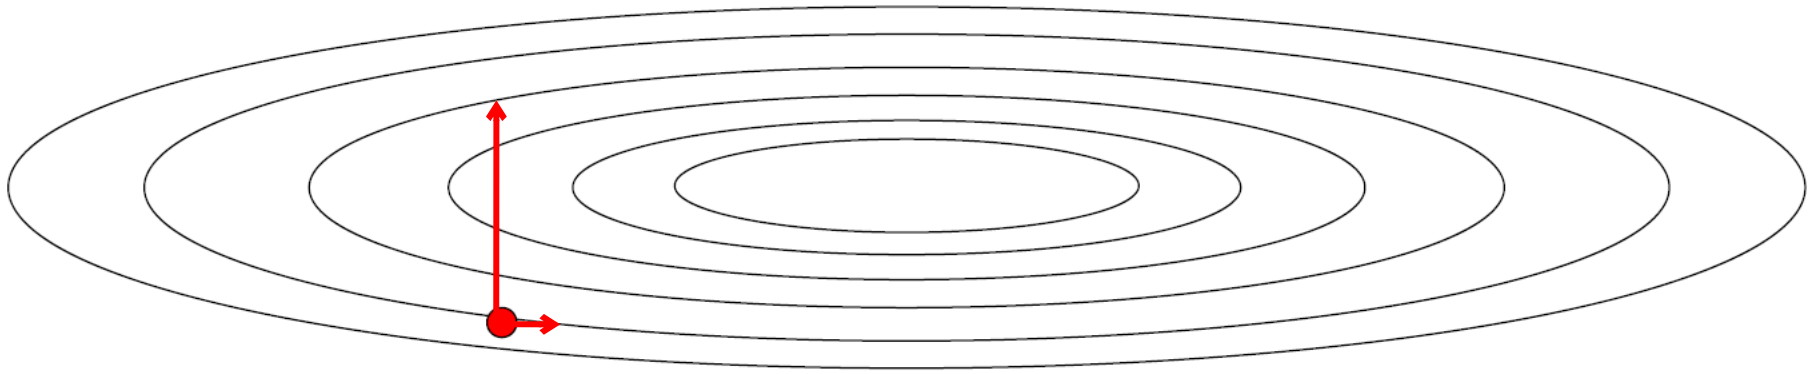
AdaGrad

cache – same size as W

Element wise scaling of the gradient based on the “historical” sum of squares in each dimension

Per parameter adaptive learning rate

AdaGrad



Equalizing effect

- Larger learning rate on “shallow” directions than on steeper directions

RMSProp

Adagrad



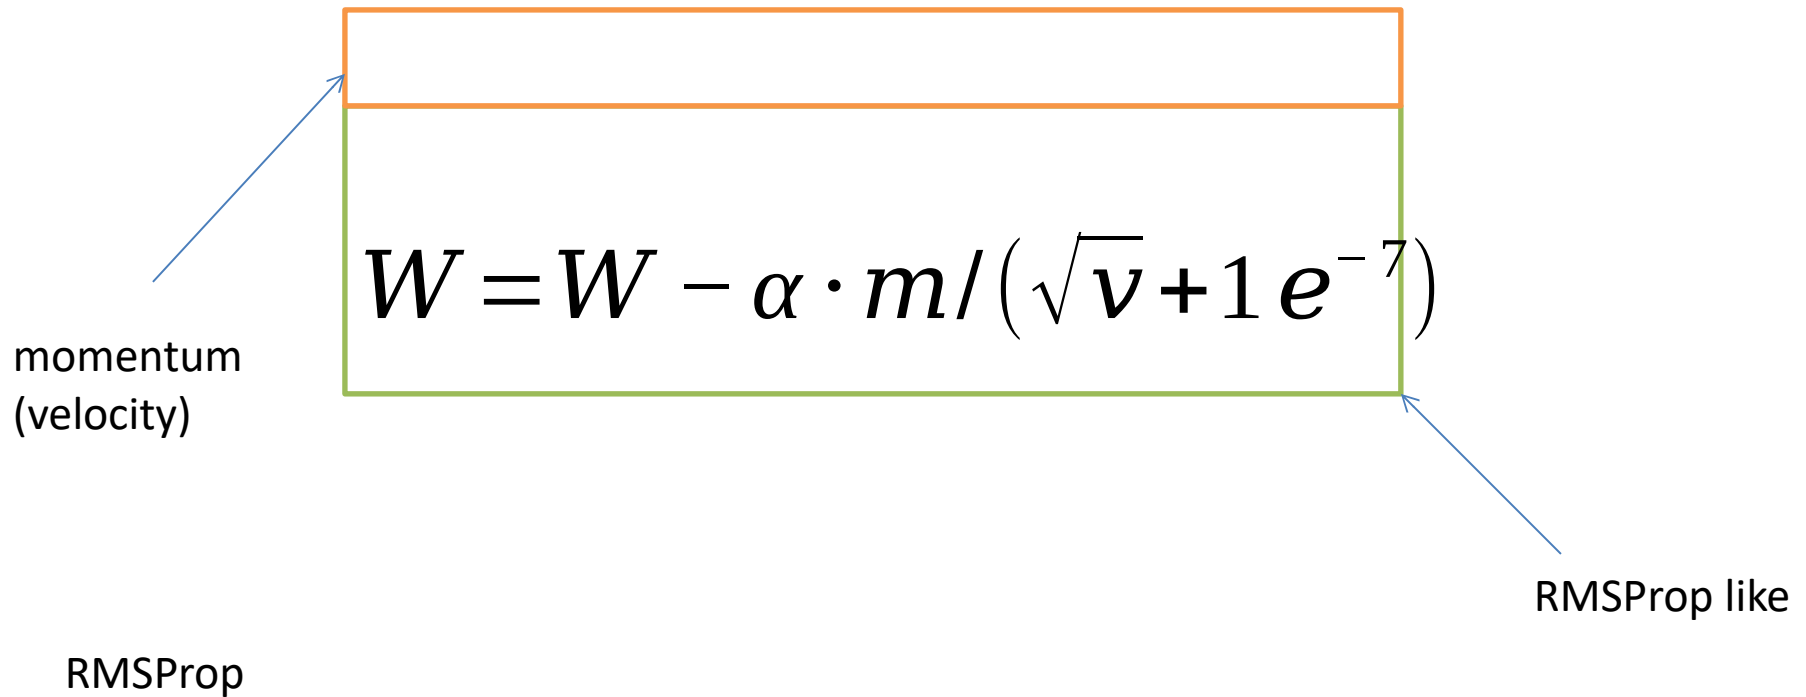
RMSProp

Introduced by Geoffrey Hinton on a lecture on Coursera:

- [14] Duchi J, Hazan E and Singer Y 2011 *The Journal of Machine Learning Research* **12** 2121
- [15] Tieleman T and Hinton E 2012 Lecture 6.5 - rmsprop, COURSERA: Neural networks for machine learning.

Adam

- Combine RMSProp with momentum



The diagram shows the Adam update equation for weights W . The equation is enclosed in a green rectangular box, which is labeled "RMSProp" at the bottom left. Above this green box is an orange rectangular box. A blue arrow points from the text "momentum (velocity)" to the top-left corner of the orange box. Another blue arrow points from the text "RMSProp like" to the bottom-right corner of the green box.

$$W = W - \alpha \cdot m / (\sqrt{v} + 1 e^{-7})$$

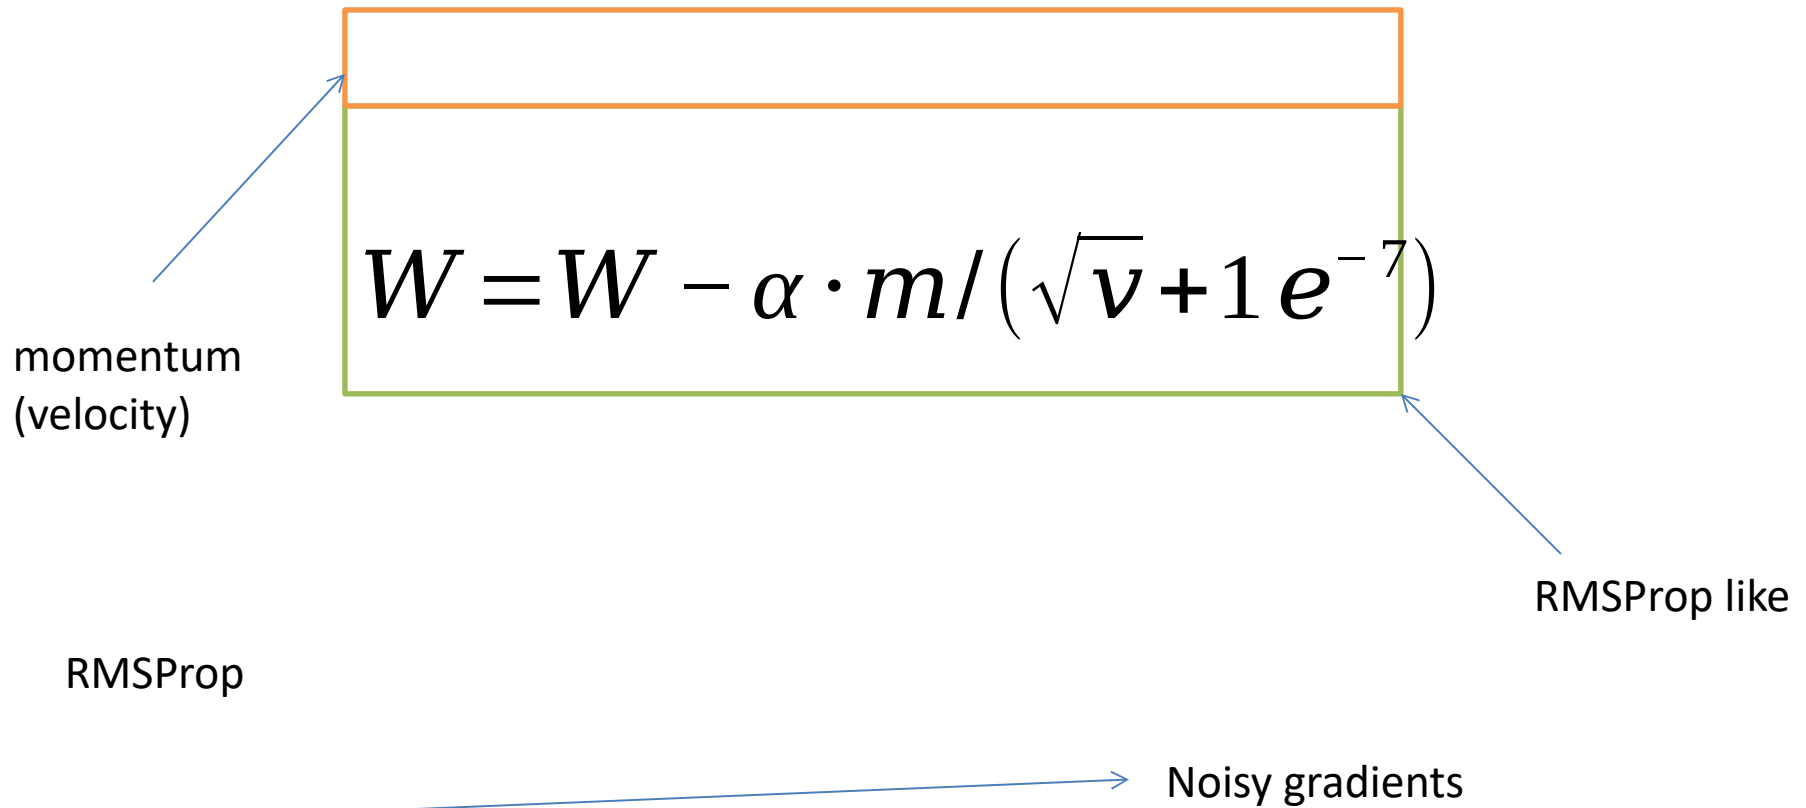
momentum
(velocity)

RMSProp

RMSProp like


Adam

- Combine RMSProp with momentum

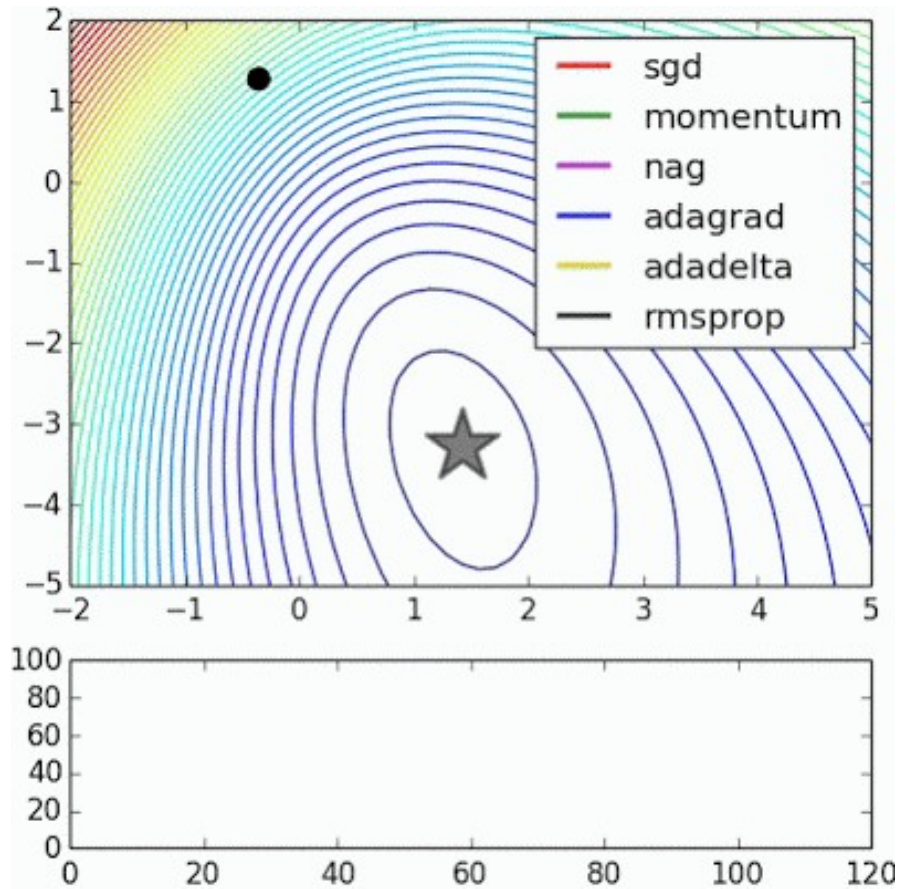


Adam

- Combine RMSProp with momentum for t in range(0, iterations):

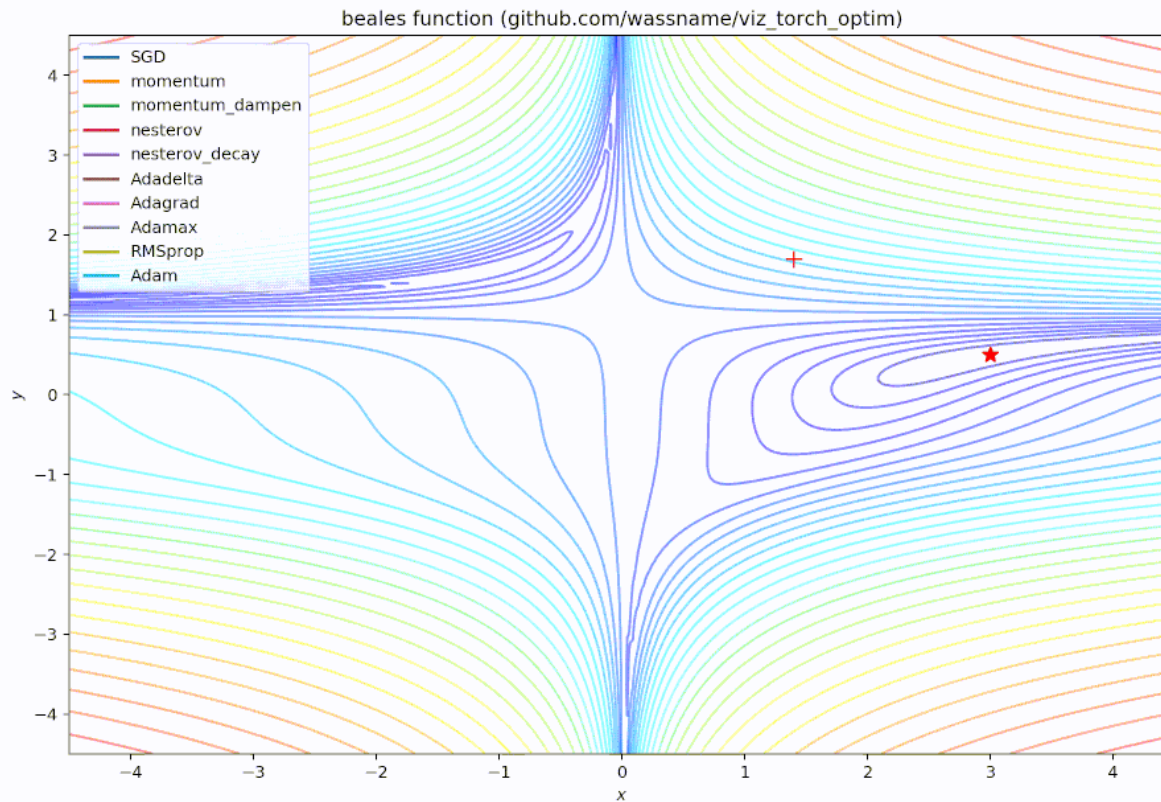
Bias correction 

$$\begin{cases} \mathbf{v} = \mathbf{v} / (1 - \beta_2^t) \\ \mathbf{W} = -\alpha \cdot \mathbf{m} / (\sqrt{\mathbf{v} + 1} e^{-7}) \end{cases}$$



[http://
www.denizyuret.com/2015/03/alec-radfords-animations-for.html](http://www.denizyuret.com/2015/03/alec-radfords-animations-for.html)

Optimization algorithms



https://awesomeopensource.com/project/3springs/viz_torch_optim

Optimizers in keras

```
tf.keras.optimizers.SGD(  
    learning_rate=0.01, momentum=0.0, nesterov=False, name="SGD", **kwargs  
)
```

```
tf.keras.optimizers.Adagrad(  
    learning_rate=0.001,  
    initial_accumulator_value=0.1,  
    epsilon=1e-07,  
    name="Adagrad",  
    **kwargs  
)
```

```
tf.keras.optimizers.Adam(  
    learning_rate=0.001,  
    beta_1=0.9,  
    beta_2=0.999,  
    epsilon=1e-07,  
    amsgrad=False,  
    name="Adam",  
    **kwargs  
)
```

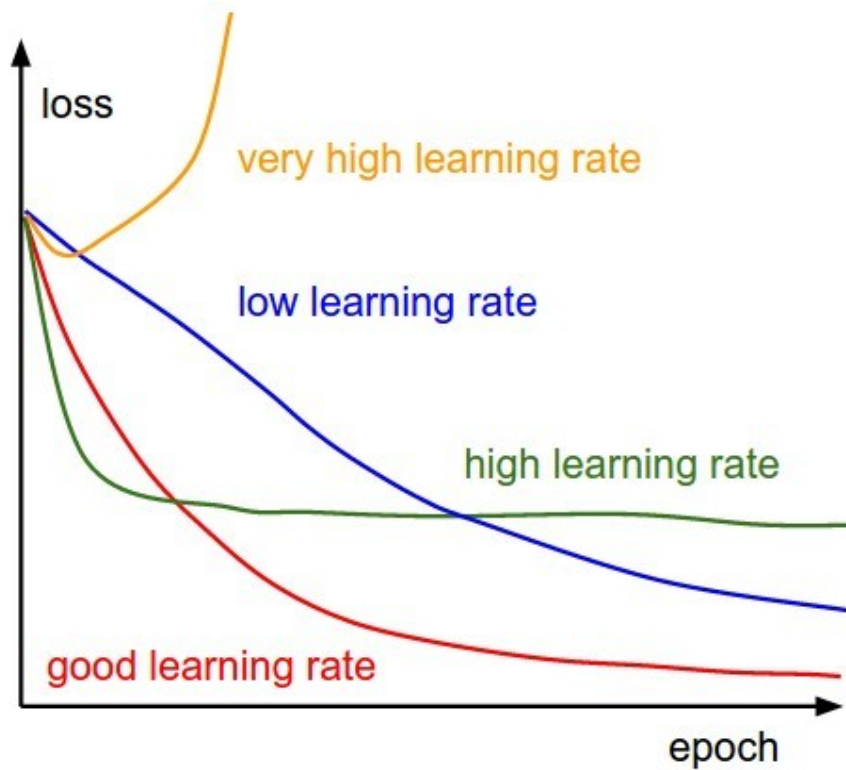
Optional additional reading

- [https://
runder.io/optimizing-gradient-descent/index.html
#otherrecentoptimizers](https://runder.io/optimizing-gradient-descent/index.html#otherrecentoptimizers)
- https://www.youtube.com/watch?v=k8fTYJPd3_I
- https://www.youtube.com/watch?v=_e-LFe_igno
- [https://
www.coursera.org/lecture/deep-neural-network
/the-problem-of-local-optima-RFANA](https://www.coursera.org/lecture/deep-neural-network/the-problem-of-local-optima-RFANA)

Training a neural network

Learning rate scheduling

Learning rate



Learning rate decay

- Learning rate decay over time!
 - Use a larger learning rate at the beginning and progressively reduce the learning rate
 - t epoch, k – decay rate
- Step decay:
 - reduce learning rate by half after some epochs
- Exponential decay

$$\alpha = \alpha_0 \cdot e^{-kt}$$

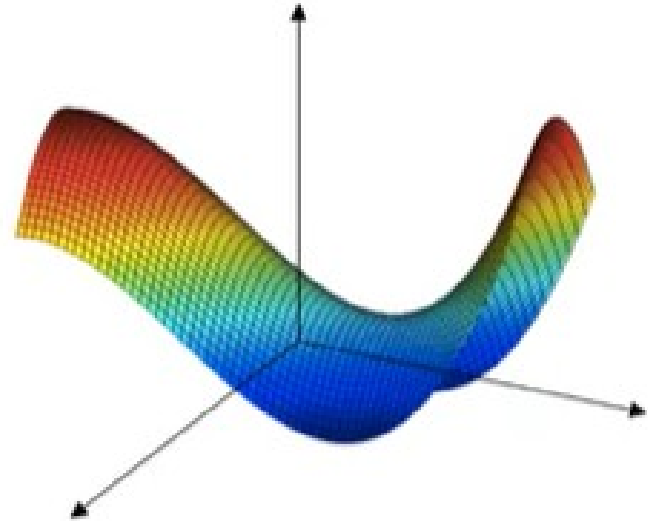
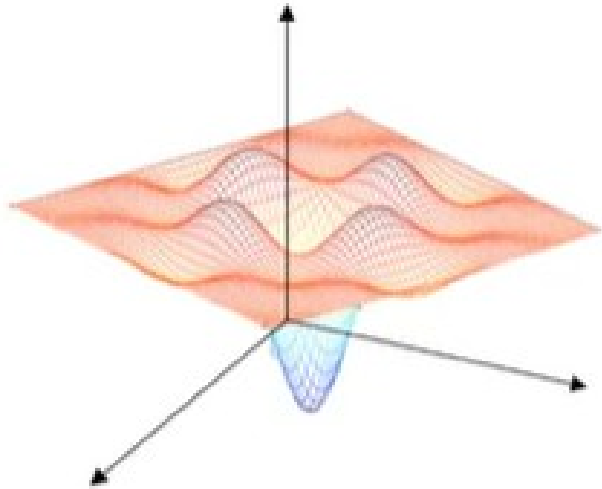
- 1/t decay

$$\alpha = \alpha_0 / (1 + kt)$$

keras optimizers schedules

```
lr_schedule = keras.optimizers.schedules.ExponentialDecay(  
    initial_learning_rate=1e-2,  
    decay_steps=10000,  
    decay_rate=0.9)  
optimizer = keras.optimizers.SGD(learning_rate=lr_schedule)
```

Problem of local optima



Gradient is zero \rightarrow in each direction it can either be a convex like a concave like function.

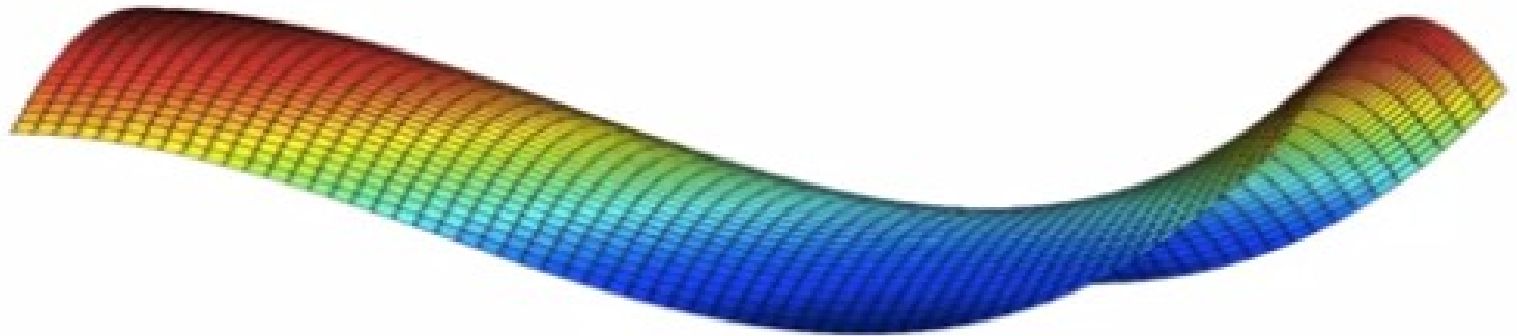
Saddle points.

Unlikely to get stuck in local optima

<https://>

www.coursera.org/lecture/deep-neural-network/the-problem-of-l

Problem of plateaus

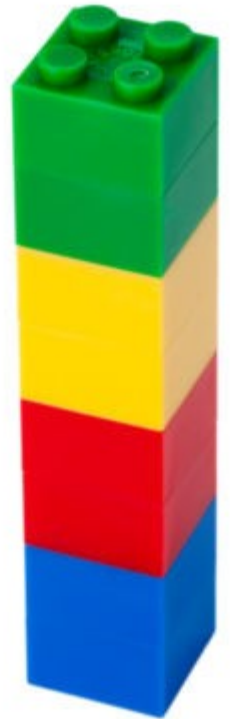


Plateaus can make the learning process too slow

Previously

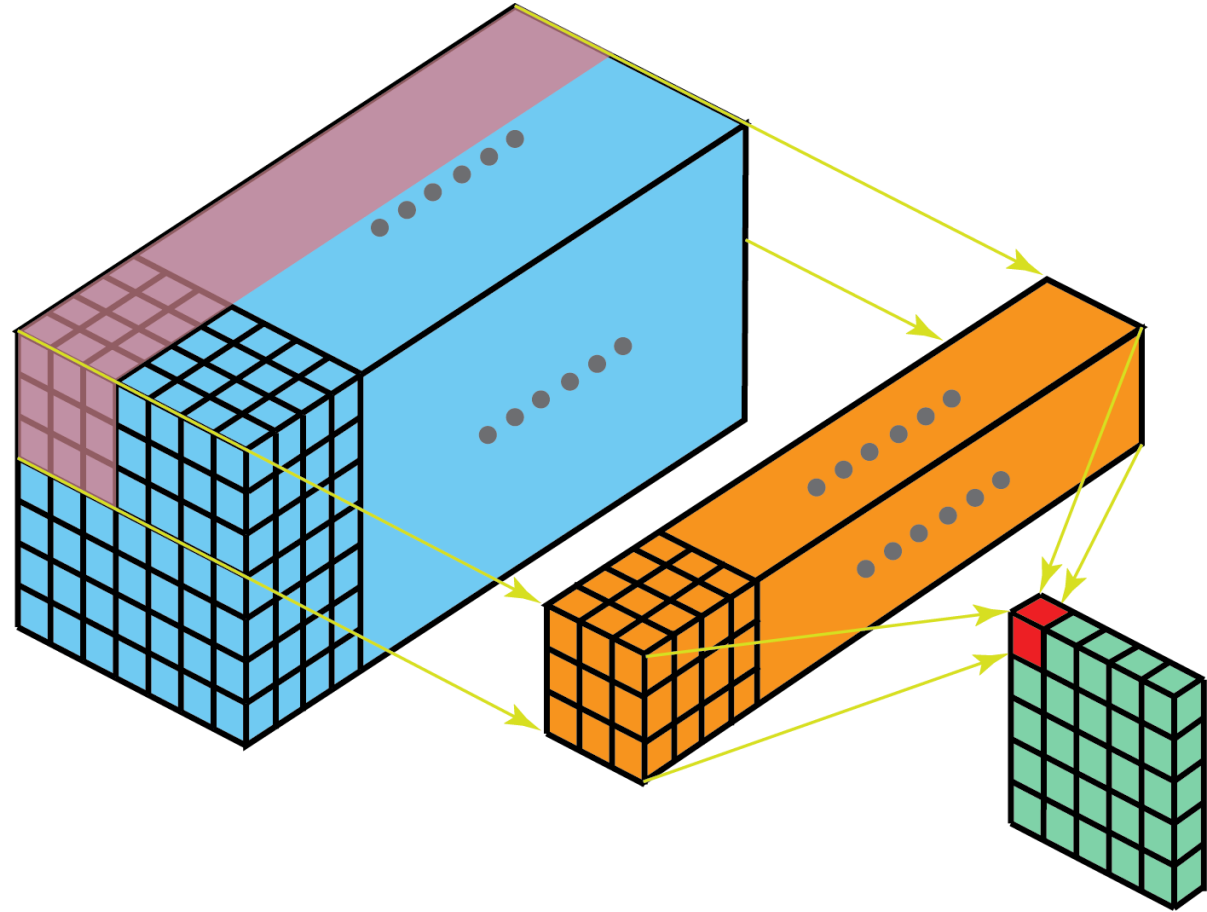
Convolutional neural networks

- Convolutional neural networks
 - Convolutional layers
 - Pooling layers
 - Fully connected layers



Convolution layers

Shared parameters
Sparsity of connections



Pooling layers

2	2	7	3
9	4	6	1
8	5	2	4
3	1	2	6

Max Pool
→
Filter - (2 x 2)
Stride - (2, 2)

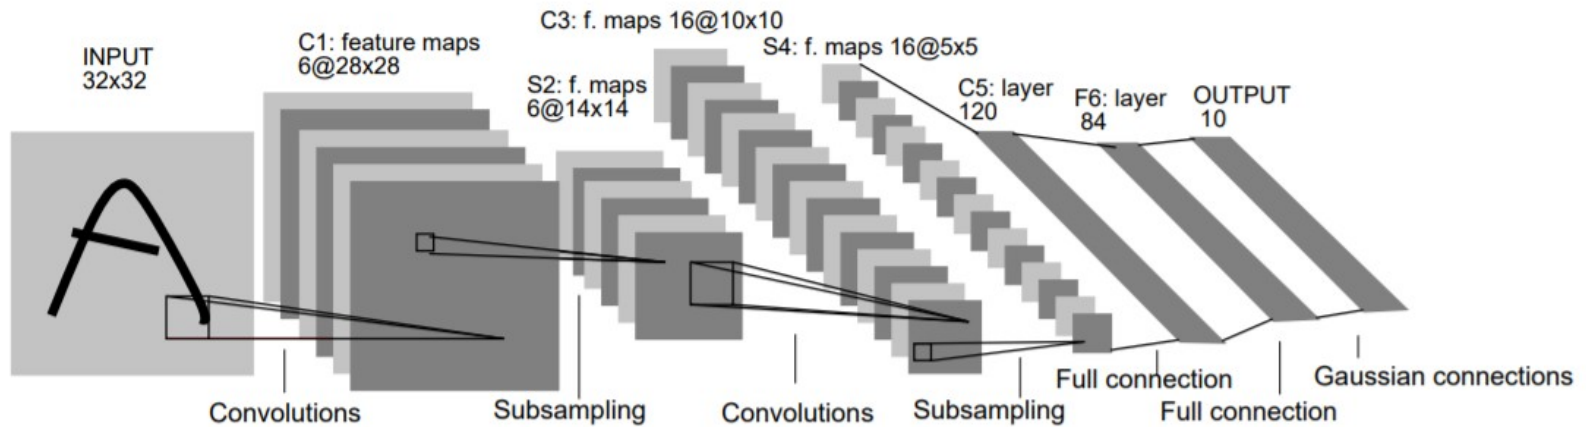
9	7
8	6

Previously

Convolutional neural networks

- Training a neural network
 - Good initialization
 - Batch normalization
 - Regularization
 - Batch training
 - Stochastic gradient descent
 - Mini-batch gradient descent
 - Batched gradient descent
 - Optimizers

LeNet 5

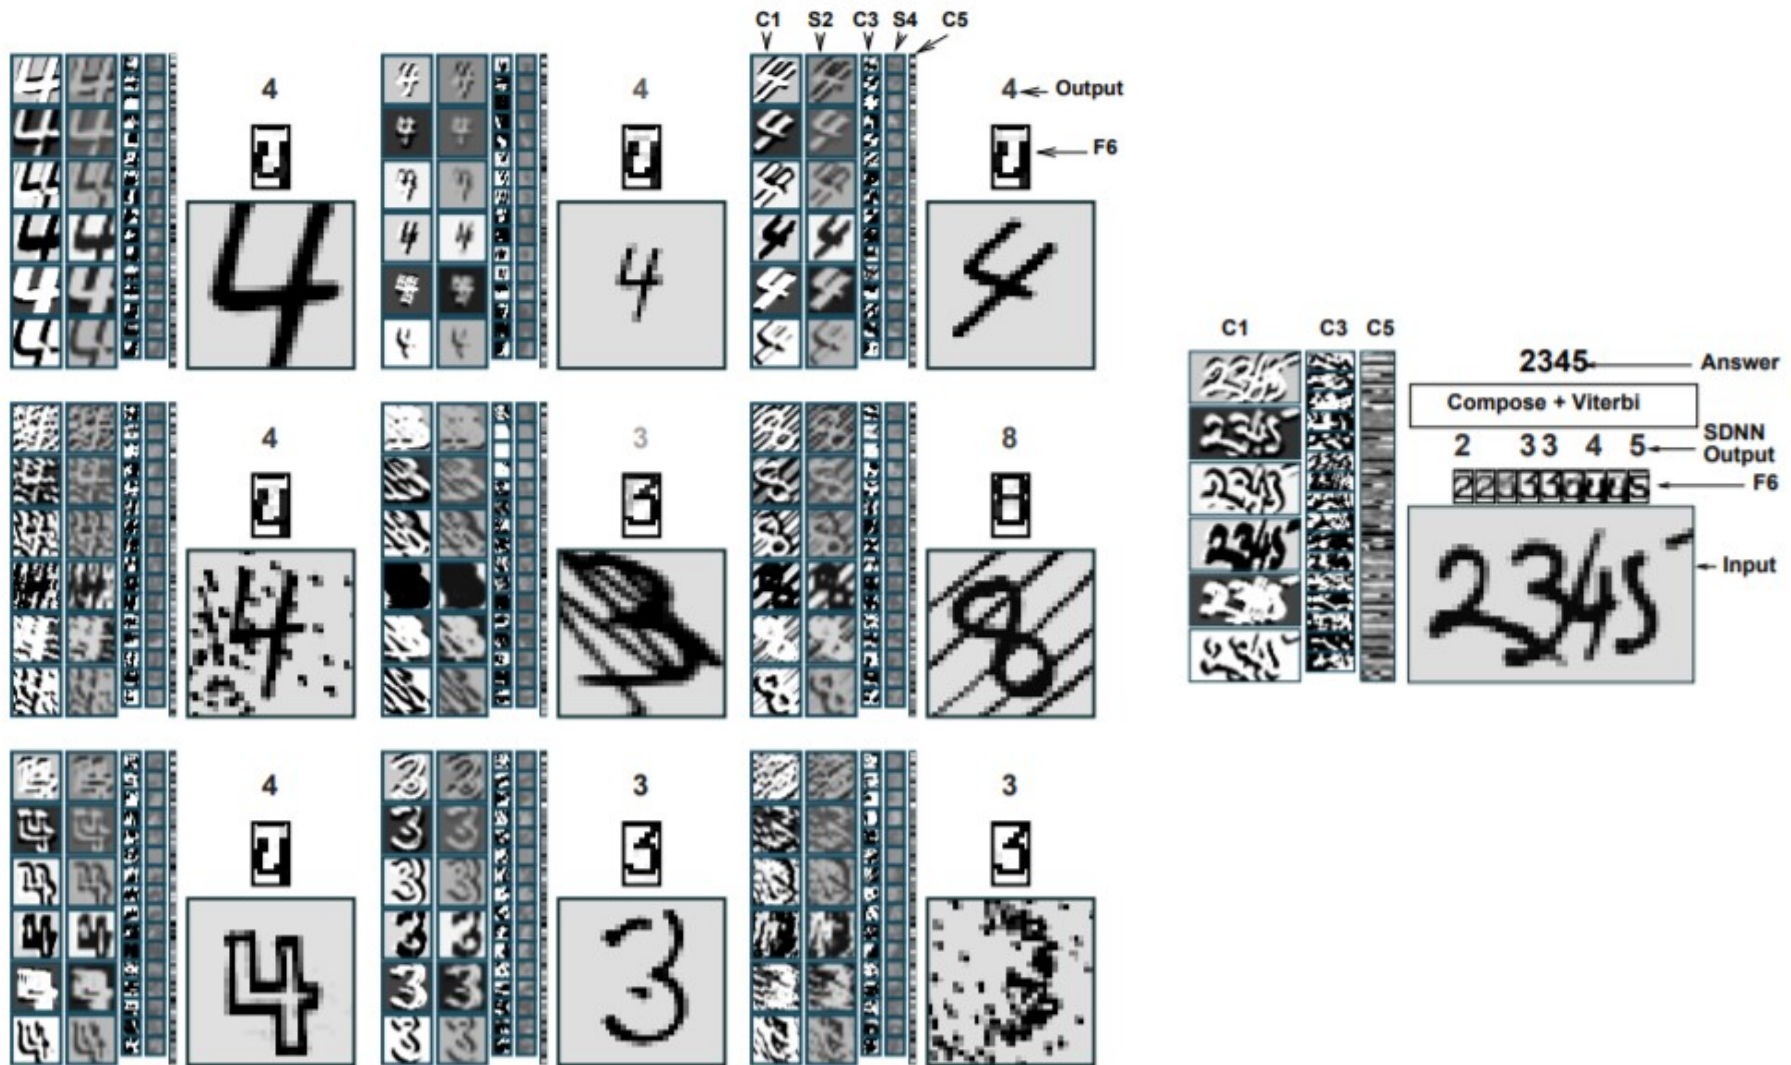


Conv filters have size 5x5 and are applied at stride 1

Subsampling (Pooling) layers have size 2x2 and are applied at stride 2

[CONV-POOL-CONV-POOL-FC-FC]

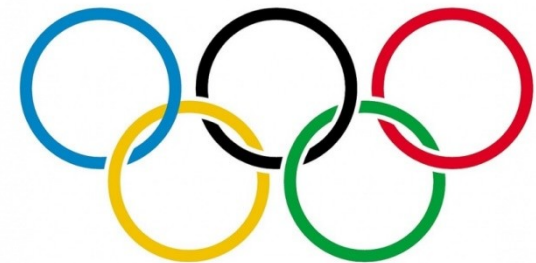
LeNet 5



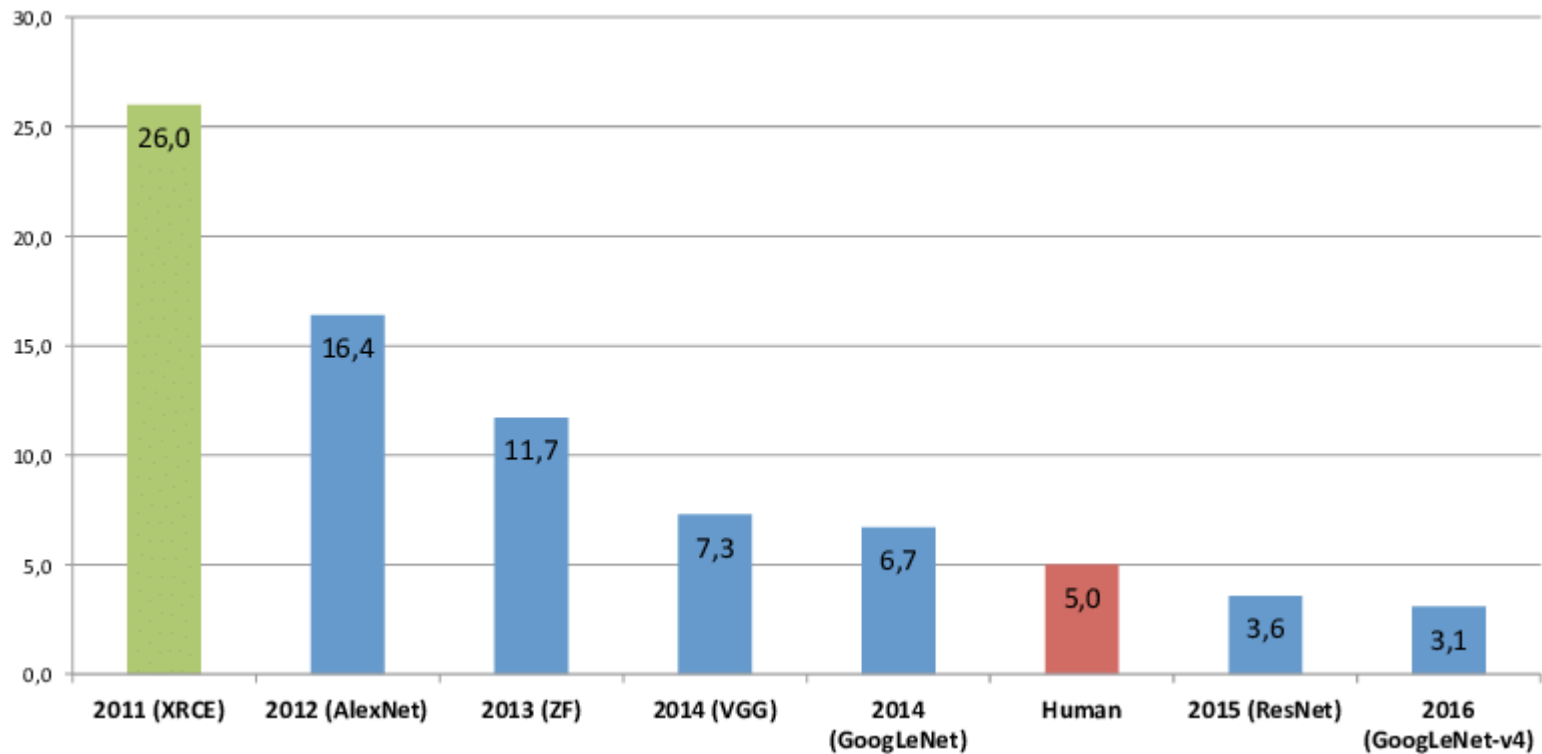
2012



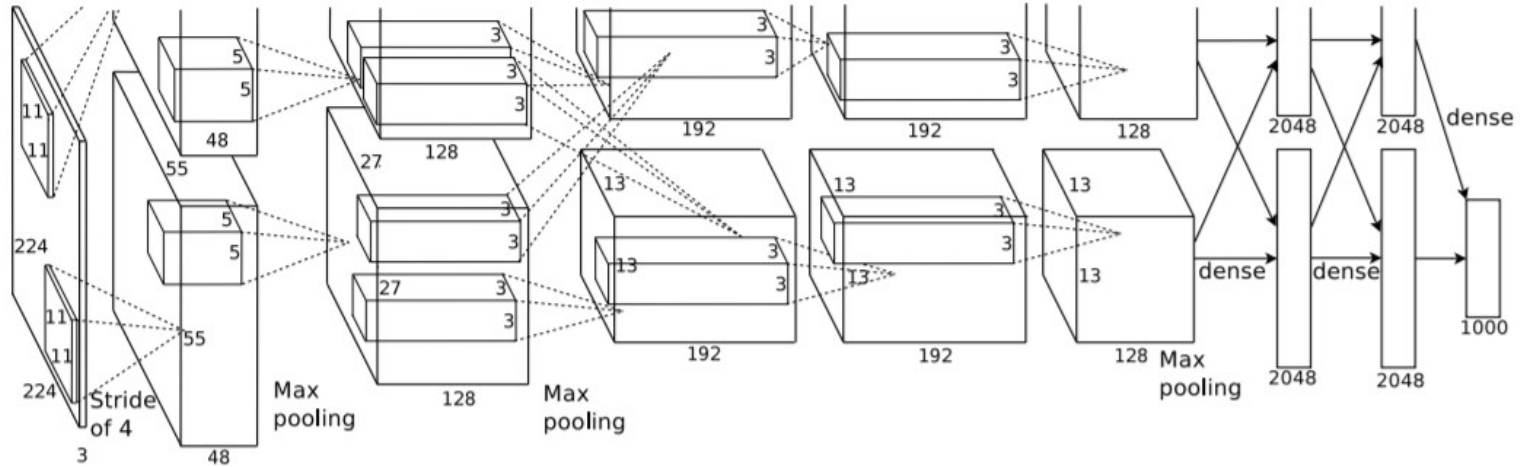
14,197,122 images
1000 categories



ImageNet Classification Error (Top 5)



Alexnet, 2012



A single GTX 580 GPU has only 3GB of memory, which limits the maximum size of the networks that can be trained on it.

Alexnet, 2012

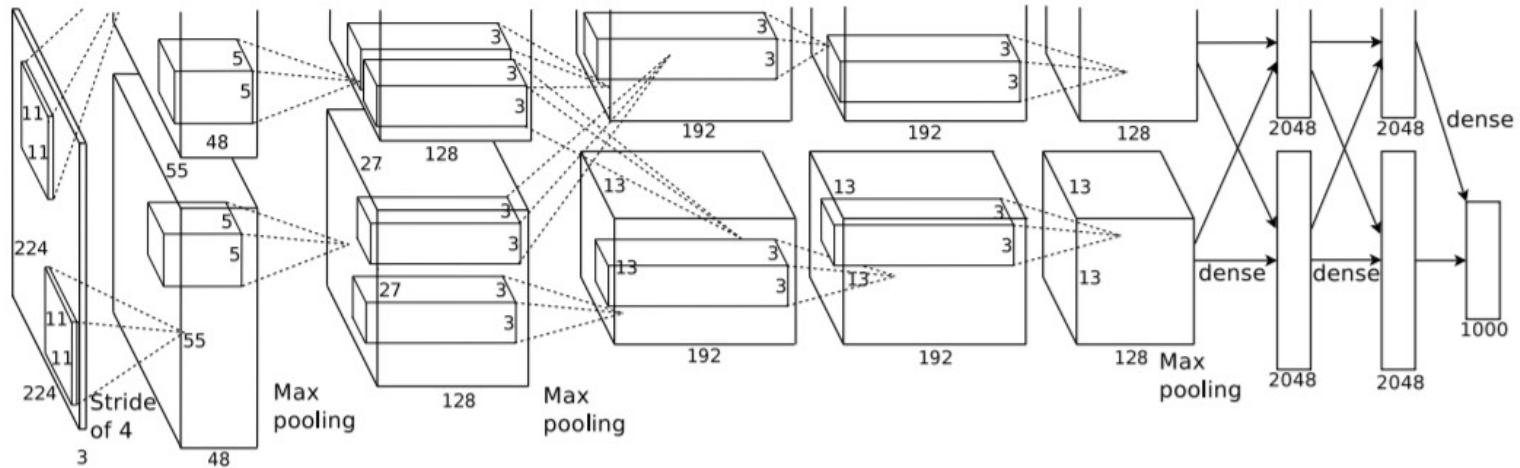


Image size: (227, 227, 3)
First layer: 96 filter of size 11x11

$$W_o = \frac{W_I - F + 2P}{S} + 1$$

What is the output volume size and the number of parameter in this layer?

Alexnet, 2012

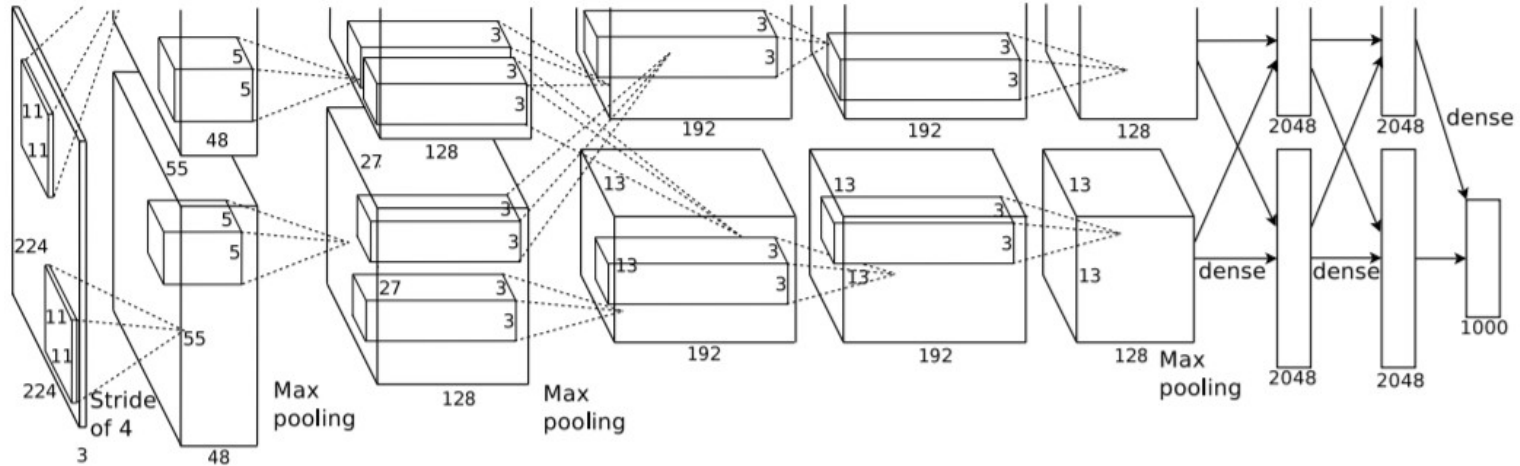


Image size: (227, 227, 3)

First layer: 96 filter of size 11x11

$$W_o = \frac{W_I - F + 2P}{S} + 1$$

Output volume: 55x55x96

Parameters: $(11*11*3)*96 \sim 35K$

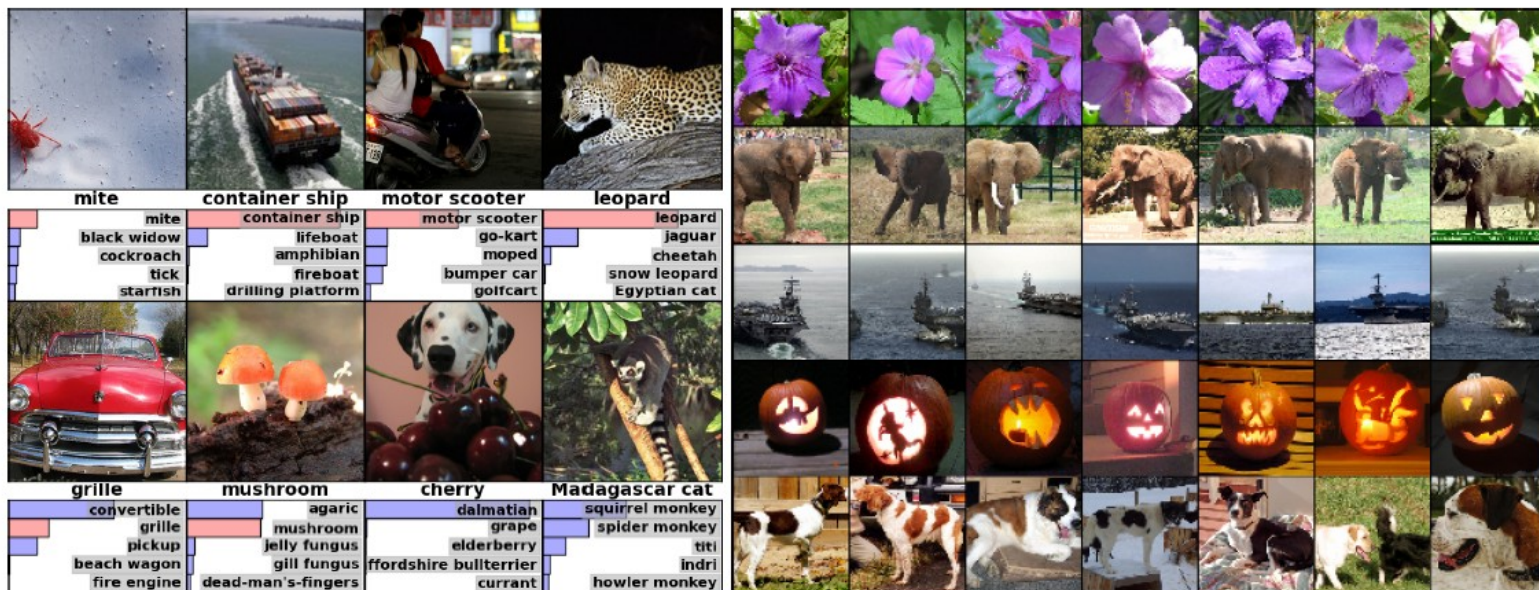
Alexnet

AlexNet Network - Structural Details														
Input			Output			Layer	Stride	Pad	Kernel size		in	out	# of Param	
227	227	3	55	55	96	conv1	4	0	11	11	3	96	34944	
55	55	96	27	27	96	maxpool1	2	0	3	3	96	96	0	
27	27	96	27	27	256	conv2	1	2	5	5	96	256	614656	
27	27	256	13	13	256	maxpool2	2	0	3	3	256	256	0	
13	13	256	13	13	384	conv3	1	1	3	3	256	384	885120	
13	13	384	13	13	384	conv4	1	1	3	3	384	384	1327488	
13	13	384	13	13	256	conv5	1	1	3	3	384	256	884992	
13	13	256	6	6	256	maxpool5	2	0	3	3	256	256	0	
						fc6			1	1	9216	4096	37752832	
						fc7			1	1	4096	4096	16781312	
						fc8			1	1	4096	1000	4097000	
Total													62,378,344	

Alexnet

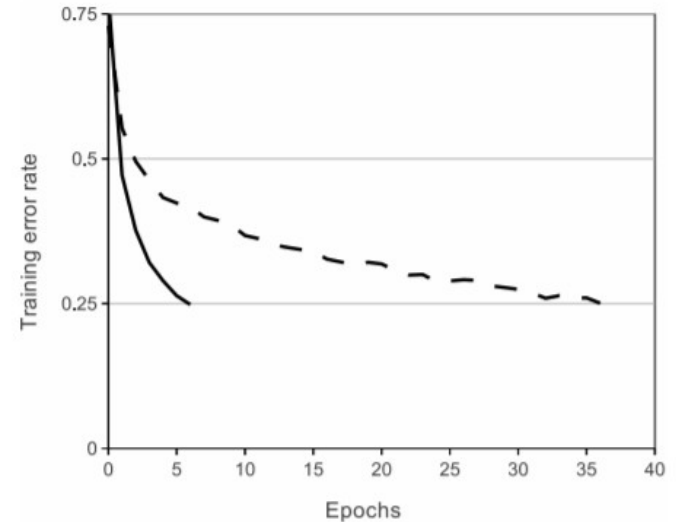
Model	Top-1 (val)	Top-5 (val)	Top-5 (test)
<i>SIFT + FVs [7]</i>	—	—	26.2%
1 CNN	40.7%	18.2%	—
5 CNNs	38.1%	16.4%	16.4%
1 CNN*	39.0%	16.6%	—
7 CNNs*	36.7%	15.4%	15.3%

Table 2: Comparison of error rates on ILSVRC-2012 validation and test sets. In *italics* are best results achieved by others. Models with an asterisk* were “pre-trained” to classify the entire ImageNet 2011 Fall release. See Section 6 for details.



Alexnet – key features

- First use of ReLU
- Overlapped max pooling
- Used normalization layers
 - Not used anymore



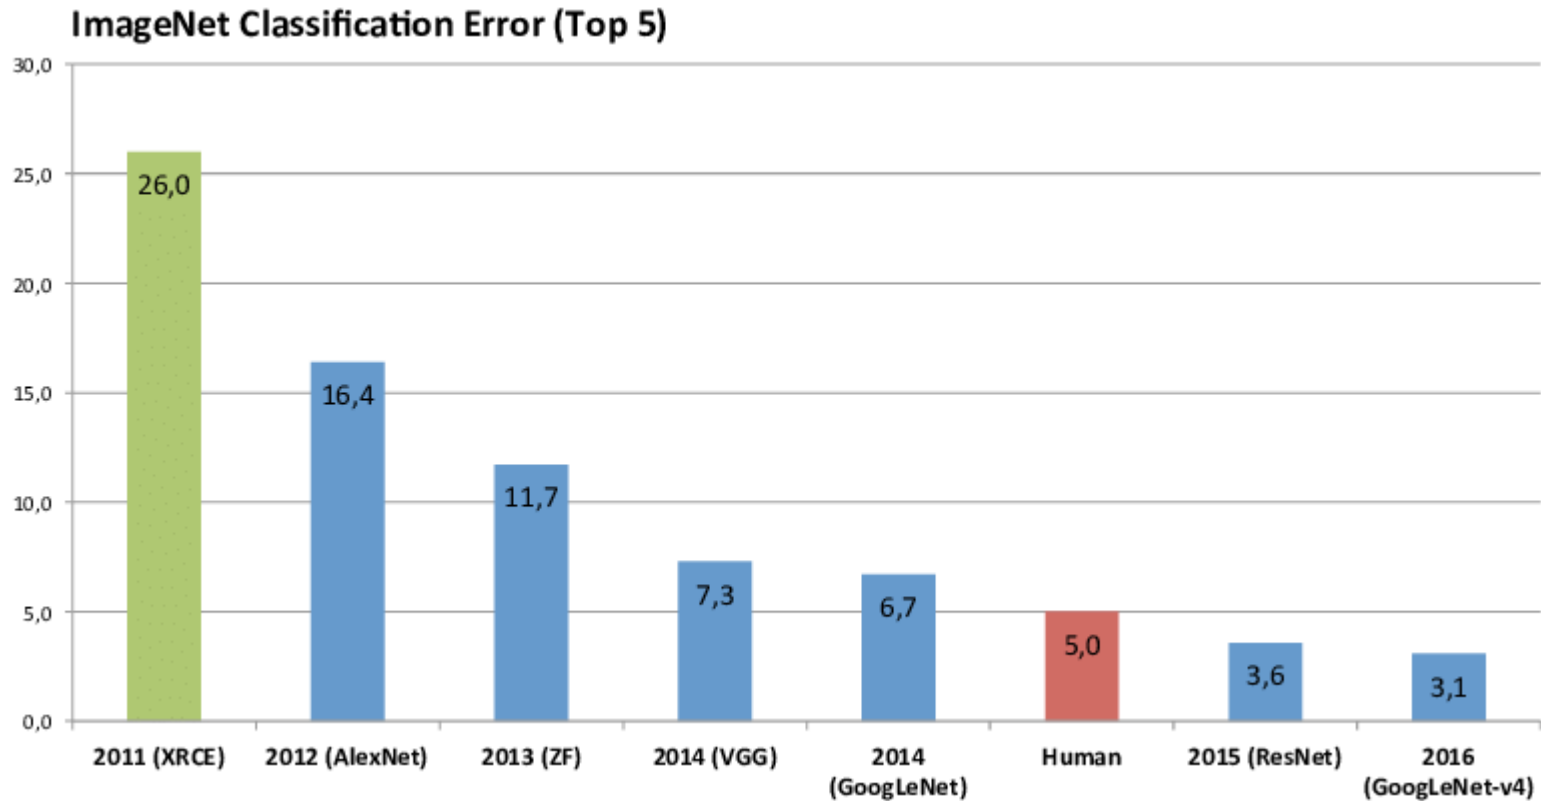
ReLU (solid line) effect on training
vs tanh (dashed line)

$$b_{x,y}^i = a_{x,y}^i / \left(k + \alpha \sum_{j=\max(0,i-n/2)}^{\min(N-1,i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

Alexnet – key features

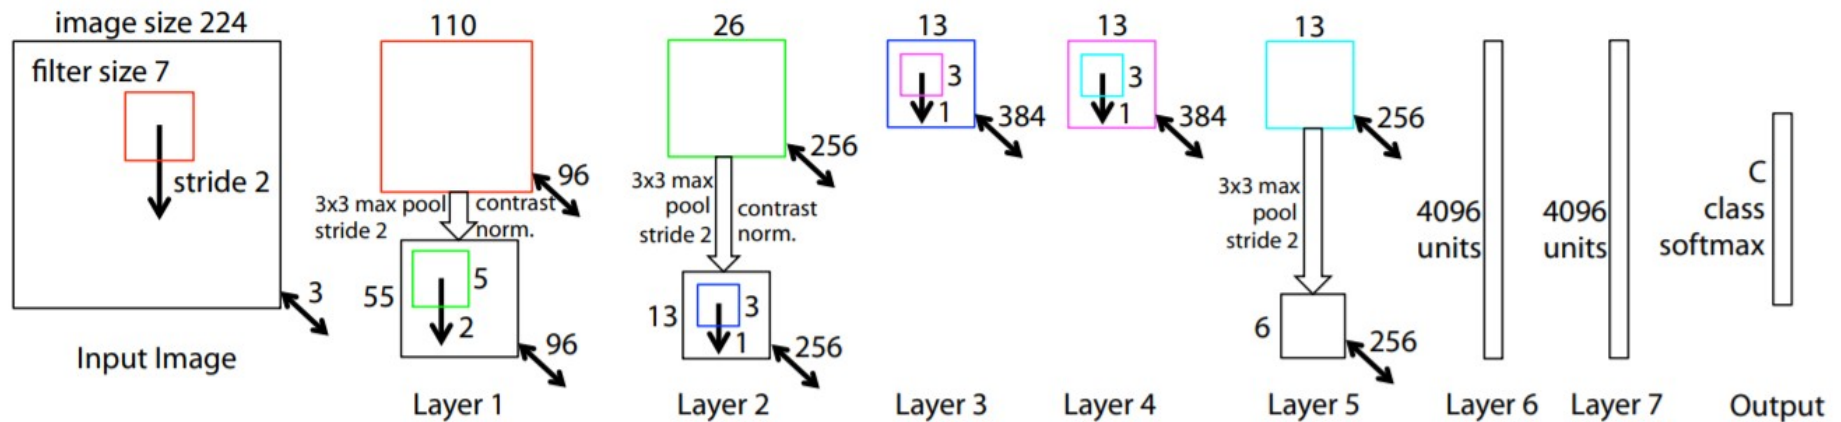
- Training setup
 - Dropout 0.5
 - Data augmentation
 - Batch size 128
 - Gradient descent with momentum ($\beta = 0.9$)
 - Initial learning rate: $1e-2$, reduced manually, when a plateau was reached
 - 7 Alexnet ensemble: 18.2% \rightarrow 15.4%

ZFNet, 2013



Built on top of Alexnet, improved hyperparameters

ZFNet



CONV1: change from (11x11 stride 4) to (7x7 stride 2)

CONV3, 4, 5: instead of 384, 384, 256 filters use 512, 1024, 512

ZFNet

Error %	Val Top-1	Val Top-5	Test Top-5
(Gunji et al., 2012)	-	-	26.2
(Krizhevsky et al., 2012), 1 convnet	40.7	18.2	--
(Krizhevsky et al., 2012), 5 convnets	38.1	16.4	16.4
(Krizhevsky et al., 2012)*, 1 convnets	39.0	16.6	--
(Krizhevsky et al., 2012)*, 7 convnets	36.7	15.4	15.3
Our replication of (Krizhevsky et al., 2012), 1 convnet	40.5	18.1	--
1 convnet as per Fig. 3	38.4	16.5	--
5 convnets as per Fig. 3 – (a)	36.7	15.3	15.3
1 convnet as per Fig. 3 but with layers 3,4,5: 512,1024,512 maps – (b)	37.5	16.0	16.1
6 convnets, (a) & (b) combined	36.0	14.7	14.8

CONV1: change from (11x11 stride 4) to (7x7 stride 2)

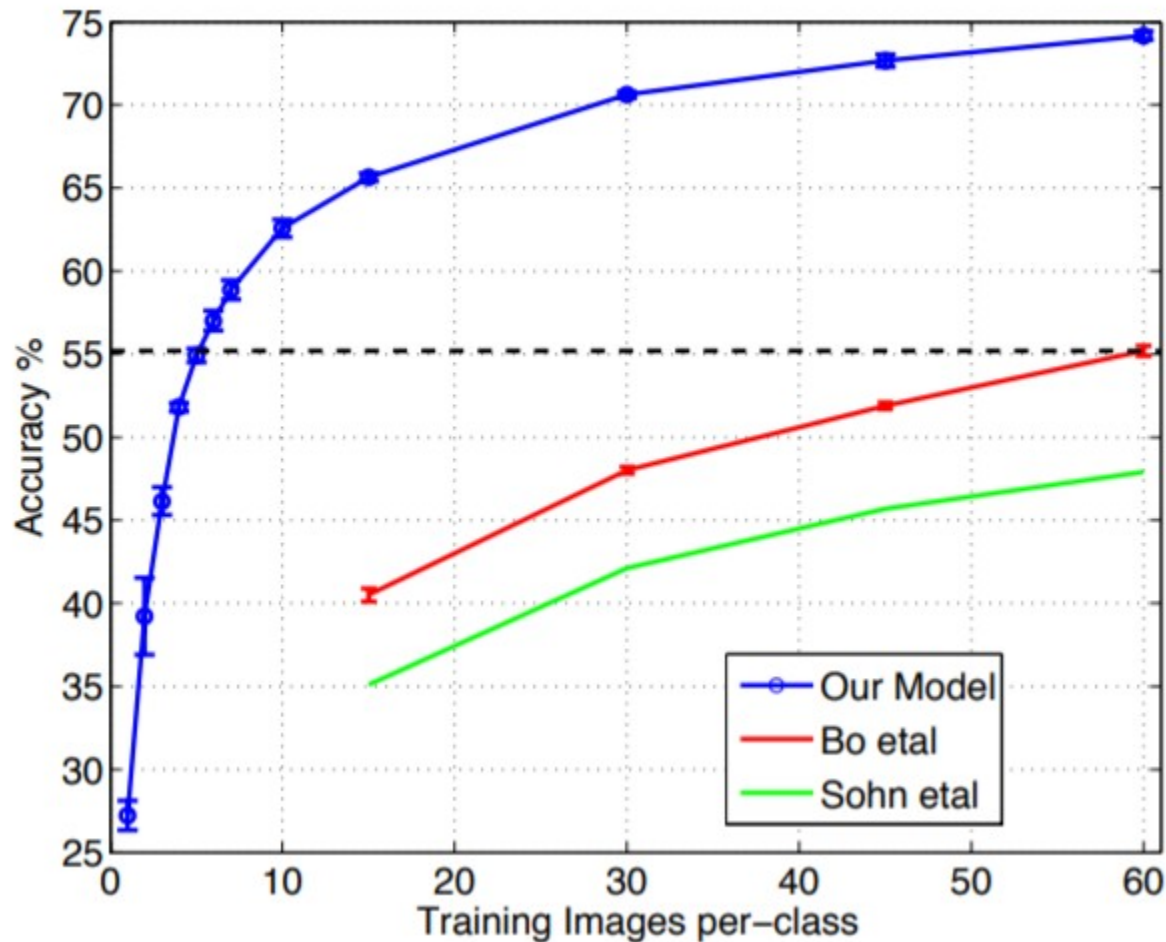
CONV3, 4, 5: instead of 384, 384, 256 filters use 512, 1024, 512

ImageNet top 5 error: 15.4% -> 14.8%

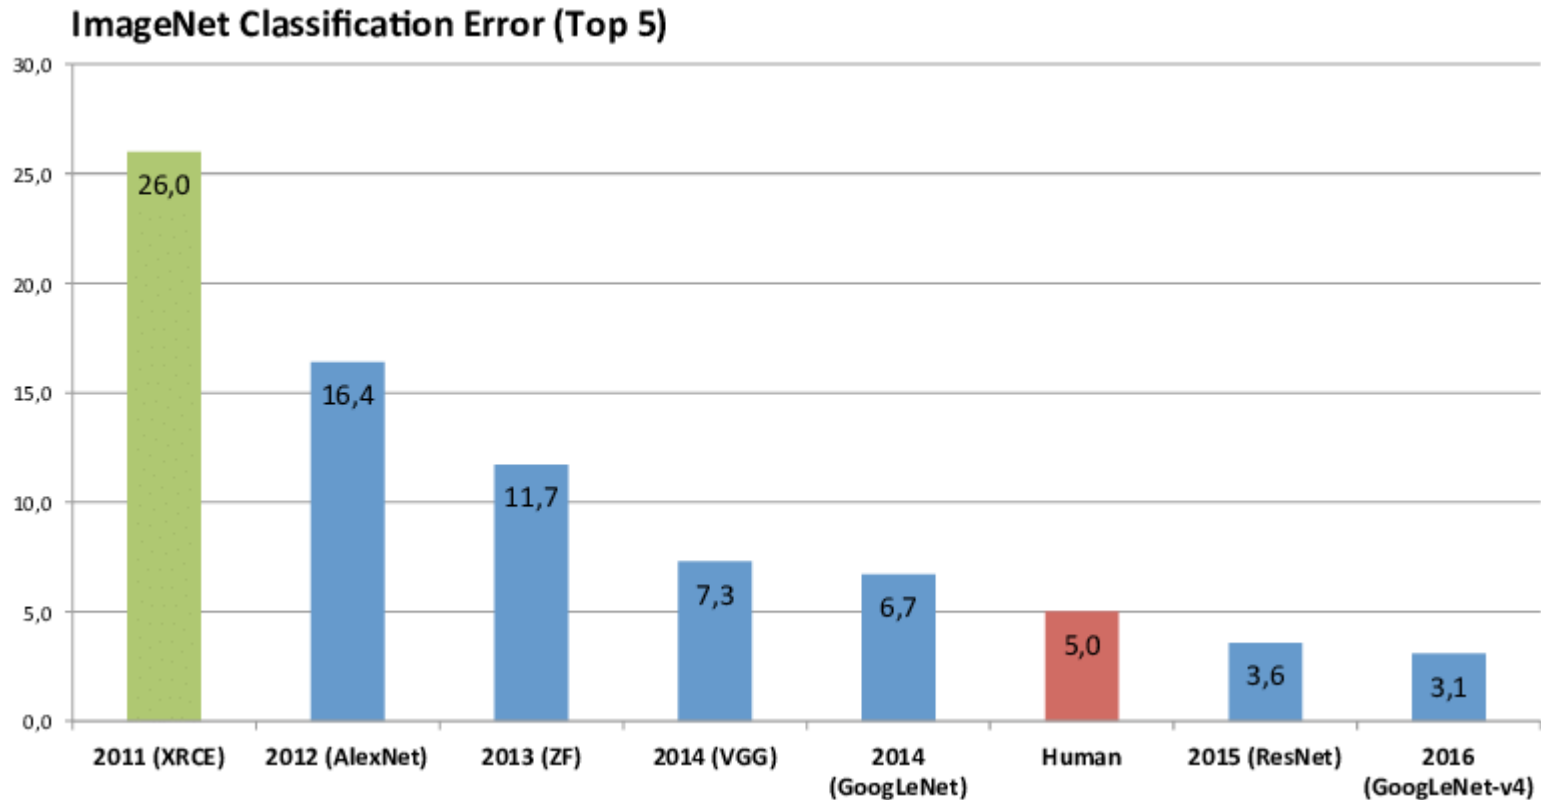


ZFNet – feature generalization

- Caltech 256



VGG, 2014



VGG

- Simplifies the network architectures
 - Always use 3x3 filters for convolutions
 - Always use max pooling of filter size 2x2 and a stride of 2

VGG

- Always use 3x3 filters for convolutions
- Always use max pooling of filter size 2x2 and a stride of 2

Top 5 accuracy: 11.7 -> 7.3

ConvNet Configuration					
A	A-LRN	B	C	D	E
11 weight layers	11 weight layers	13 weight layers	16 weight layers	16 weight layers	19 weight layers
input (224×224 RGB image)					
conv3-64	conv3-64 LRN	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64	conv3-64 conv3-64
maxpool					
conv3-128	conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128	conv3-128 conv3-128
maxpool					
conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256	conv3-256 conv3-256 conv1-256	conv3-256 conv3-256 conv3-256	conv3-256 conv3-256 conv3-256 conv3-256
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512	conv3-512 conv3-512 conv1-512	conv3-512 conv3-512 conv3-512	conv3-512 conv3-512 conv3-512 conv3-512
maxpool					
FC-4096					
FC-4096					
FC-1000					
soft-max					

Table 2: **Number of parameters** (in millions).

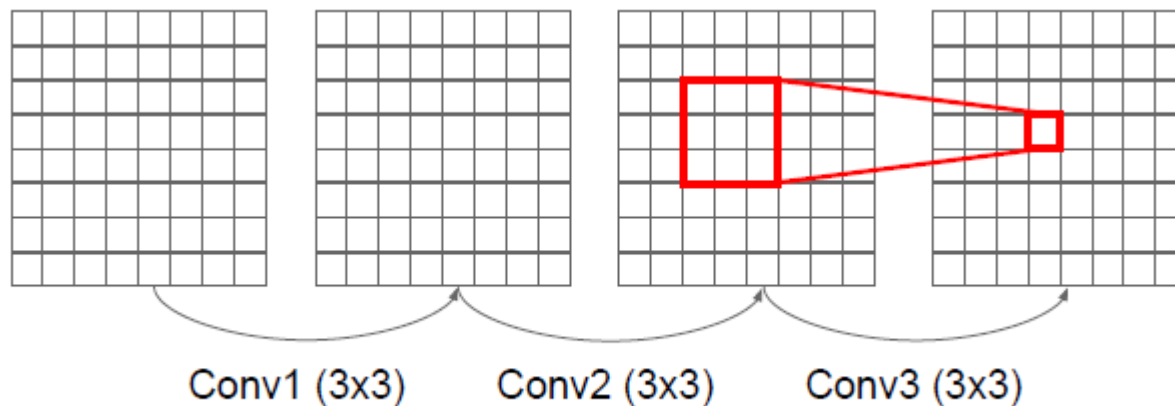
Network	A,A-LRN	B	C	D	E
Number of parameters	133	133	134	138	144

VGG

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 224, 224, 64)	1792
conv2d_2 (Conv2D)	(None, 224, 224, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 112, 112, 64)	0
conv2d_3 (Conv2D)	(None, 112, 112, 128)	73856
conv2d_4 (Conv2D)	(None, 112, 112, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 56, 56, 128)	0
conv2d_5 (Conv2D)	(None, 56, 56, 256)	295168
conv2d_6 (Conv2D)	(None, 56, 56, 256)	590080
conv2d_7 (Conv2D)	(None, 56, 56, 256)	590080
max_pooling2d_3 (MaxPooling2D)	(None, 28, 28, 256)	0
conv2d_8 (Conv2D)	(None, 28, 28, 512)	1180160
conv2d_9 (Conv2D)	(None, 28, 28, 512)	2359808
conv2d_10 (Conv2D)	(None, 28, 28, 512)	2359808
max_pooling2d_4 (MaxPooling2D)	(None, 14, 14, 512)	0
conv2d_11 (Conv2D)	(None, 14, 14, 512)	2359808
conv2d_12 (Conv2D)	(None, 14, 14, 512)	2359808
conv2d_13 (Conv2D)	(None, 14, 14, 512)	2359808
max_pooling2d_5 (MaxPooling2D)	(None, 7, 7, 512)	0
flatten_1 (Flatten)	(None, 25088)	0
dense_1 (Dense)	(None, 4096)	102764544
dropout_1 (Dropout)	(None, 4096)	0
dense_2 (Dense)	(None, 4096)	16781312
dropout_2 (Dropout)	(None, 4096)	0
dense_3 (Dense)	(None, 2)	8194
Total params: 134,268,738		
Trainable params: 134,268,738		
Non-trainable params: 0		

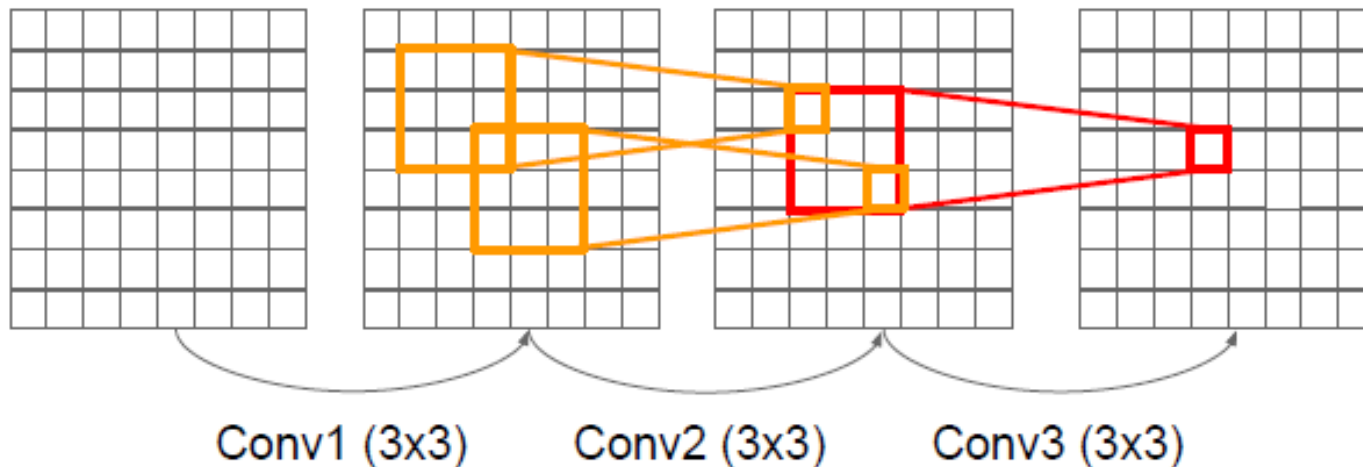
VGG

- “It is easy to see that a stack of two 3×3 conv. layers has an effective receptive field of 5×5 ; three such layers have a 7×7 effective receptive field.”*



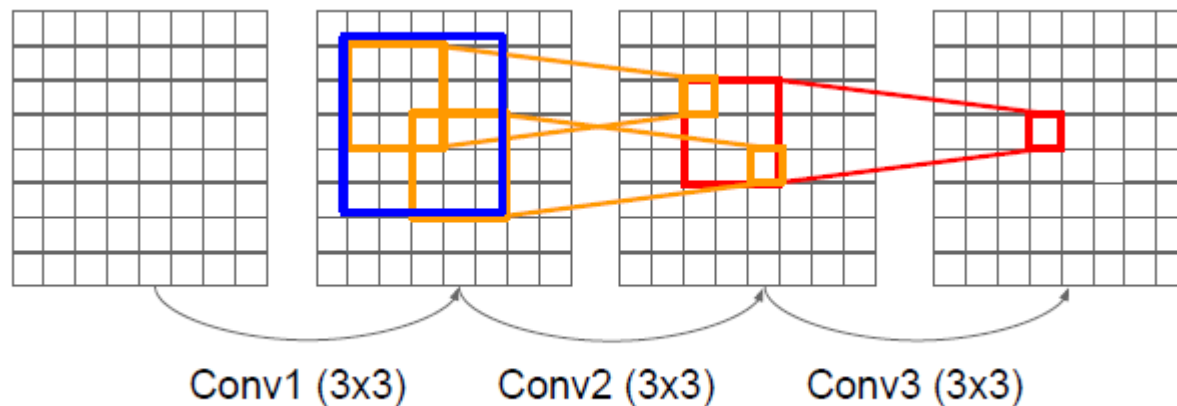
VGG

- “It is easy to see that a stack of two 3×3 conv. layers has an effective receptive field of 5×5 ; three such layers have a 7×7 effective receptive field.”*



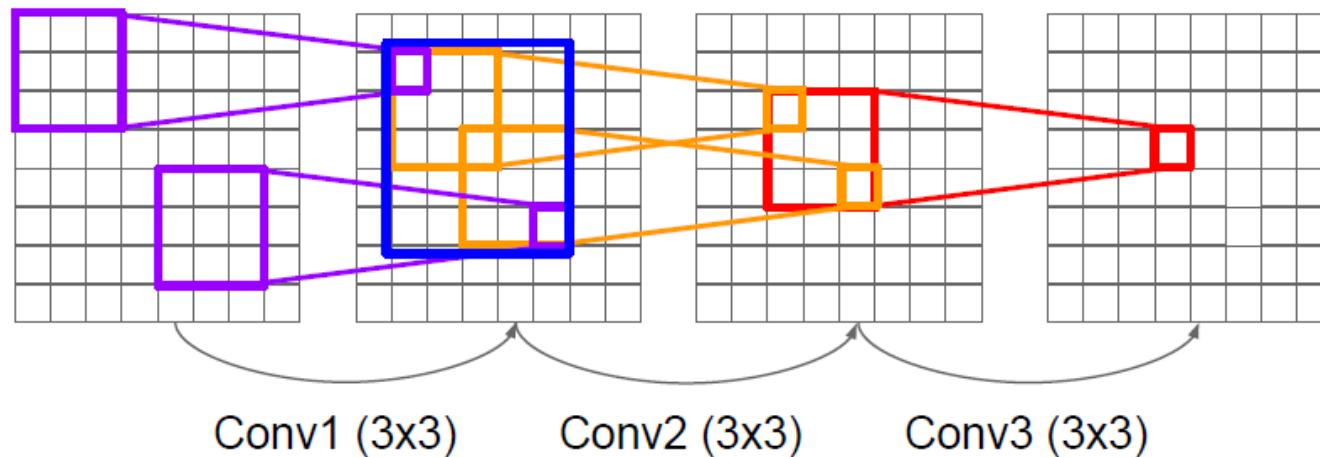
VGG

- “It is easy to see that a stack of two 3×3 conv. layers has an effective receptive field of 5×5 ; three such layers have a 7×7 effective receptive field.”*



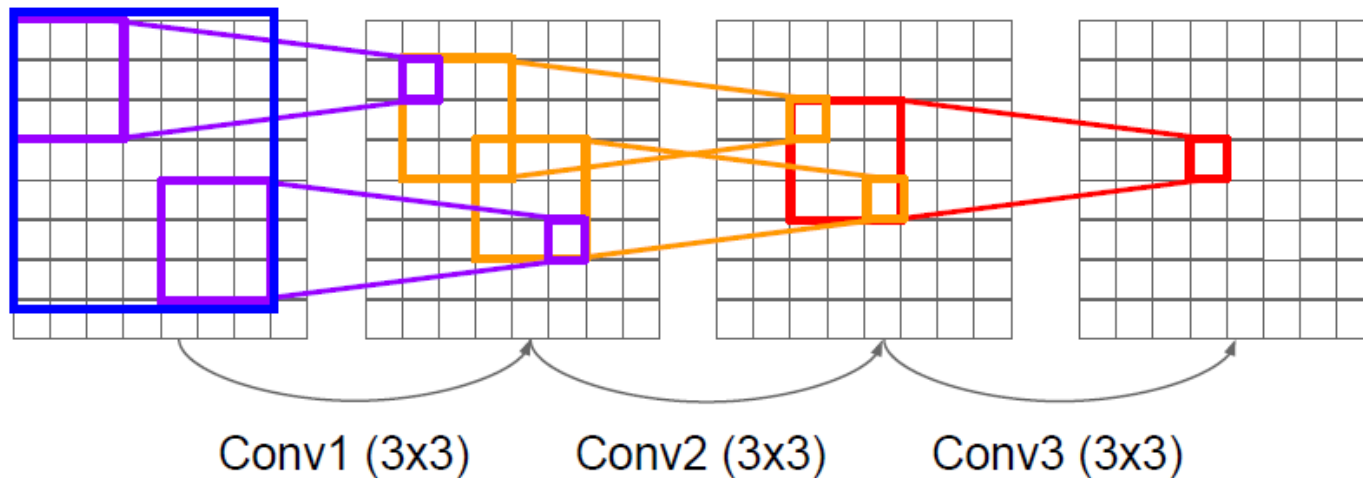
VGG

- “It is easy to see that a stack of two 3×3 conv. layers has an effective receptive field of 5×5 ; three such layers have a 7×7 effective receptive field.”*



VGG

- “It is easy to see that a stack of two 3×3 conv. layers has an effective receptive field of 5×5 ; three such layers have a 7×7 effective receptive field.”*



VGG

- *“It is easy to see that a stack of two 3×3 conv. layers has an effective receptive field of 5×5 ; three such layers have a 7×7 effective receptive field.”*
- What is the effect of this replacement in terms on the number of parameters?
- What about non-linearities?

VGG

- *“It is easy to see that a stack of two 3×3 conv. layers has an effective receptive field of 5×5 ; three such layers have a 7×7 effective receptive field.”*
- What is the effect of this replacement in terms on the number of parameters?

$$3 \cdot (3^2 C) = 27C \text{ vs } 7^2 C = 49C$$

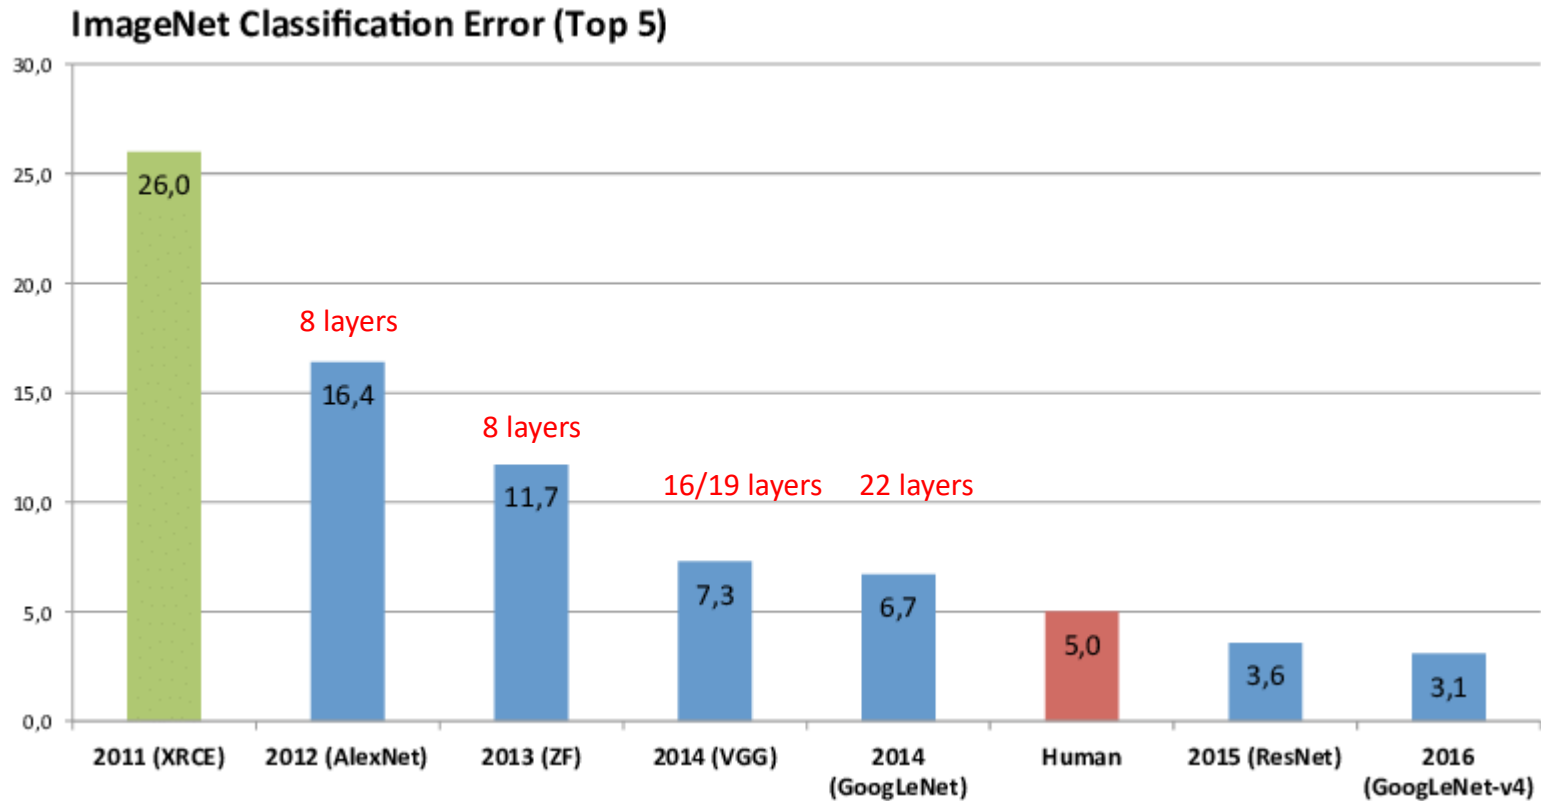
- What about non-linearities?

VGG – key features

- Did not use local response normalization layers
- Use ensembles to boost performance
- Use VGG 16 or VGG 19 (VGG 19 brings only a small increase in performance, but it requires more memory)
- Similar training procedure as in AlexNet



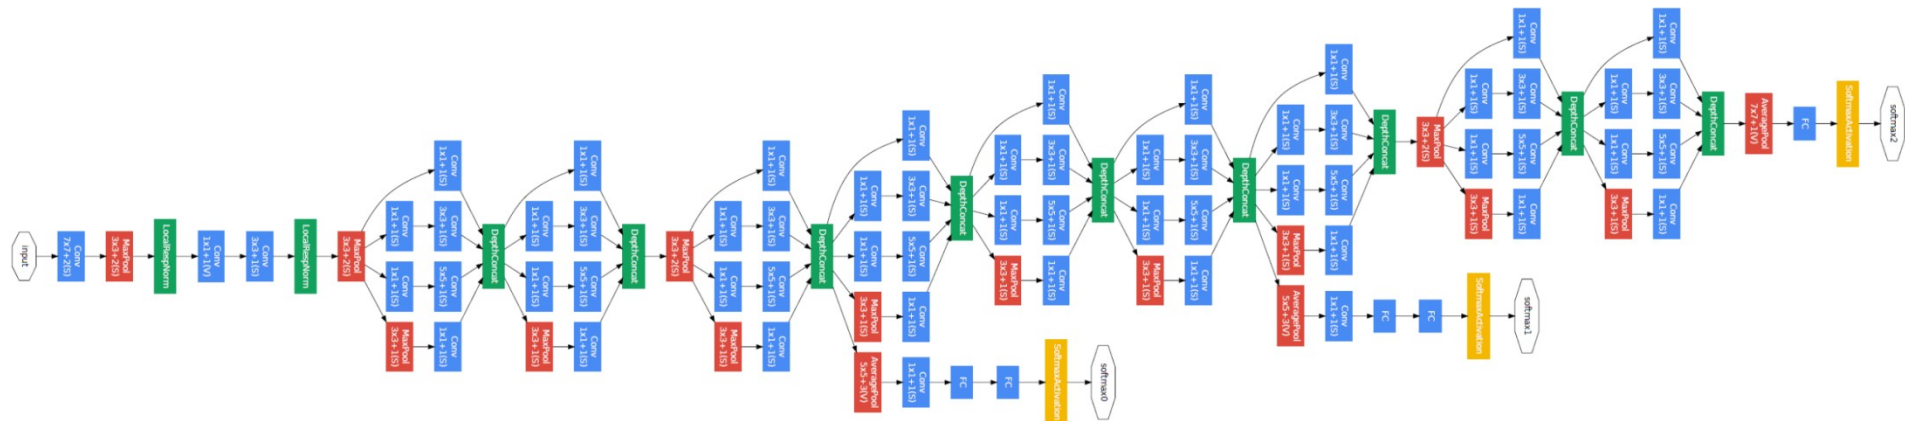
GoogLeNet, 2014



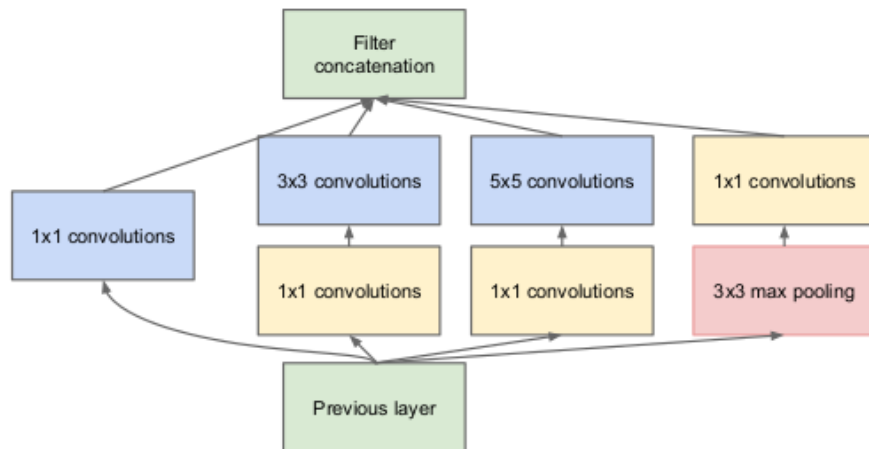
6.7 top-5 accuracy

GoogLeNet

This name is an homage to Yann LeCun's pioneering LeNet 5 network



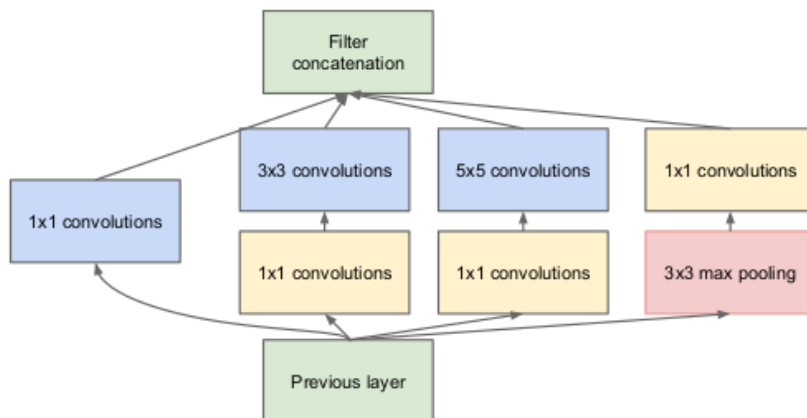
Inception module



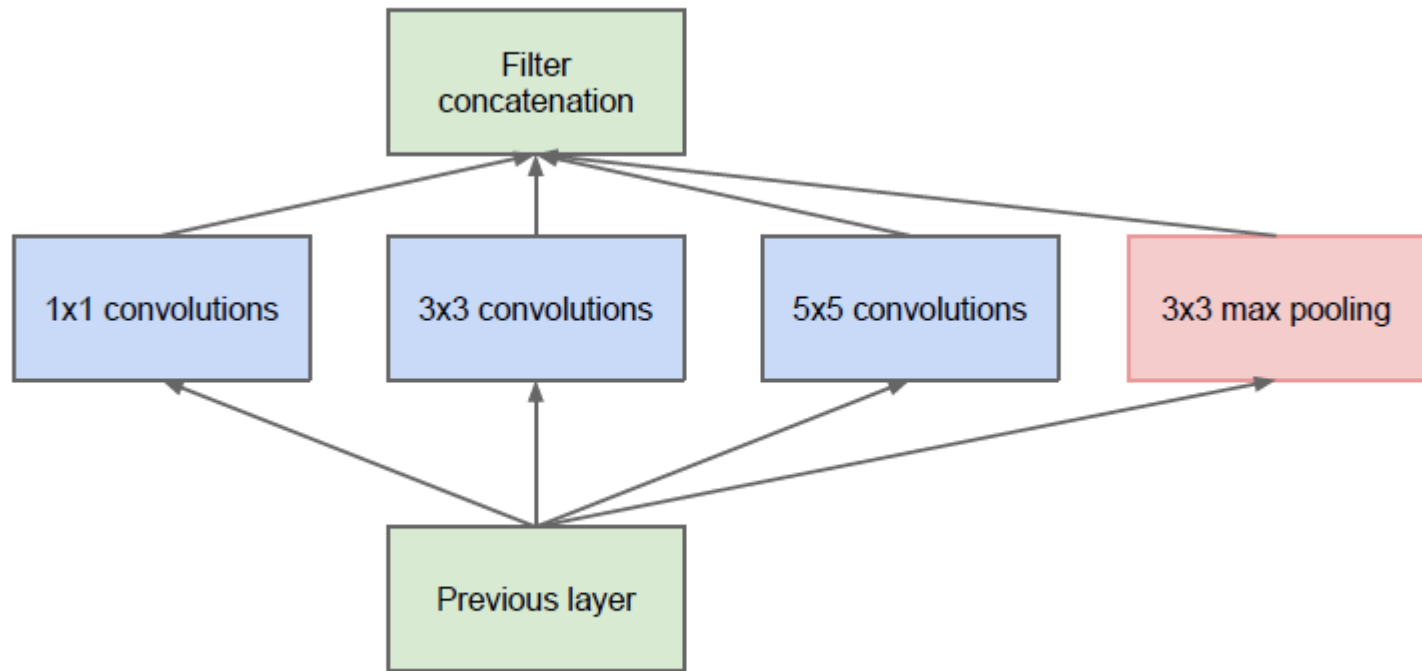
type	patch size/ stride	output size	depth	#1×1	#3×3 reduce	#3×3	#5×5 reduce	#5×5	pool proj	params	ops
convolution	7×7/2	112×112×64	1							2.7K	34M
max pool	3×3/2	56×56×64	0								
convolution	3×3/1	56×56×192	2		64	192				112K	360M
max pool	3×3/2	28×28×192	0								
inception (3a)		28×28×256	2	64	96	128	16	32	32	159K	128M
inception (3b)		28×28×480	2	128	128	192	32	96	64	380K	304M
max pool	3×3/2	14×14×480	0								
inception (4a)		14×14×512	2	192	96	208	16	48	64	364K	73M
inception (4b)		14×14×512	2	160	112	224	24	64	64	437K	88M
inception (4c)		14×14×512	2	128	128	256	24	64	64	463K	100M
inception (4d)		14×14×528	2	112	144	288	32	64	64	580K	119M
inception (4e)		14×14×832	2	256	160	320	32	128	128	840K	170M
max pool	3×3/2	7×7×832	0								
inception (5a)		7×7×832	2	256	160	320	32	128	128	1072K	54M
inception (5b)		7×7×1024	2	384	192	384	48	128	128	1388K	71M
avg pool	7×7/1	1×1×1024	0								
dropout (40%)		1×1×1024	0								
linear		1×1×1000	1							1000K	1M
softmax		1×1×1000	0								

Inception module

- You don't need to "pick" the filter sizes, instead you let the network "choose" between several values
 - Filters of different sizes
 - Pooling layer
 - Concatenate activations depth-wise



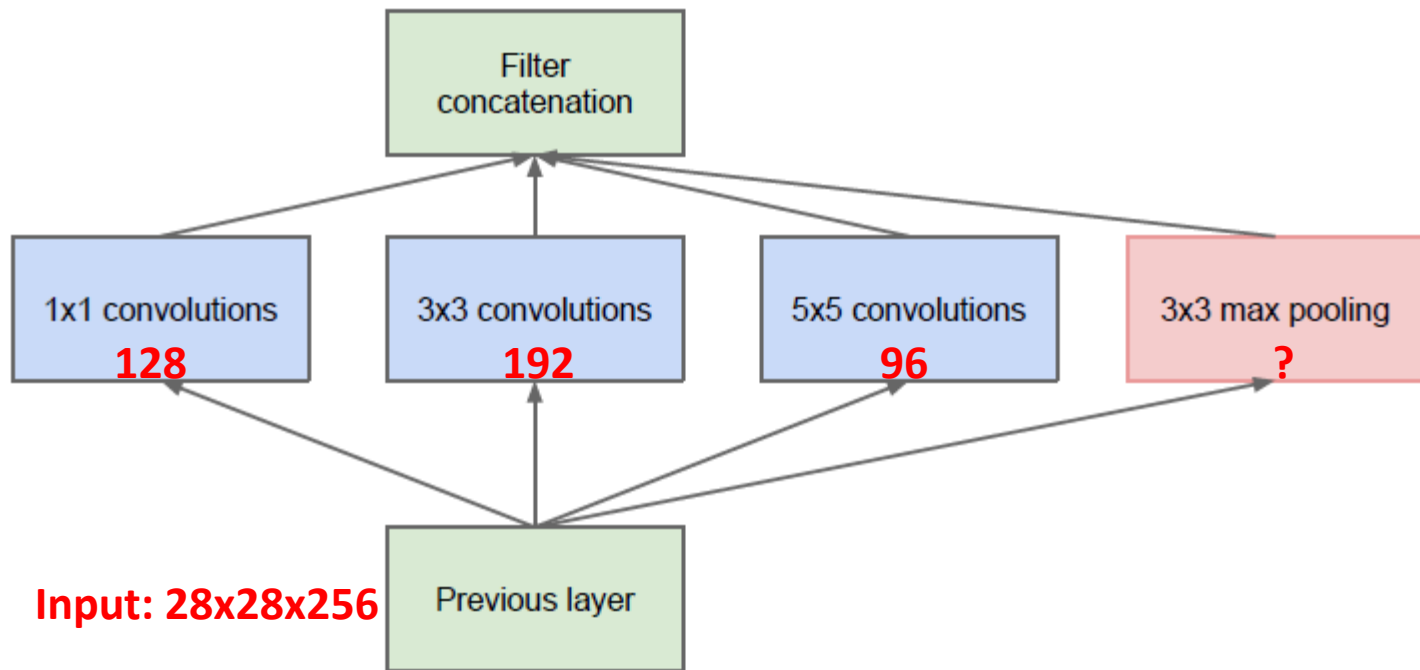
Inception module



(a) Inception module, naïve version

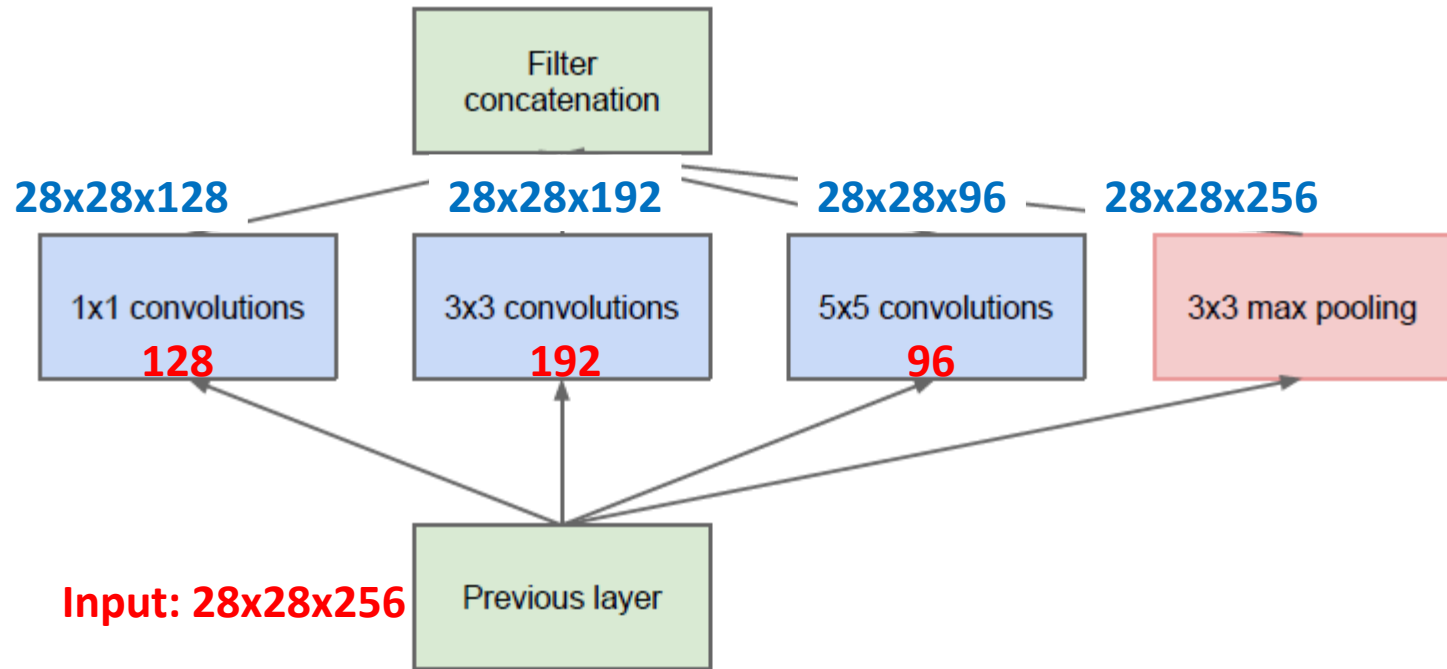
Inception module

What are the output sizes of all these operations?



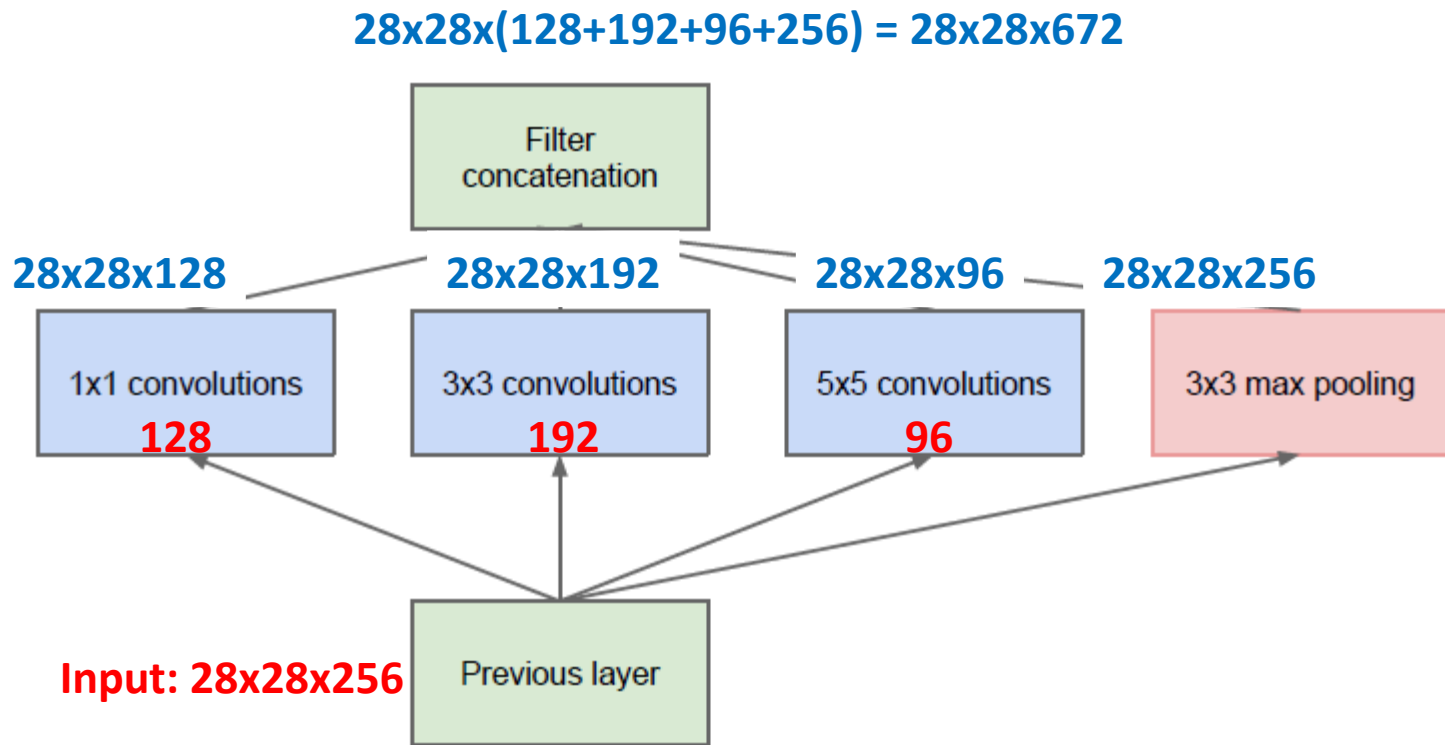
(a) Inception module, naïve version

Inception module



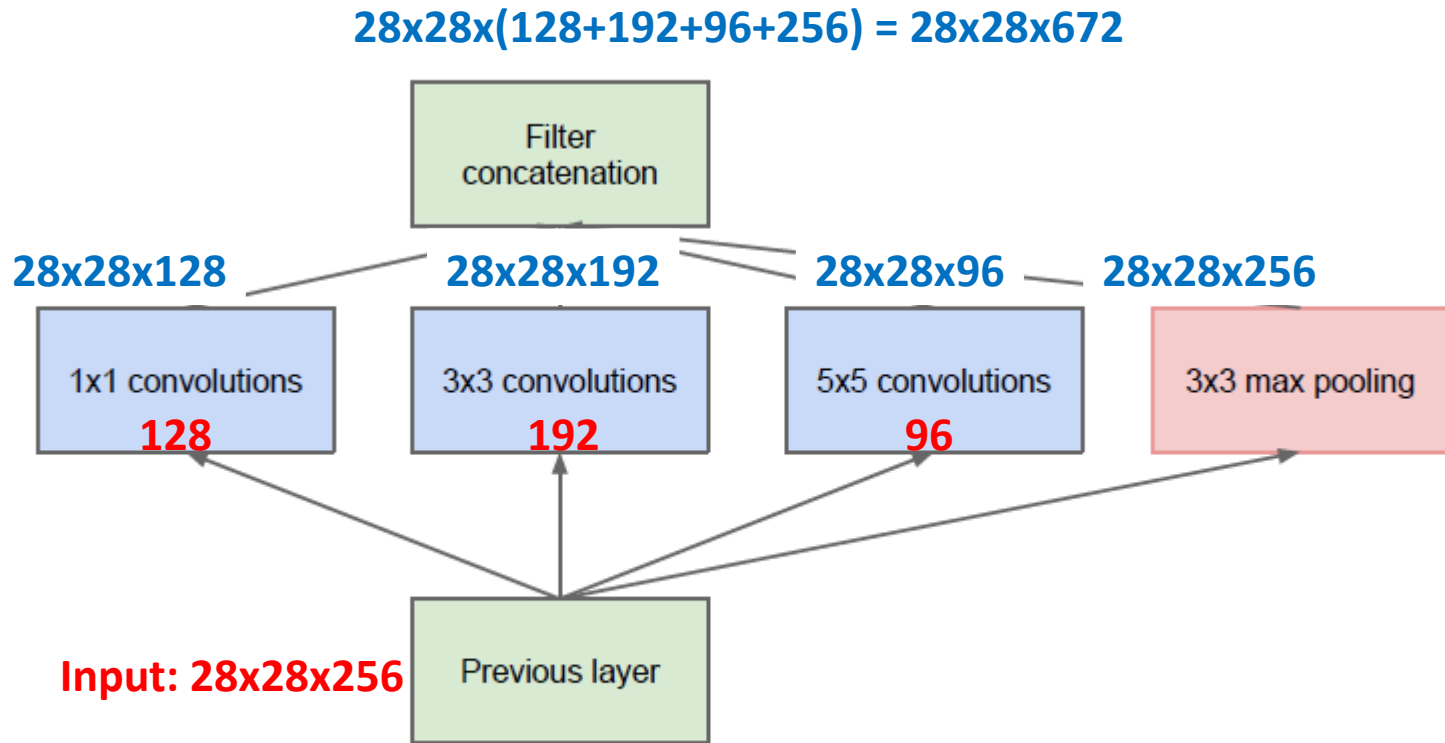
(a) Inception module, naïve version

Inception module



(a) Inception module, naïve version

Inception module

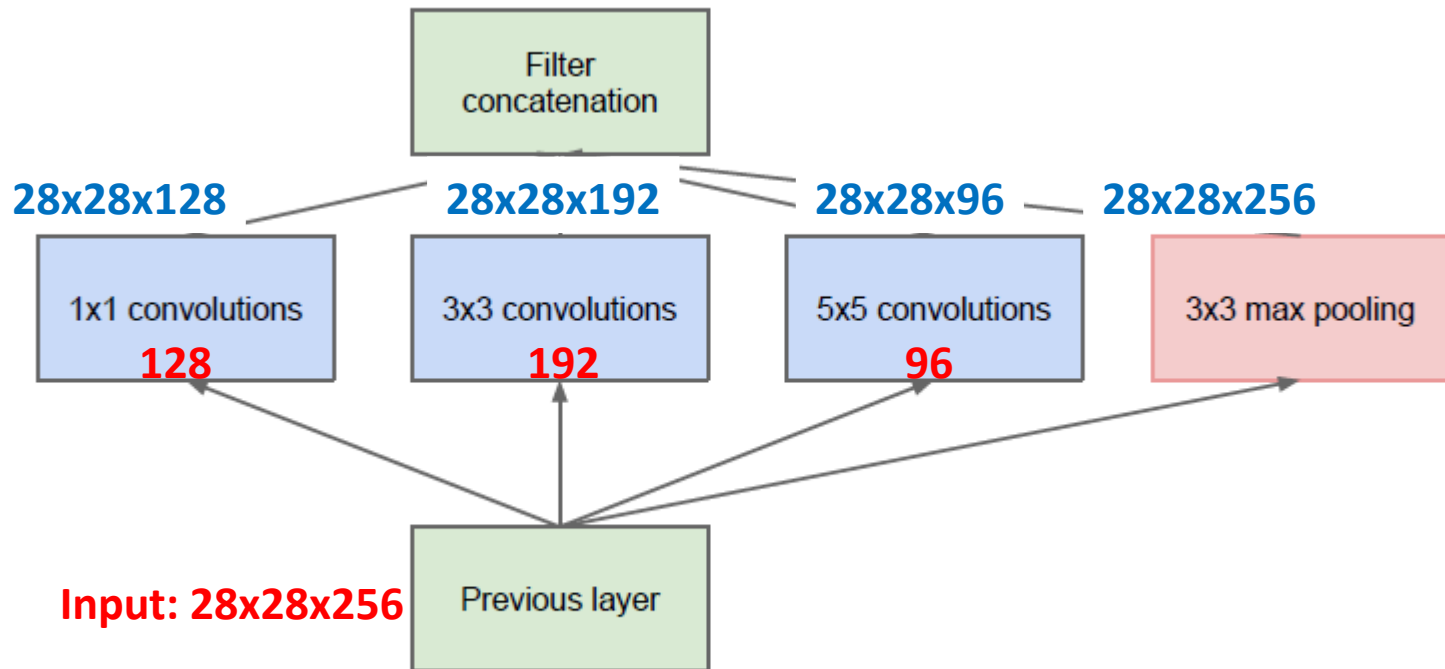


(a) Inception module, naïve version

[1x1 conv, 128 filters]: ? convolution operations
[3x3 conv, 192 filters]: ? convolution operations
[5x5 conv, 96 filters]: ? convolution operations

Inception module

$$28 \times 28 \times (128 + 192 + 96 + 256) = 28 \times 28 \times 672 = 526848$$

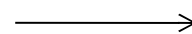


(a) Inception module, naïve version

[1x1 conv, 128 filters]: $28 \times 28 \times 128 \times 1 \times 1 \times 256$

[3x3 conv, 192 filters]: $28 \times 28 \times 192 \times 3 \times 3 \times 256$

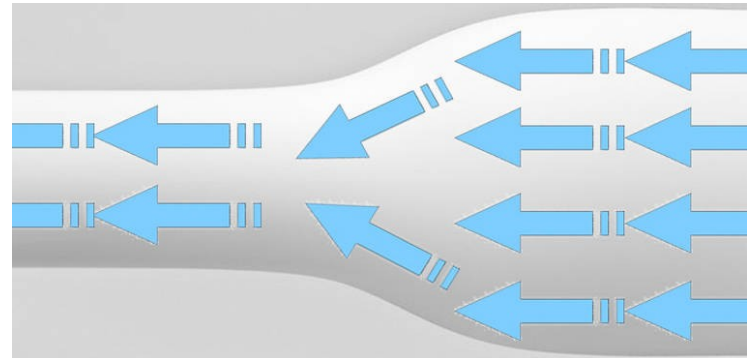
[5x5 conv, 96 filters]: $28 \times 28 \times 96 \times 5 \times 5 \times 256$



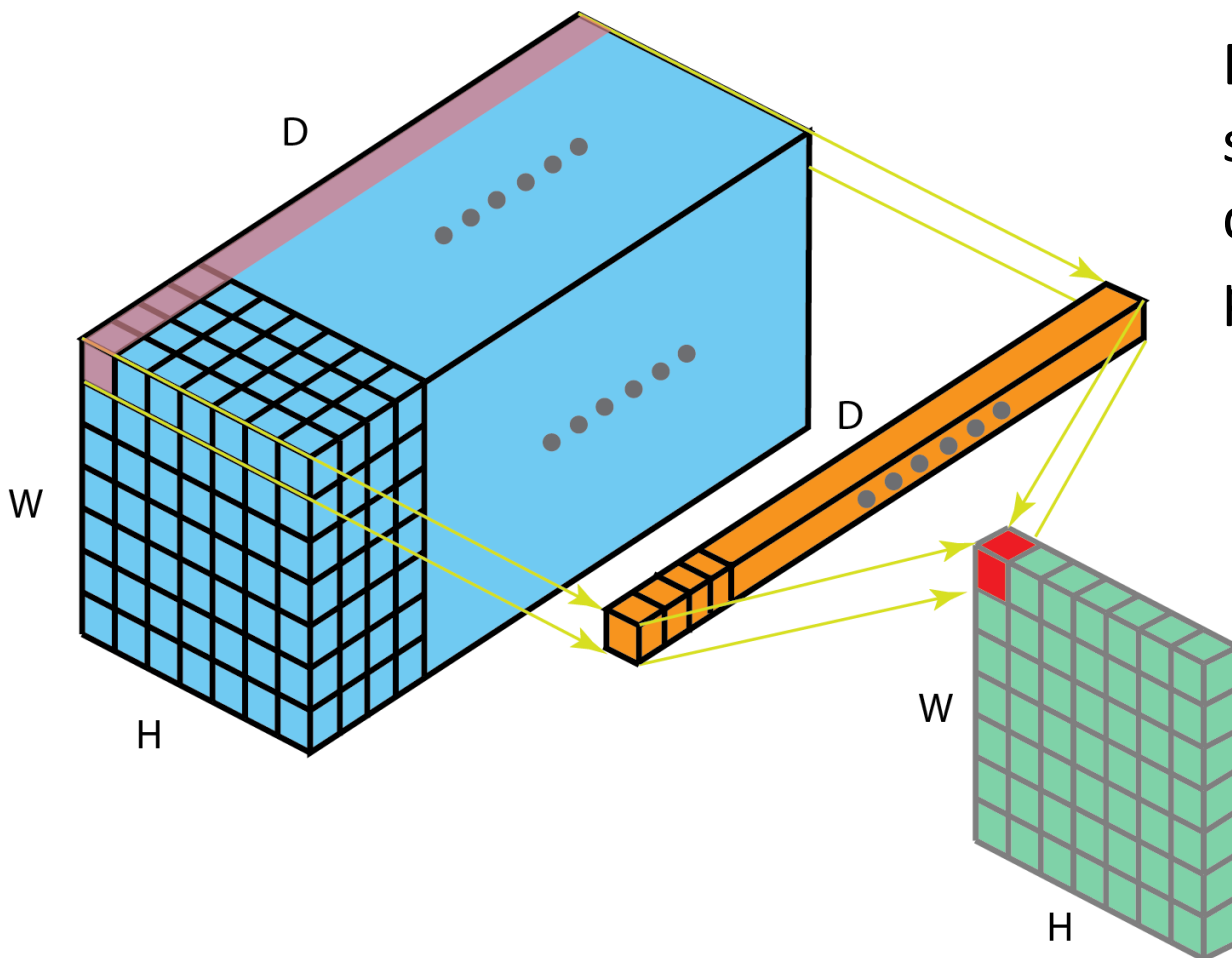
>850M

Bottleneck

- Use *bottleneck* layers to reduce the depth dimensions of the activations
 - use 1x1 convolutions

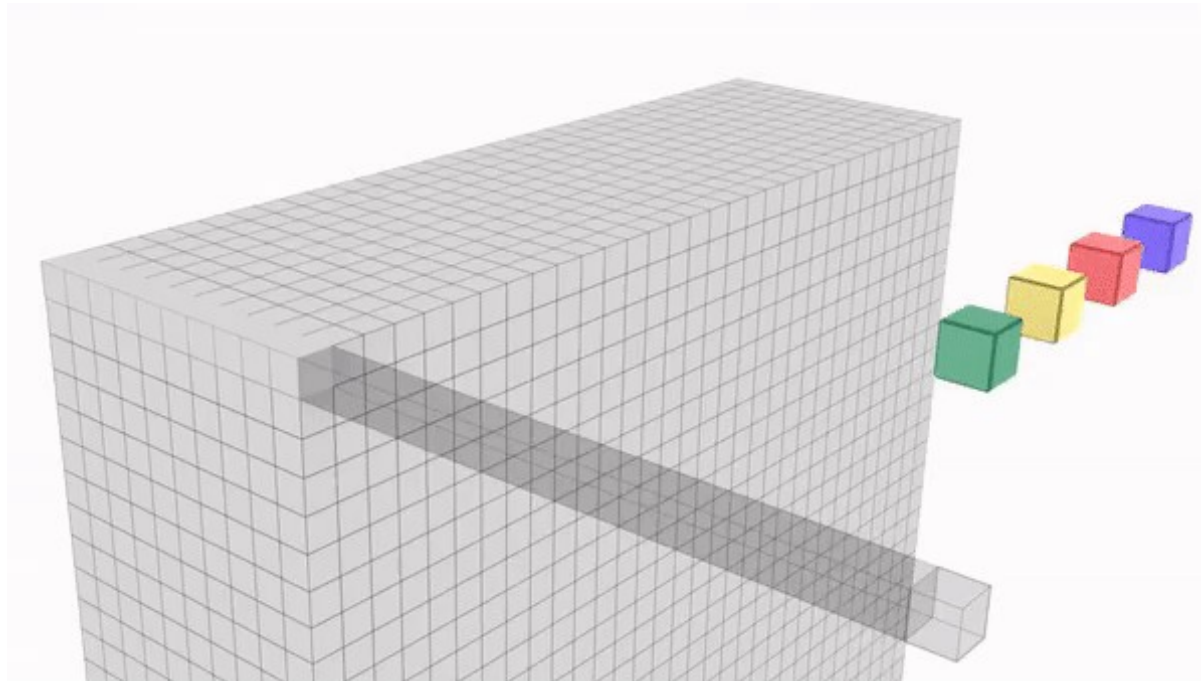


1x1 convolutions



Each filter has a $1 \times 1 \times D$ size and performs a D -dimensional dot product

1x1 convolutions



1x1 convolution



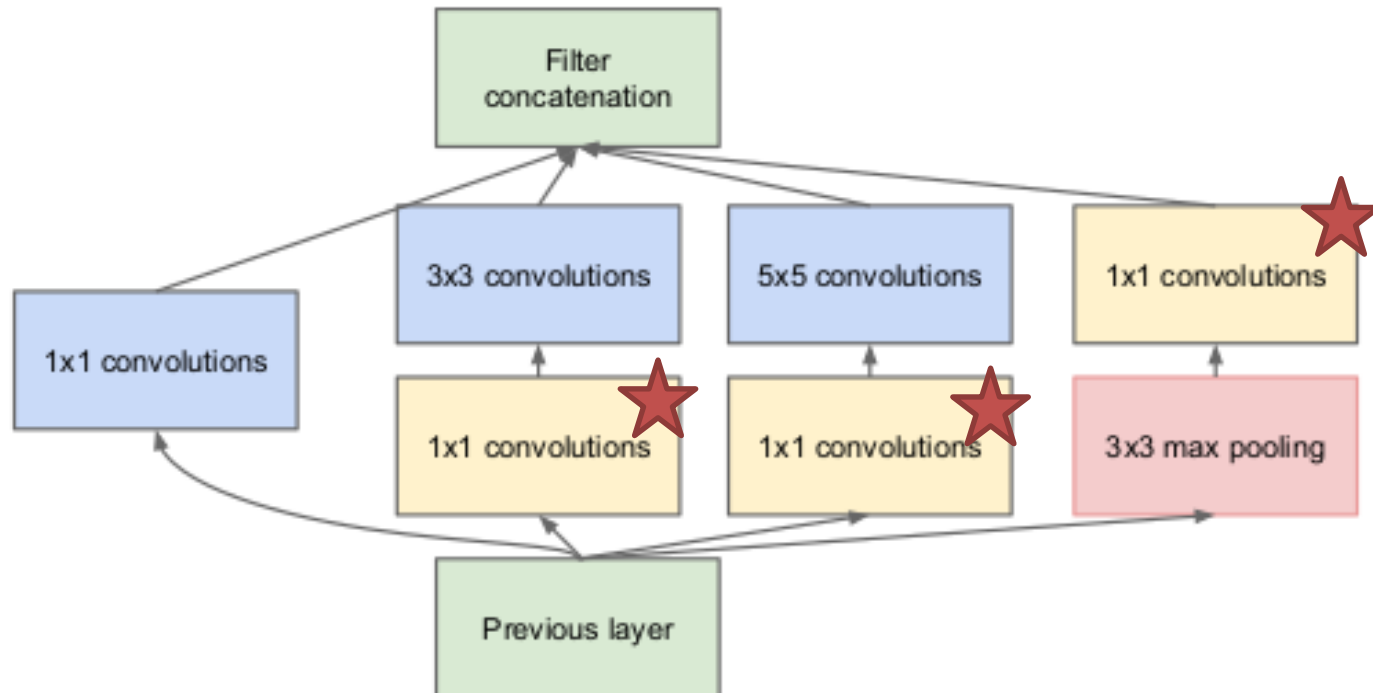
Yann LeCun

April 6, 2015 · 🌐

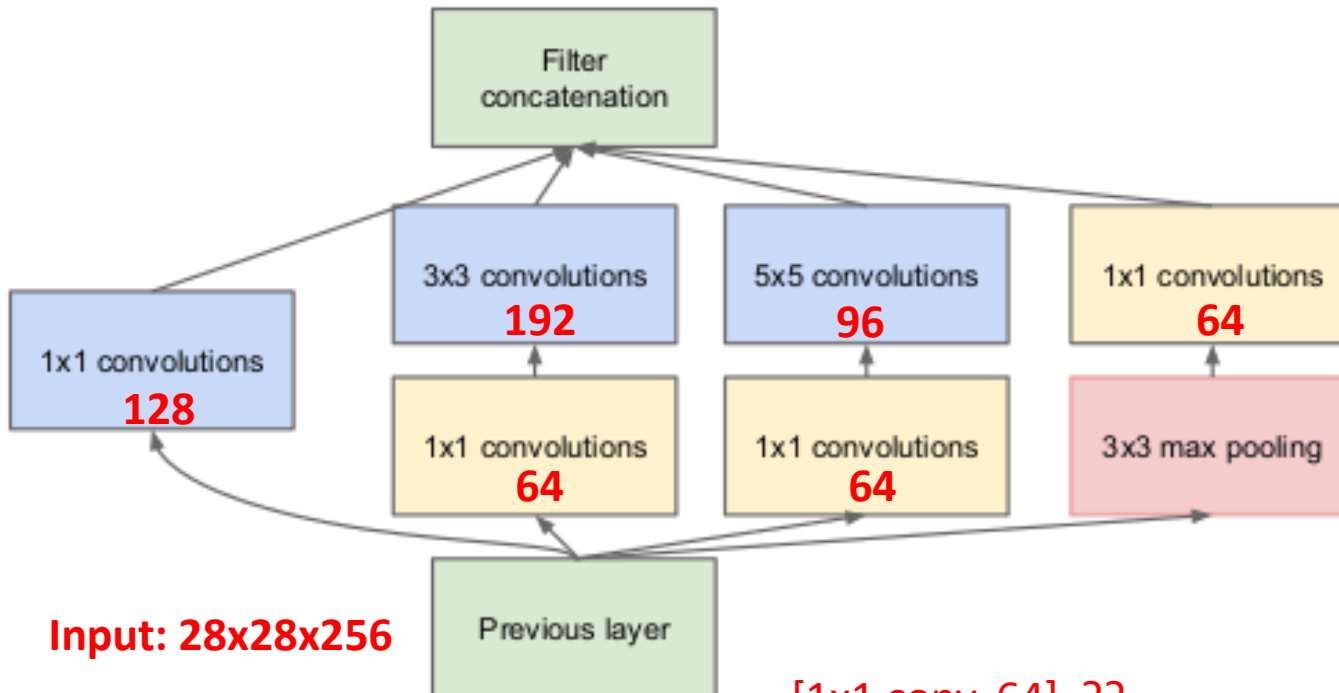
In Convolutional Nets, there is no such thing as "fully-connected layers". There are only convolution layers with 1x1 convolution kernels and a full connection table.

It's a too-rarely-understood fact that ConvNets don't need to have a fixed-size input. You can train them on inputs that happen to produce a single output vector (with no spatial extent), and then apply them to larger images. Instead of a single output vector, you then get a spatial map of output vectors. Each vector sees input windows at different locations on the input. In that scenario, the "fully connected layers" really act as 1x1 convolutions.

The real Inception module

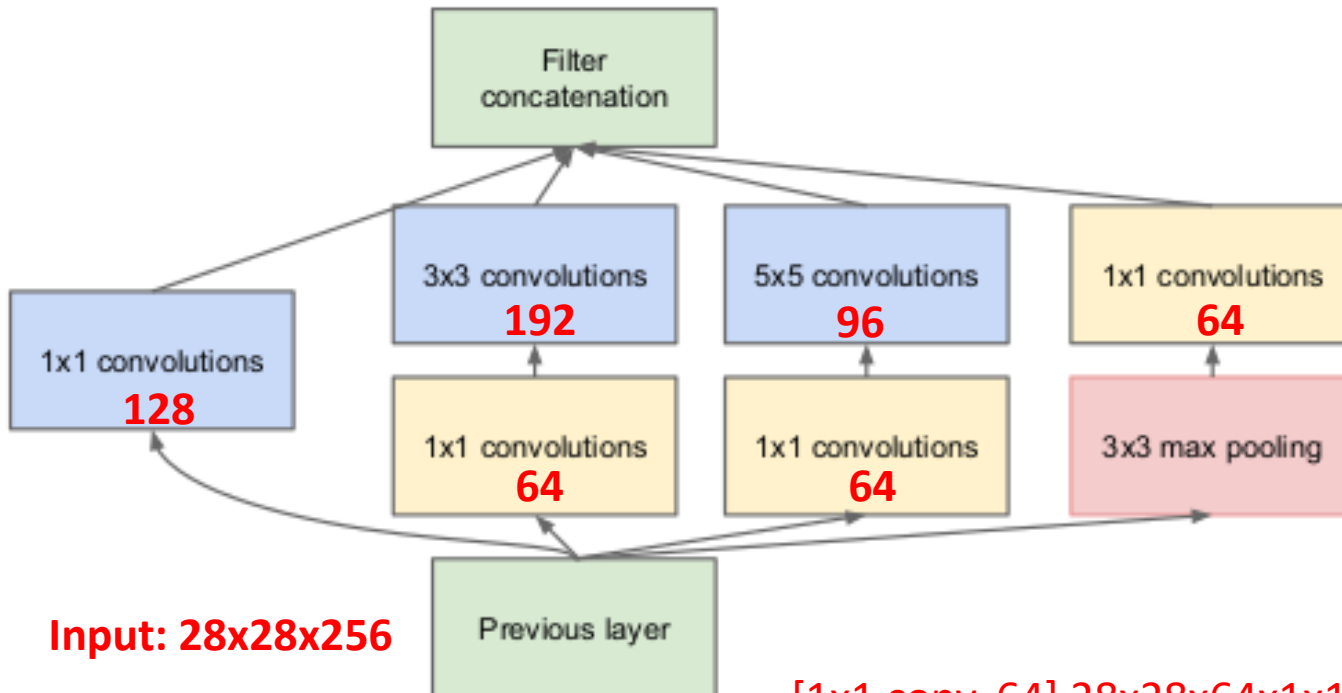


The real Inception module



[1x1 conv, 64] ??
[1x1 conv, 64] ??
[1x1 conv, 128] ??
[3x3 conv, 192] ??
[5x5 conv, 96] ??
[1x1 conv, 64] ??

The real Inception module



Input: 28x28x256

~~>850M~~

358M

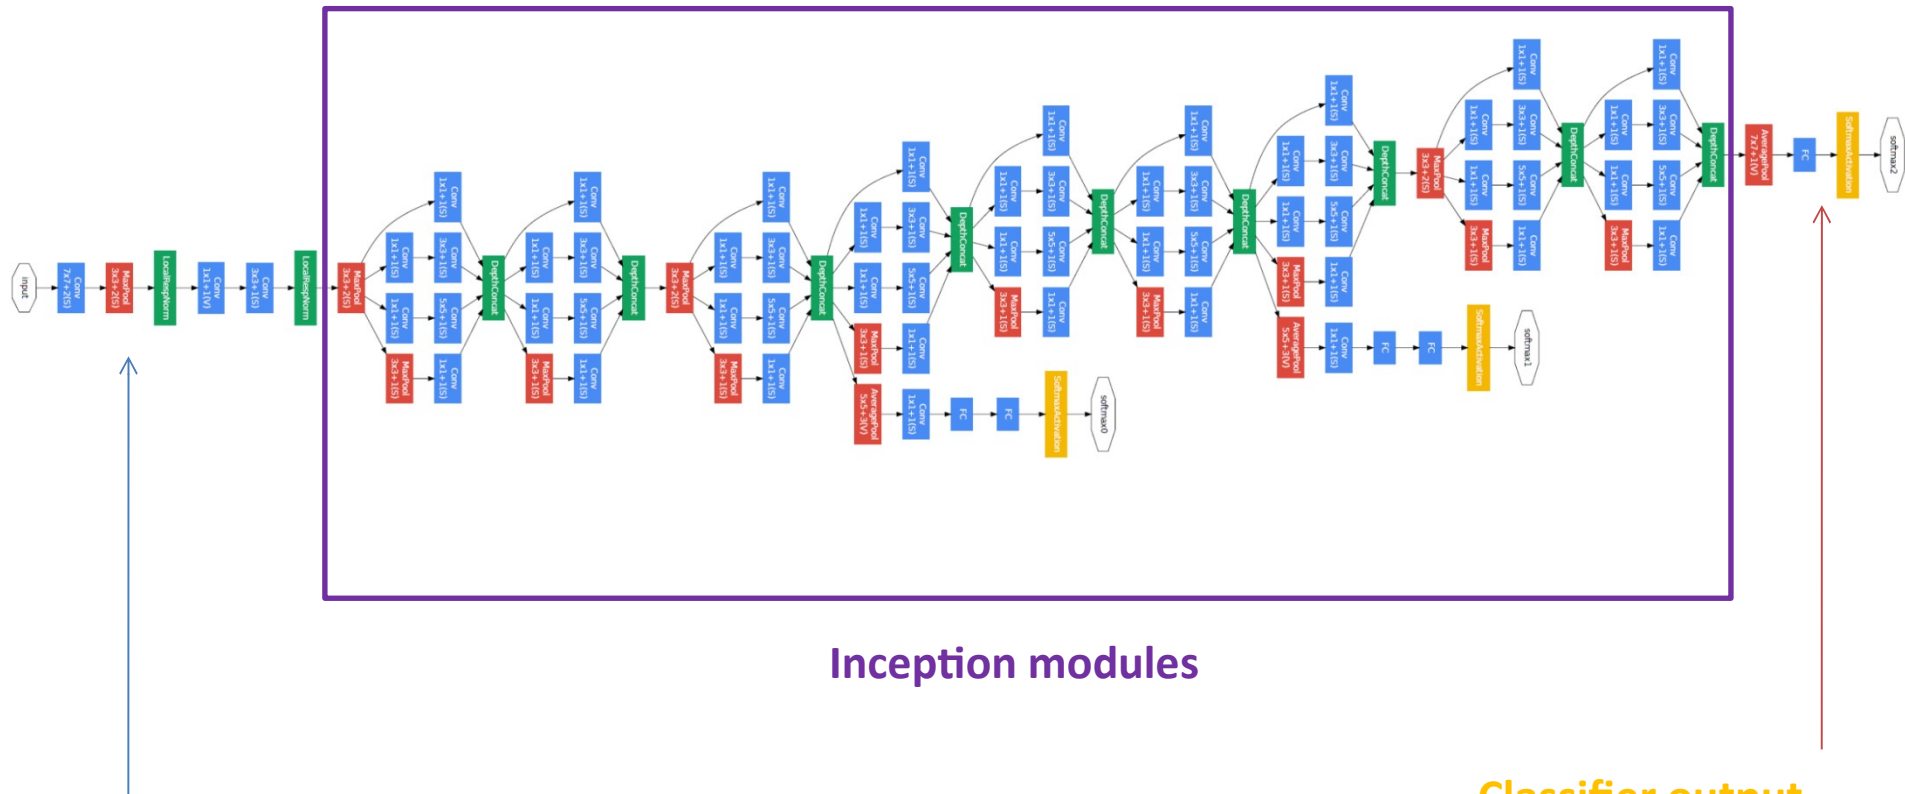
[1x1 conv, 64] 28x28x64x1x1x256
[1x1 conv, 64] 28x28x64x1x1x256
[1x1 conv, 128] 28x28x128x1x1x256
[3x3 conv, 192] 28x28x192x3x3x64
[5x5 conv, 96] 28x28x96x5x5x64
[1x1 conv, 64] 28x28x64x1x1x256

GoogLe Net

Stack Inception modules on top of each other

type	patch size/ stride	output size	depth	#1×1	#3×3 reduce	#3×3	#5×5 reduce	#5×5	pool proj	params	ops
convolution	7×7/2	112×112×64	1							2.7K	34M
max pool	3×3/2	56×56×64	0								
convolution	3×3/1	56×56×192	2		64	192				112K	360M
max pool	3×3/2	28×28×192	0								
inception (3a)		28×28×256	2	64	96	128	16	32	32	159K	128M
inception (3b)		28×28×480	2	128	128	192	32	96	64	380K	304M
max pool	3×3/2	14×14×480	0								
inception (4a)		14×14×512	2	192	96	208	16	48	64	364K	73M
inception (4b)		14×14×512	2	160	112	224	24	64	64	437K	88M
inception (4c)		14×14×512	2	128	128	256	24	64	64	463K	100M
inception (4d)		14×14×528	2	112	144	288	32	64	64	580K	119M
inception (4e)		14×14×832	2	256	160	320	32	128	128	840K	170M
max pool	3×3/2	7×7×832	0								
inception (5a)		7×7×832	2	256	160	320	32	128	128	1072K	54M
inception (5b)		7×7×1024	2	384	192	384	48	128	128	1388K	71M
avg pool	7×7/1	1×1×1024	0								
dropout (40%)		1×1×1024	0								
linear		1×1×1000	1							1000K	1M
softmax		1×1×1000	0								

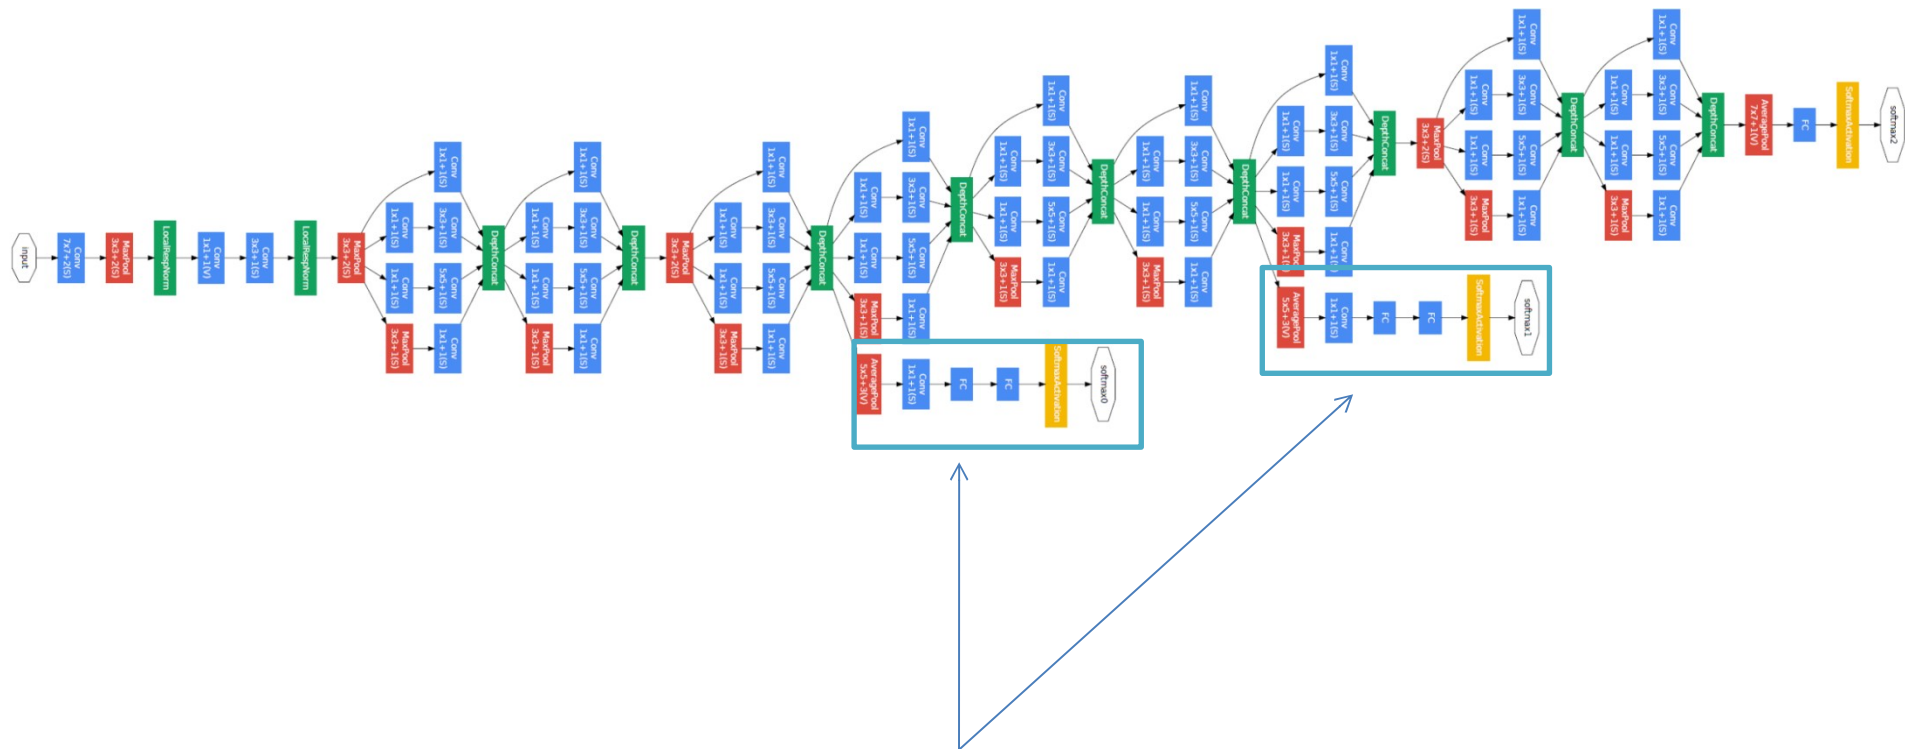
GoogLe Net



Stem network:
Conv – Pool – 2x Conv-Pool

Classifier output
removes FC layers

GoogLe Net



Output layers to inject additional gradient at lower layers
AvgPool – 1x1 Conv – FC – FC - Softmax

GoogLe Net

Stack Inception modules on top of each other

type	patch size/ stride	output size	depth	#1×1	#3×3 reduce	#3×3	#5×5 reduce	#5×5	pool proj	params	ops
convolution	7×7/2	112×112×64	1							2.7K	34M
max pool	3×3/2	56×56×64	0								
convolution	3×3/1	56×56×192	2		64	192				112K	360M
max pool	3×3/2	28×28×192	0								
inception (3a)		28×28×256	2	64	96	128	16	32	32	159K	128M
inception (3b)		28×28×480	2	128	128	192	32	96	64	380K	304M
max pool	3×3/2	14×14×480	0								
inception (4a)		14×14×512	2	192	96	208	16	48	64	364K	73M
inception (4b)		14×14×512	2	160	112	224	24	64	64	437K	88M
inception (4c)		14×14×512	2	128	128	256	24	64	64	463K	100M
inception (4d)		14×14×528	2	112	144	288	32	64	64	580K	119M
inception (4e)		14×14×832	2	256	160	320	32	128	128	840K	170M
max pool	3×3/2	7×7×832	0								
inception (5a)		7×7×832	2	256	160	320	32	128	128	1072K	54M
inception (5b)		7×7×1024	2	384	192	384	48	128	128	1388K	71M
avg pool	7×7/1	1×1×1024	0								
dropout (40%)		1×1×1024	0								
linear		1×1×1000	1							1000K	1M
softmax		1×1×1000	0								



GoogLe Net

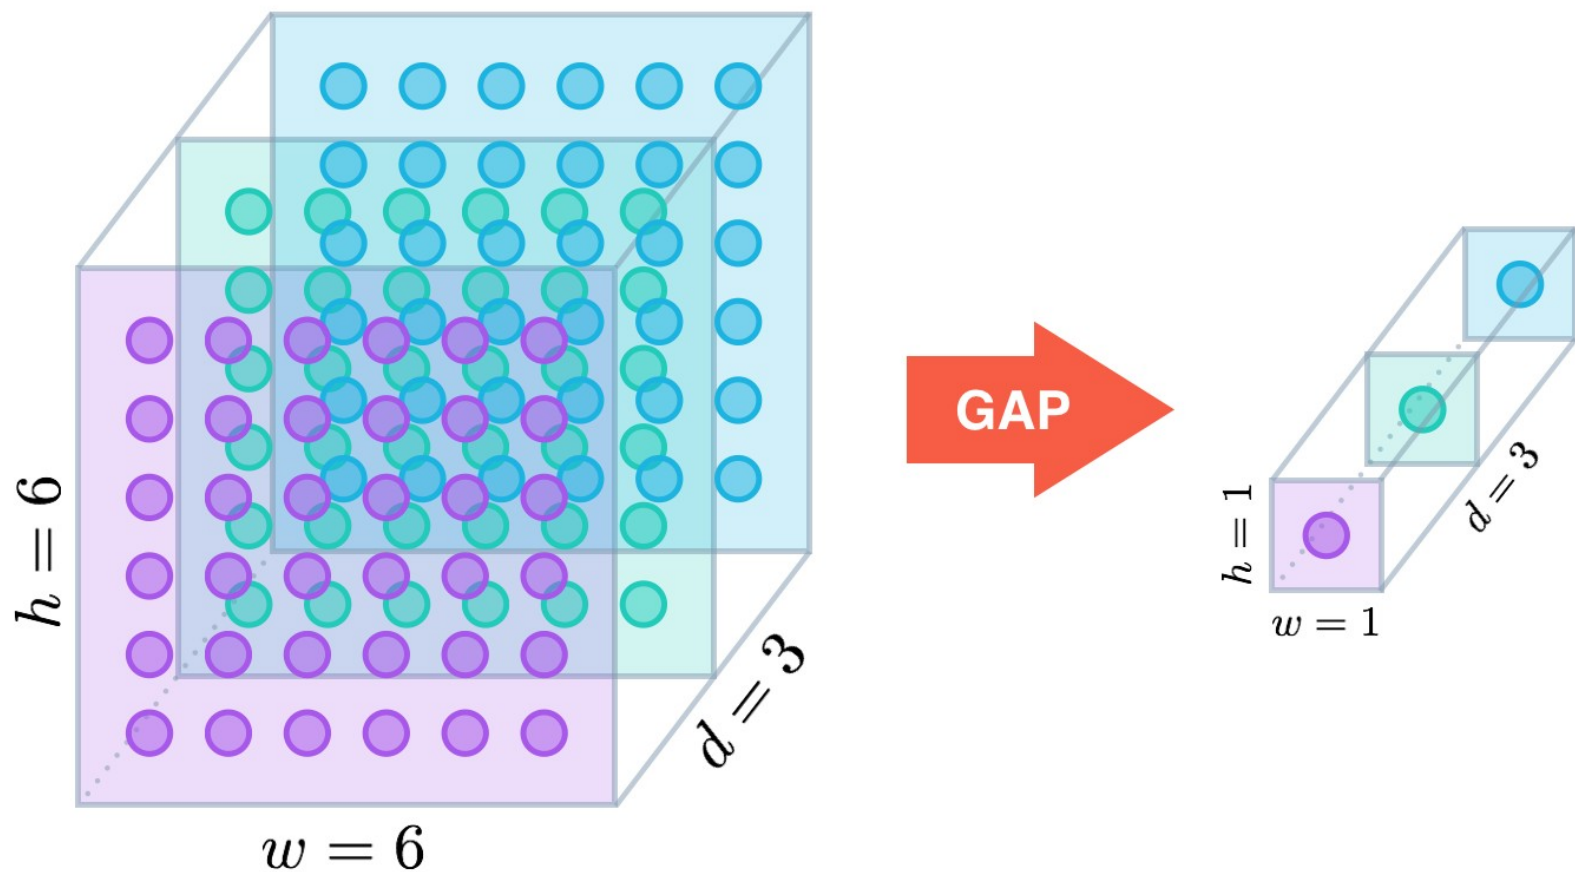
- Removes fully connected layers
- *It was found that a move from fully connected layers to average pooling improved the top-1 accuracy by about 0.6%.*

Global average pooling

Network In Network, 2014

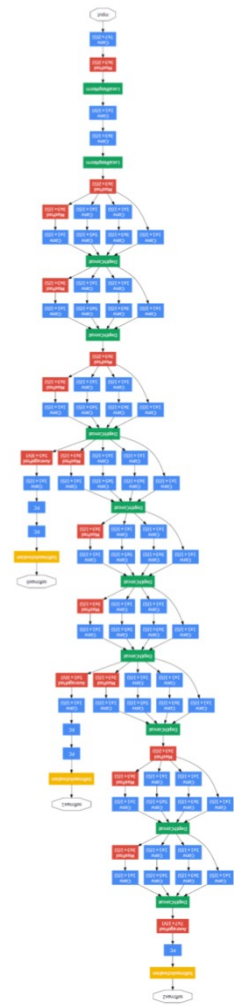
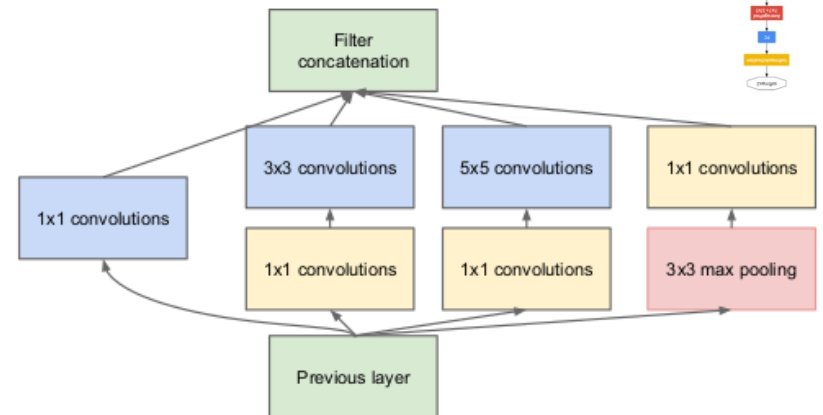
- Instead of adding fully connected layers on top of the feature maps, take **the average** of each feature map, and the resulting vector is fed directly into the softmax layer.
 - is more native to the convolution structure by enforcing correspondences between feature maps and categories (feature maps can be easily interpreted as categories confidence maps)
 - no parameter to optimize in the global average pooling thus overfitting is avoided at this layer
 - sums up spatial information -> more robust to spatial translations of the input

Global average pooling (GAP)

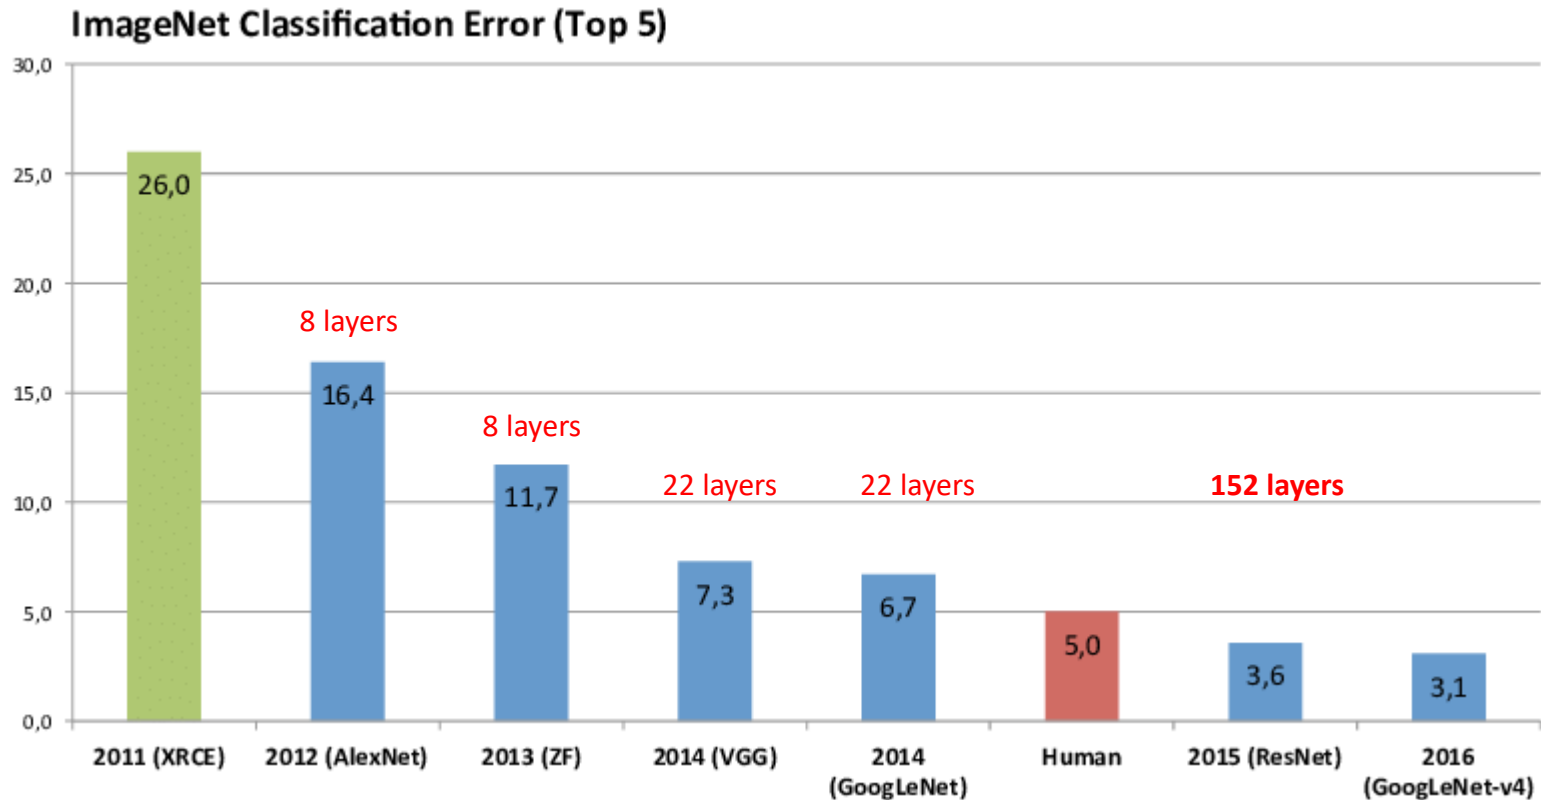


GoogLe Net

- 22 layers
- Inception module
- 12x less parameters than AlexNet
- **Top 5 accuracy: 6.7%**
- Removes fully connected layers as in NIN
 - implementation differs in that we use an extra linear layer



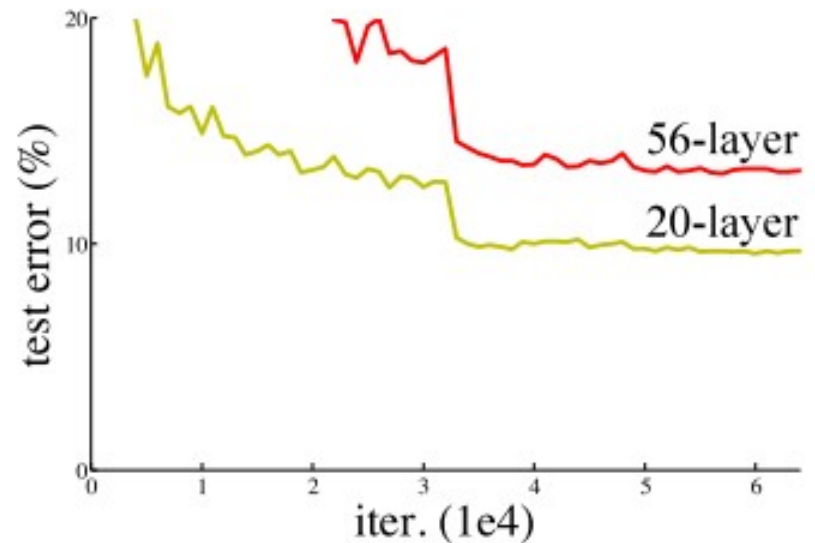
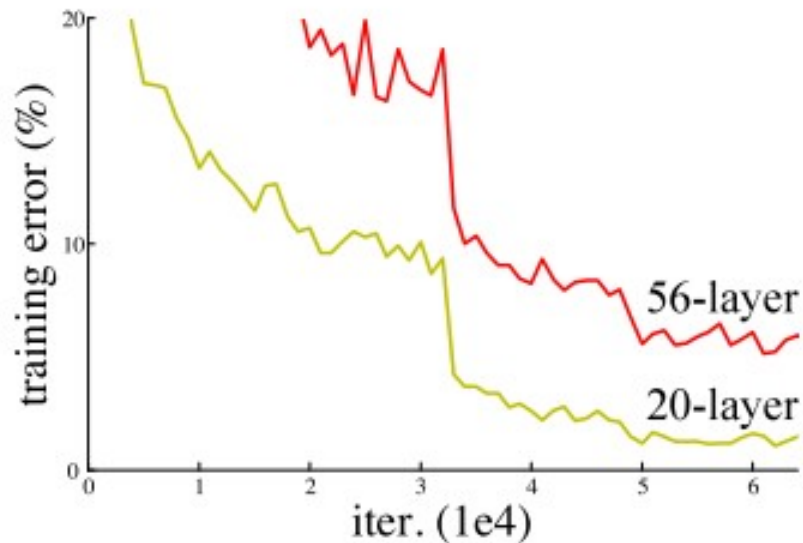
ResNet, 2015



6.7 top-5 accuracy

Res Net

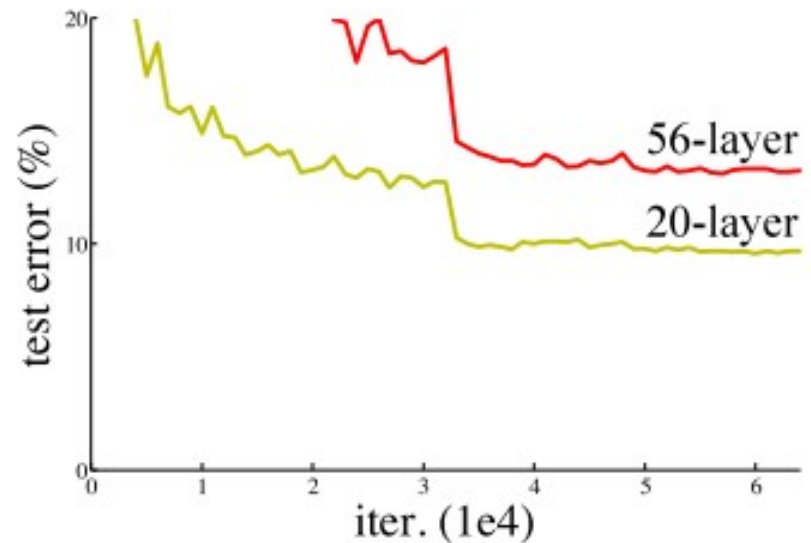
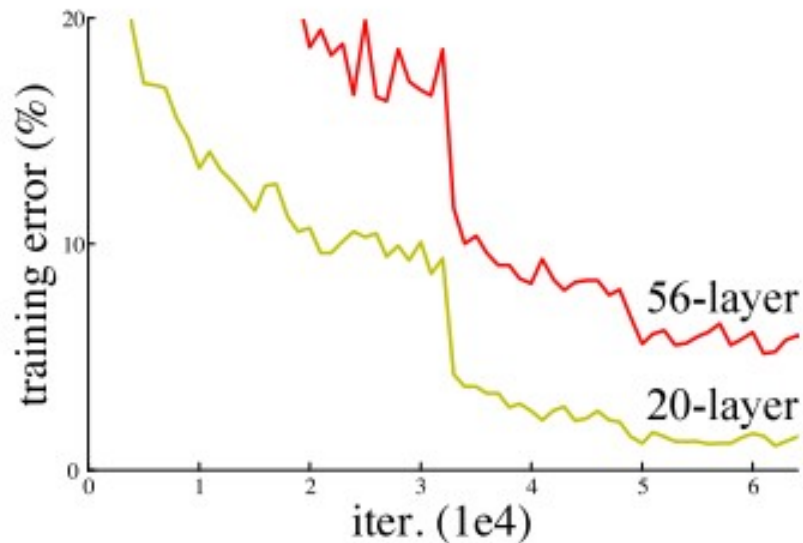
Increasing a network depth



Overfitting?

Res Nets

Increasing a network depth



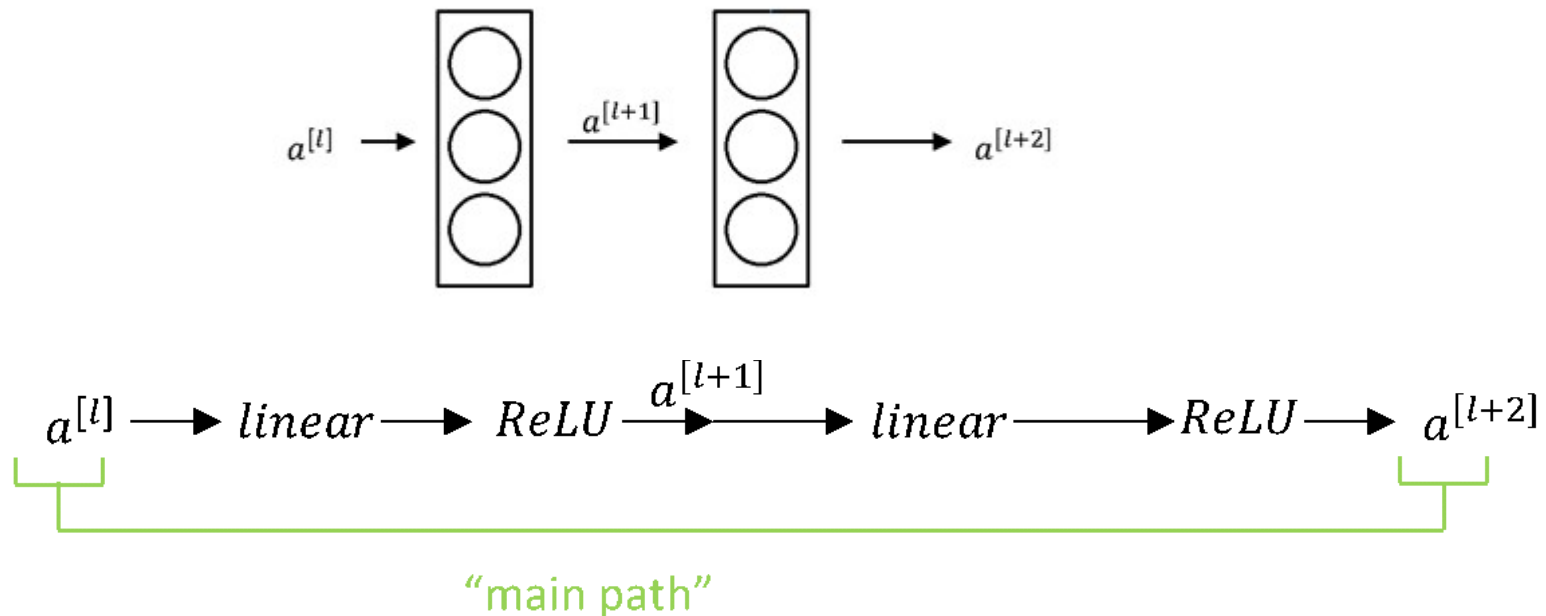
Overfitting? NO!

Hypothesis: Deeper models are harder to optimize

Res Nets

- Deeper models shouldn't hurt performance, they should be at least as good as shallower ones
- Copy the learned layers from the shallower layer and set additional layers to identity mapping

Classical neural network block



$$a^{[l]} \rightarrow \text{linear} \rightarrow \text{ReLU} \rightarrow a^{[l+1]} \rightarrow \text{linear} \rightarrow \text{ReLU} \rightarrow a^{[l+2]}$$

$$z^{[l+1]} = W^{[l+1]}a^{[l]} + b^{[l+1]}$$

$$a^{[l+1]} = g(z^{[l+1]})$$

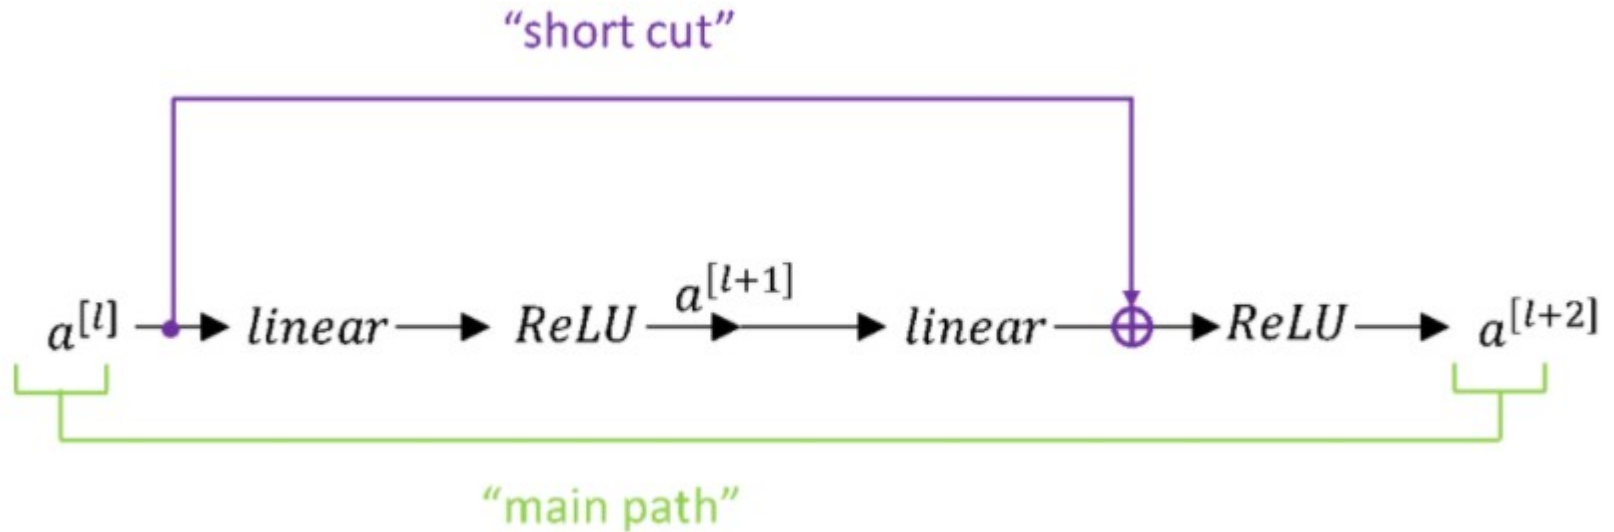
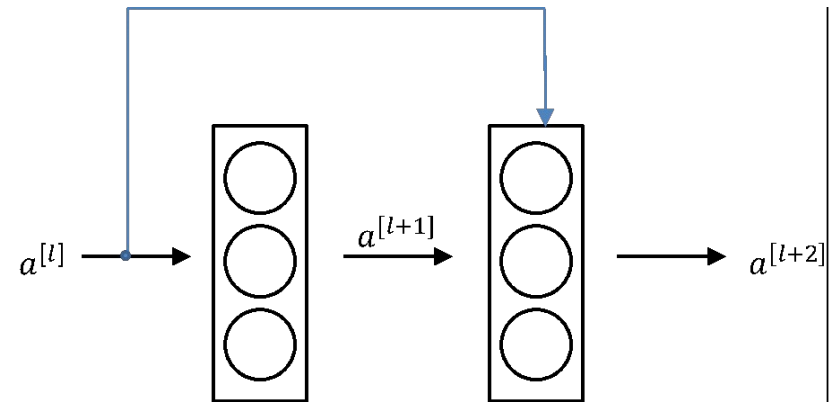
$$z^{[l+2]} = W^{[l+2]}a^{[l+1]} + b^{[l+2]}$$

$$a^{[l+2]} = g(z^{[l+2]})$$

Image source: <http://datahacker.rs/deep-learning-residual-networks/>

Example credit: Andrew Ng

Residual block



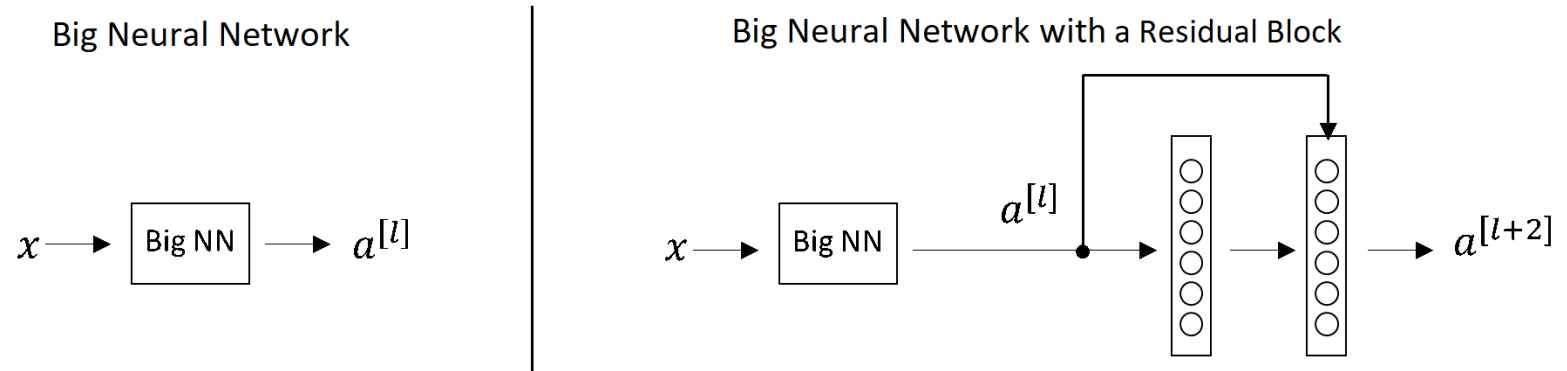
$$z^{[l+1]} = W^{[l+1]}a^{[l]} + b^{[l+1]}$$

$$a^{[l+1]} = g(z^{[l+1]})$$

$$z^{[l+2]} = W^{[l+2]}a^{[l+1]} + b^{[l+2]}$$

$$a^{[l+2]} = g(z^{[l+2]} + a^{[l]})$$

Residual block



Equations for the neural network with a residual block are:

$$a^{[l+2]} = g(z^{[l+2]} + a^{[l]})$$

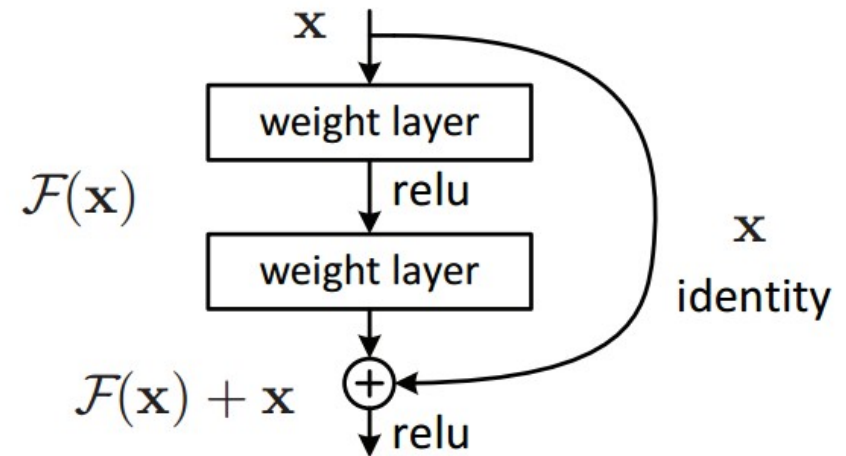
$$a^{[l+2]} = g(W^{[l+2]}a^{[l+1]} + b^{[l+2]} + a^{[l]})$$

If we have $W^{[l+2]} = 0$ and $b = 0$ then:

$$a^{[l+2]} = g(W^{[l+2]}a^{[l+1]} + b^{[l+2]} + a^{[l]}) = g(a^{[l]}) = a^{[l]}$$

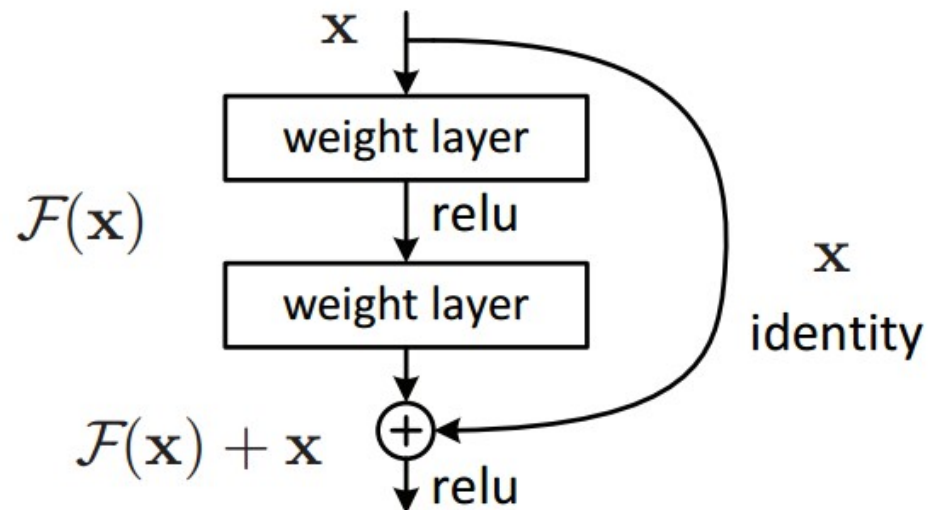
Res Net

- Try to fit a residual mapping instead of directly trying to fit a desired underlying mapping



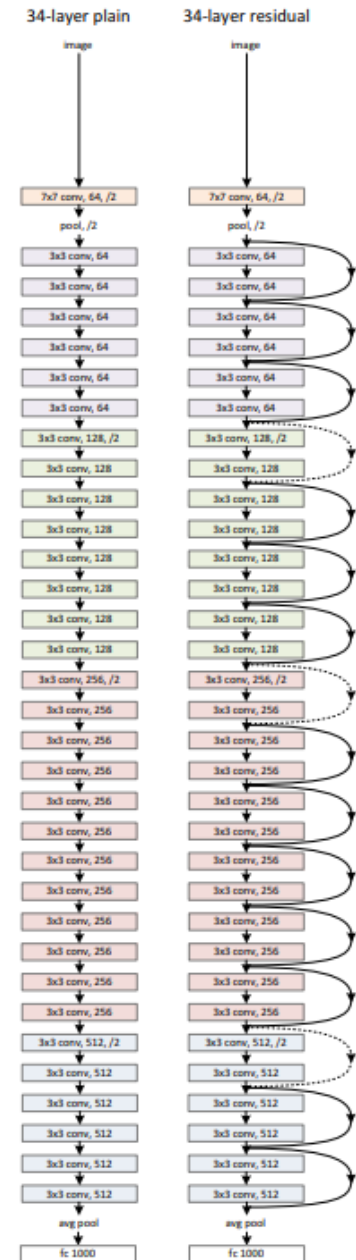
Residual blocks

- Residual blocks
 - Main path
 - Shortcut



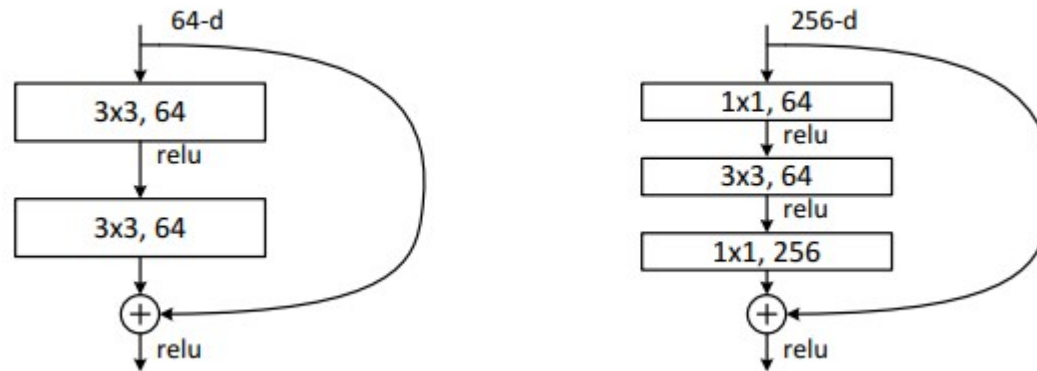
Res Net

- Start with a CONV layer
- Stack multiple residual blocks on top of each other
 - Every residual block has two 3x3 conv layers
 - Periodically, double number of filters and downsample spatially using stride 2
- No fully connected layers



Res Net > 50 layers

- Use “bottlenecks” just like in Inception networks



A deeper residual function F for ImageNet.
Left: a building block (on 56×56 feature maps) as in Fig. 3 for ResNet34. Right: a “bottleneck” building block for ResNet-50/101/152.

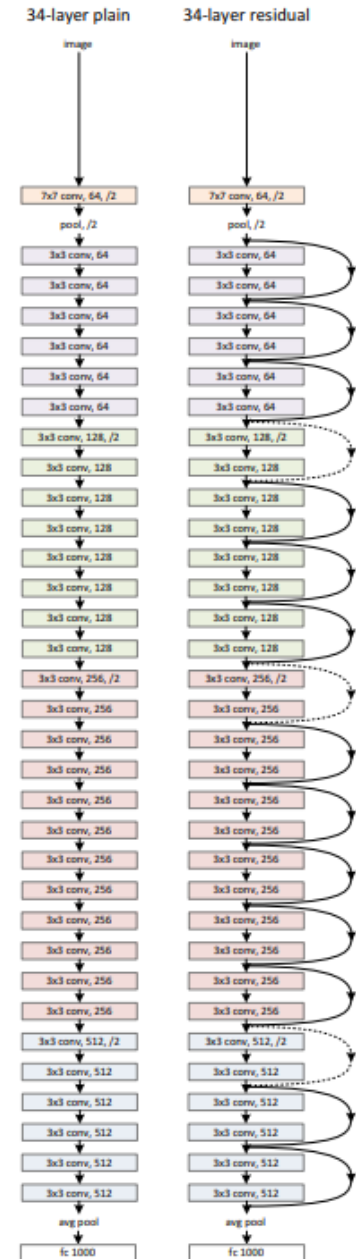
Res Net

- Allows us to train really deep networks
 - “Identity Mappings in Deep Residual Networks”, 2016 – trained a **1001-layer deep** ResNet to outperform its shallower counterparts
- Helps us with vanishing and exploding gradients
- It is easy for ResNet to learn the identity function

Res Net

1st places on:

- ImageNet detection
- ImageNet localization
- COCO detection
- COCO segmentation



Res Net

- Batch Normalization after every CONV layer
- He initialization
- Momentum gradient descent (0.9)
- Learning rate: 0.1, divided by 10 when validation error plateaus
- Mini-batch size 256
- Weight decay of $1e-5$
- No dropout used

Conv nets – the big picture

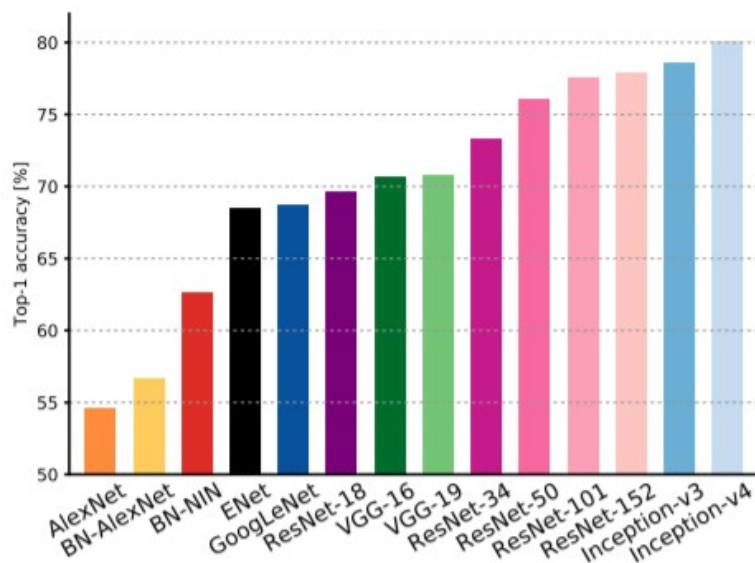


Figure 1: **Top1 vs. network.** Single-crop top-1 validation accuracies for top scoring single-model architectures. We introduce with this chart our choice of colour scheme, which will be used throughout this publication to distinguish effectively different architectures and their correspondent authors. Notice that networks of the same group share the same hue, for example ResNet are all variations of pink.

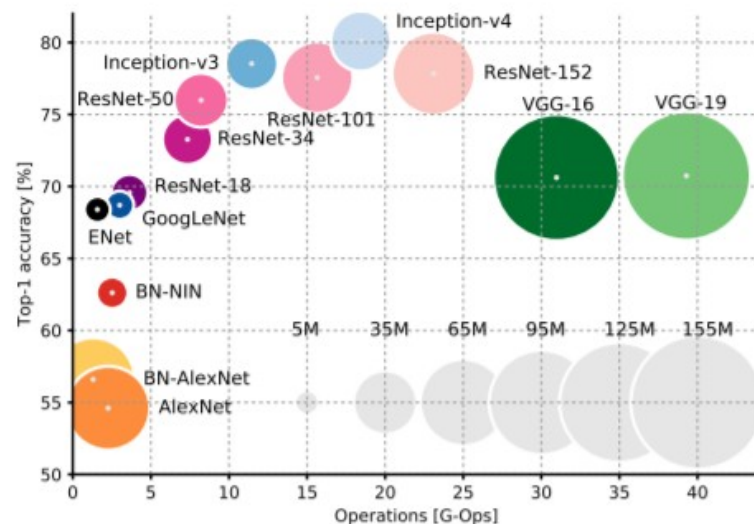
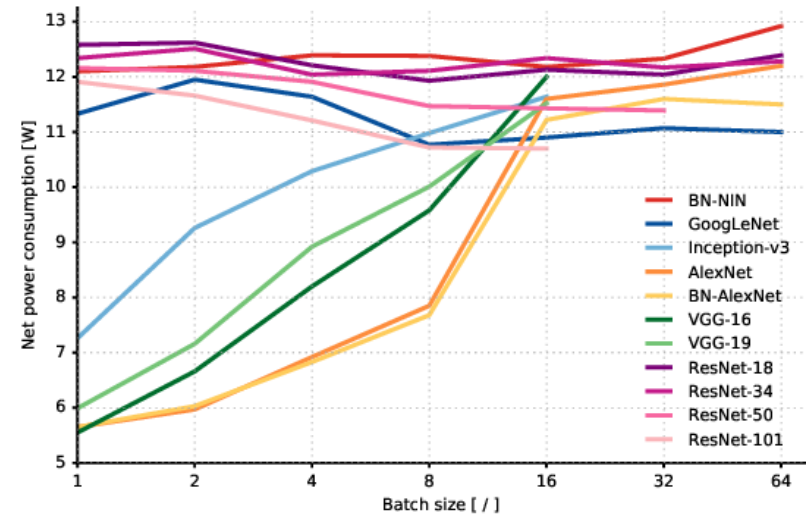
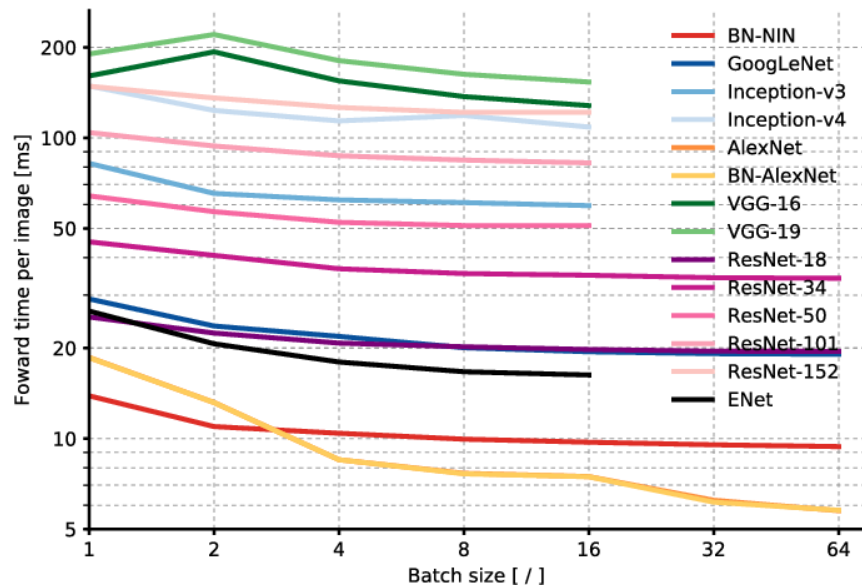


Figure 2: **Top1 vs. operations, size \propto parameters.** Top-1 one-crop accuracy versus amount of operations required for a single forward pass. The size of the blobs is proportional to the number of network parameters; a legend is reported in the bottom right corner, spanning from 5×10^6 to 155×10^6 params. Both these figures share the same y-axis, and the grey dots highlight the centre of the blobs.

Conv nets – the big picture



How to read a research paper?

1. Read the Title, the abstract and the figures

- get a general sense of the concepts in the paper

2. Read the introduction + conclusions + figures + skim the rest

- author(s) try to summarize their work carefully to clarify for the reviewer why their paper should be accepted for publication

3. Read the paper but skip the math

4. Read the whole thing but skip the parts that don't make sense

- some of it doesn't make sense (it's not unusual), it's okay to skim it initially.